



Why MongoDB

Architecture and Technology Overview

Safe Harbor Statement

This presentation contains “forward-looking statements” within the meaning of Section 27A of the Securities Act of 1933, as amended, and Section 21E of the Securities Exchange Act of 1934, as amended. Such forward-looking statements are subject to a number of risks, uncertainties, assumptions and other factors that could cause actual results and the timing of certain events to differ materially from future results expressed or implied by the forward-looking statements. Factors that could cause or contribute to such differences include, but are not limited to, those identified our filings with the Securities and Exchange Commission. You should not rely upon forward-looking statements as predictions of future events. Furthermore, such forward-looking statements speak only as of the date of this presentation.

In particular, the development, release, and timing of any features or functionality described for MongoDB products remains at MongoDB’s sole discretion. This information is merely intended to outline our general product direction and it should not be relied on in making a purchasing decision nor is this a commitment, promise or legal obligation to deliver any material, code, or functionality. Except as required by law, we undertake no obligation to update any forward-looking statements to reflect events or circumstances after the date of such statements.

MongoDB's **Intelligent Data Platform**

is the best way of making your data simple
to **Organize, Use, & Enrich**
Anywhere

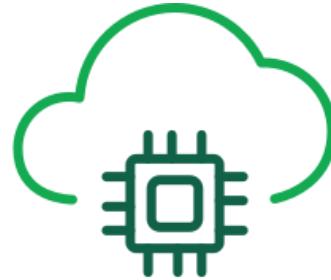
Data Is Everywhere

\$4.5tn
eCommerce sales



by 2021

7.5bn connected IoT devices



by 2025

83% see AI as strategic priority



in 2019

20x faster, 120x lower latency



by 2021

3.8bn smartphone users



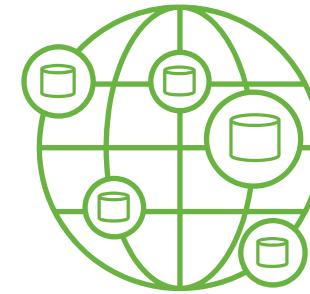
by 2021

\$6tn+ in cyber-crime damage

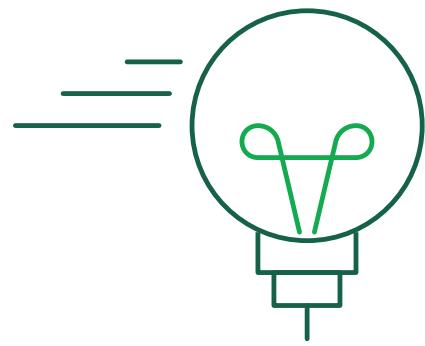


by 2021

Every company
must become a
data &
software
company



Demands from the Business



Innovate Faster

01010101
000101010
010001011
001101010
100101010

Data Driven



Lower Cost

Developers are key....

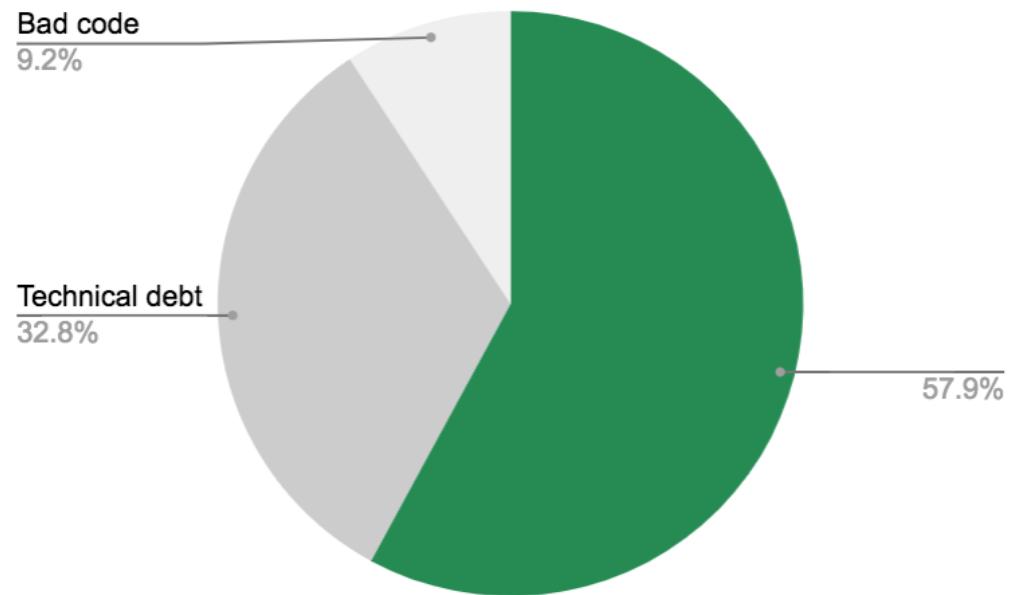
Developers are key....but

Access to developers is a bigger constraint on growth than access to capital, according to survey of thousands of C-level executives.

The Developer Coefficient, Stripe

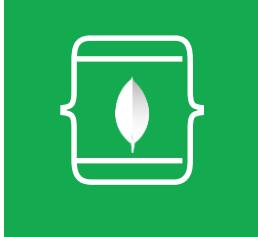
stripe

Developers spend **42% of their work week** on maintenance issues and fixing bad code.



MongoDB: Built for Developer Productivity

Intelligent Data Platform



**Best way to work
with data**



**Intelligently put data
where you need it**



**Freedom
to run anywhere**

The evolution of MongoDB

Document Validation \$lookup Fast Failover Simpler Scalability Aggregation ++ Encryption At Rest In-Memory Storage Engine BI Connector MongoDB Compass APM Integration Auto Index Builds Backups to File System	Linearizable reads Intra-cluster compression Read only views Log Redaction Graph Processing Decimal Collations Faceted Navigation Aggregation ++ Auto-balancing ++ ARM, Power, zSeries BI & Spark Connectors ++ Compass ++ LDAP Authorization Encrypted Backups Cloud Foundry Integration	Change Streams Retryable Writes Expressive Array Updates Query Expressivity Causal Consistency Consistent Sharded Sec. Reads Ops Manager ++ Query Advisor Schema Validation End to End Compression IP Whitelisting Default Bind to Localhost Sessions WiredTiger 1m+ Collections Expressive \$lookUp R Driver Atlas Cross Region Replication Atlas Auto Storage Scaling	Replica Set Transactions Atlas Global Clusters Atlas HIPAA Atlas LDAP Atlas Audit Atlas Enc. Storage Engine Atlas Backup Snapshots Type Conversions 40% Faster Shard Migrations Snapshot Reads Non-Blocking Sec. Reads SHA-2 TLS 1.1+ Compass Agg Pipeline Builder Compass Export to Code Charts Beta Free Monitoring Cloud Service Ops Manager K8s Beta MongoDB Stitch GA MongoDB Mobile Beta	Distributed Transactions Global Point in Time Reads Large Transactions Mutable Shard Key Values Atlas Data Lake (Beta) Atlas Auto Scaling (Beta) Atlas Full-Text Search (Beta) Atlas ISO Compliance Atlas Service Broker Field Level Encryption (Beta) Multi-CAs & Online Rotation On-Demand Materialized Views Wildcard Indexes Agg Pipeline ++ Expressive Updates Apache Kafka Connector (Beta) MongoDB Charts GA Retryable Reads & Writes New Index Builds 10x Faster stepDown Storage Node Watchdog Zstandard Compression Ops Manager Headless Backup Ops Manager K8s GA Ops Manager Single Agent
--	--	--	---	---

3.2

3.4

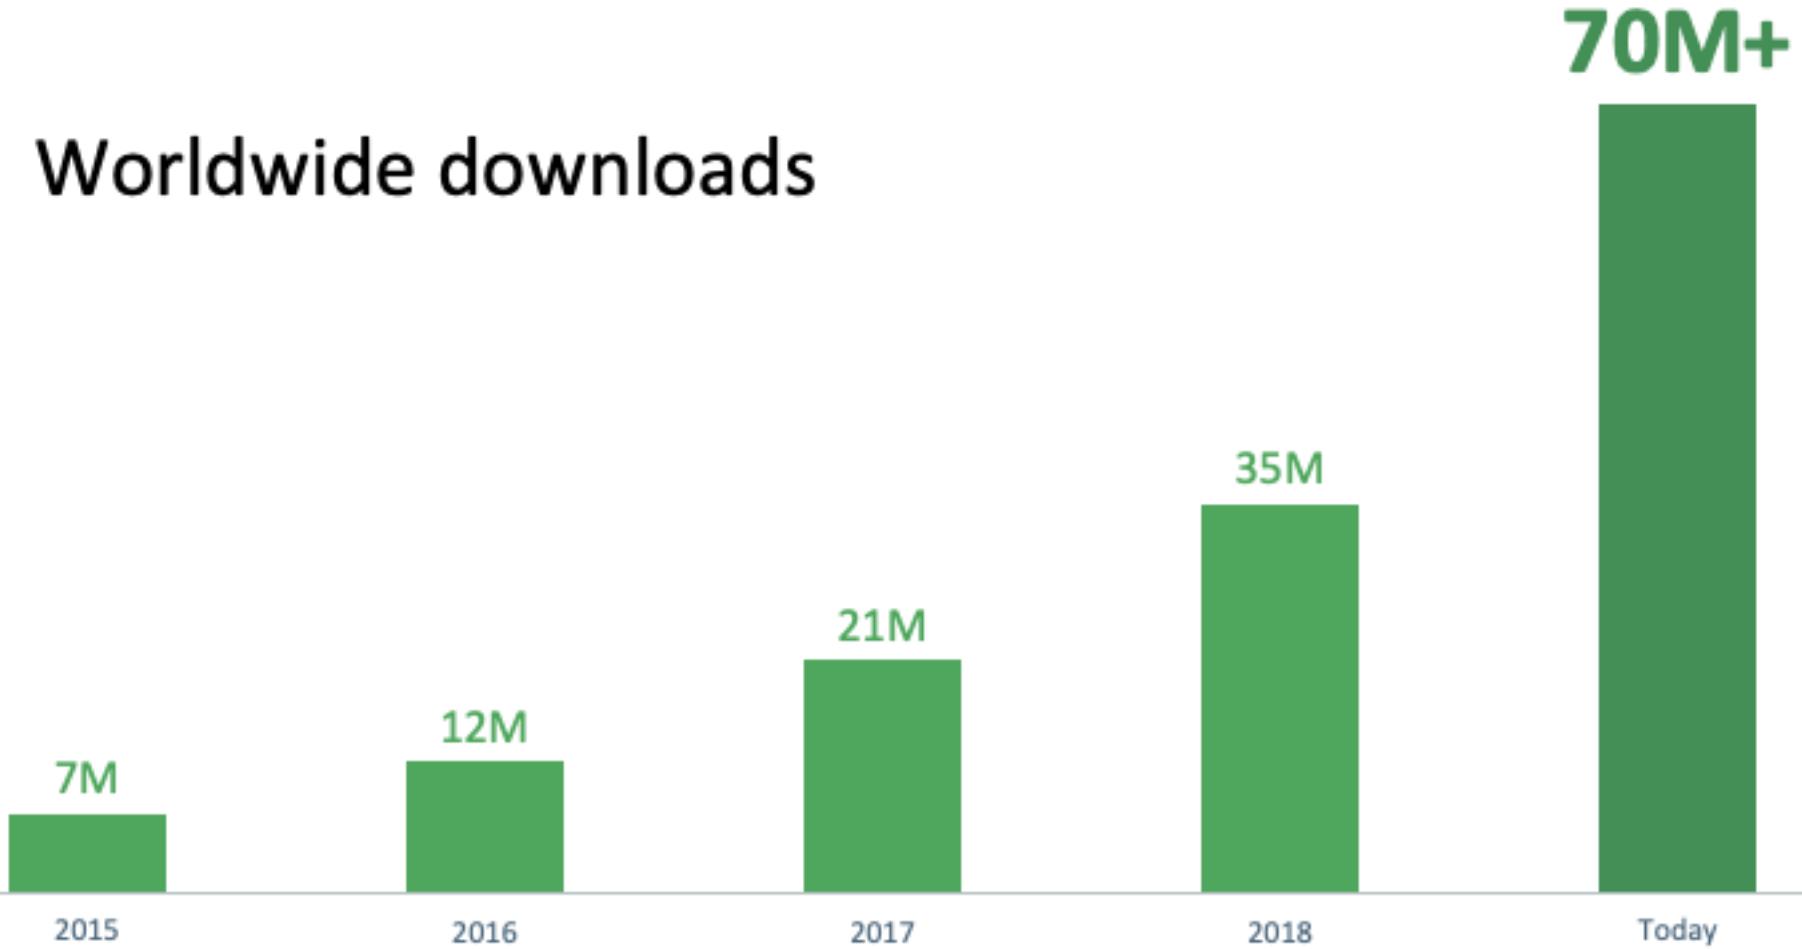
3.6

4.0

4.2

Increasingly the Database of Choice

Worldwide downloads



Stack Overflow Survey: Most Wanted Database



MongoDB Data Platform

Client-Side Database



Mobile



Web



IoT / Embedded (coming)

Application Development



Rules



Serverless Functions



Code Deployment



Sync

Data Layer



MongoDB
Server



MongoDB
Atlas

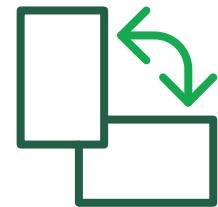
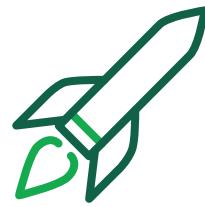
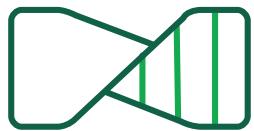


Full-Text
Search



MongoDB Atlas
Data Lake

Best way to work with data



Easy:

Work with data in a natural, intuitive way, fully transactional

Flexible:

Adapt and make changes quickly

Fast:

Get great performance with less code

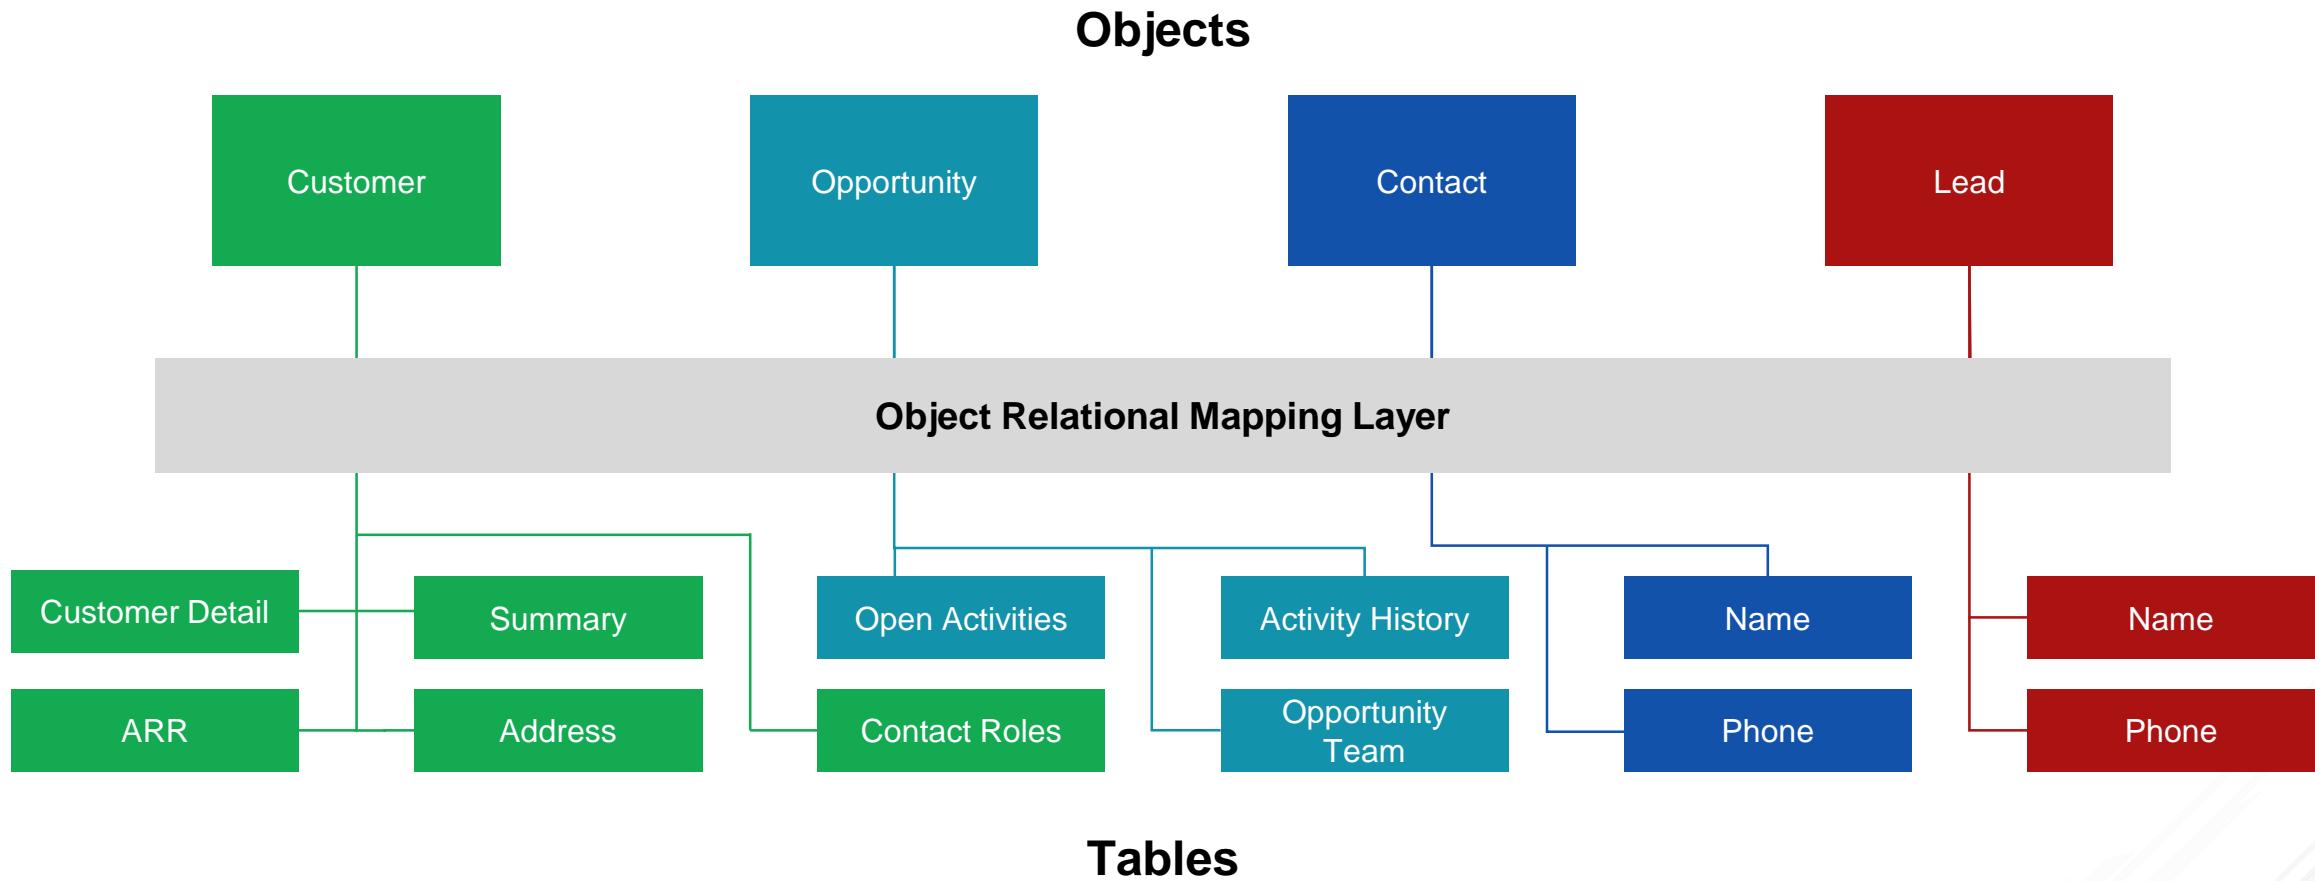
Versatile:

Supports a wide variety of data models and queries

You probably have thousands of tables

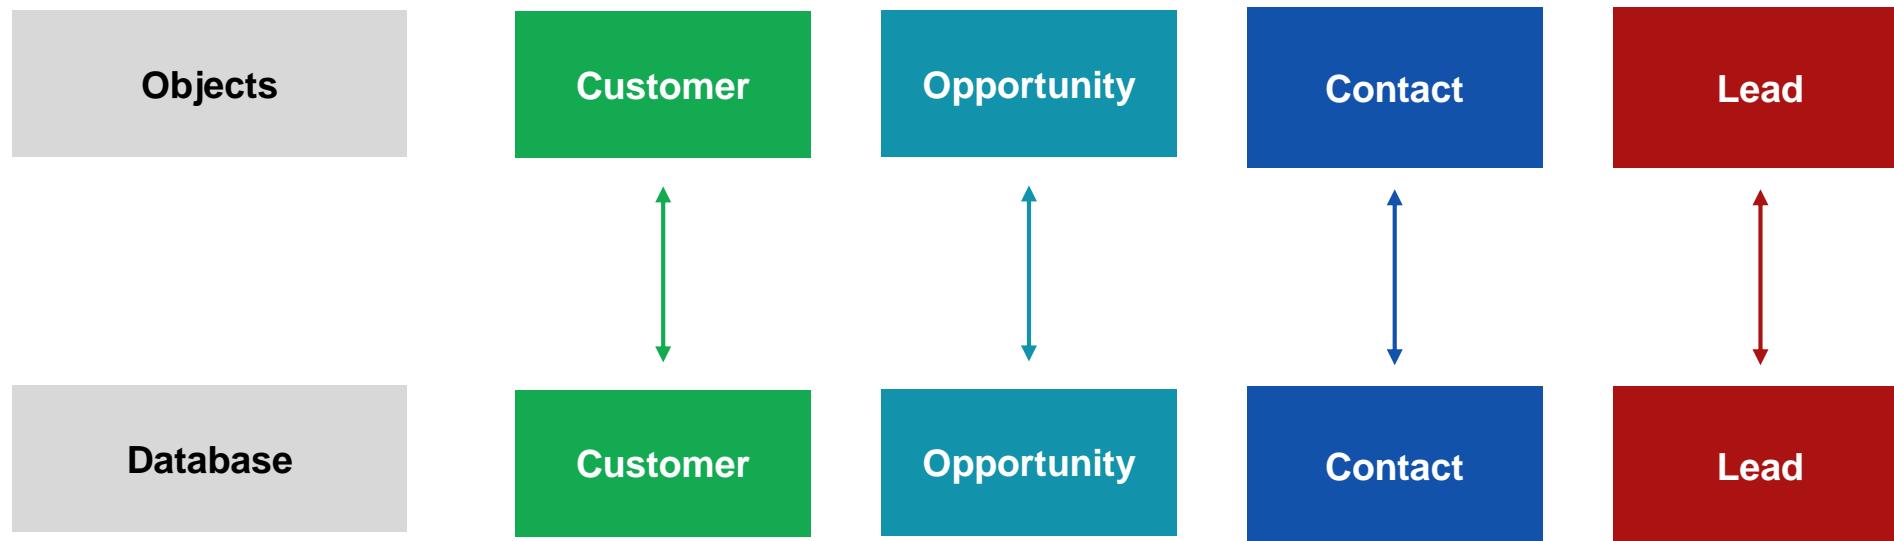


Go from this....





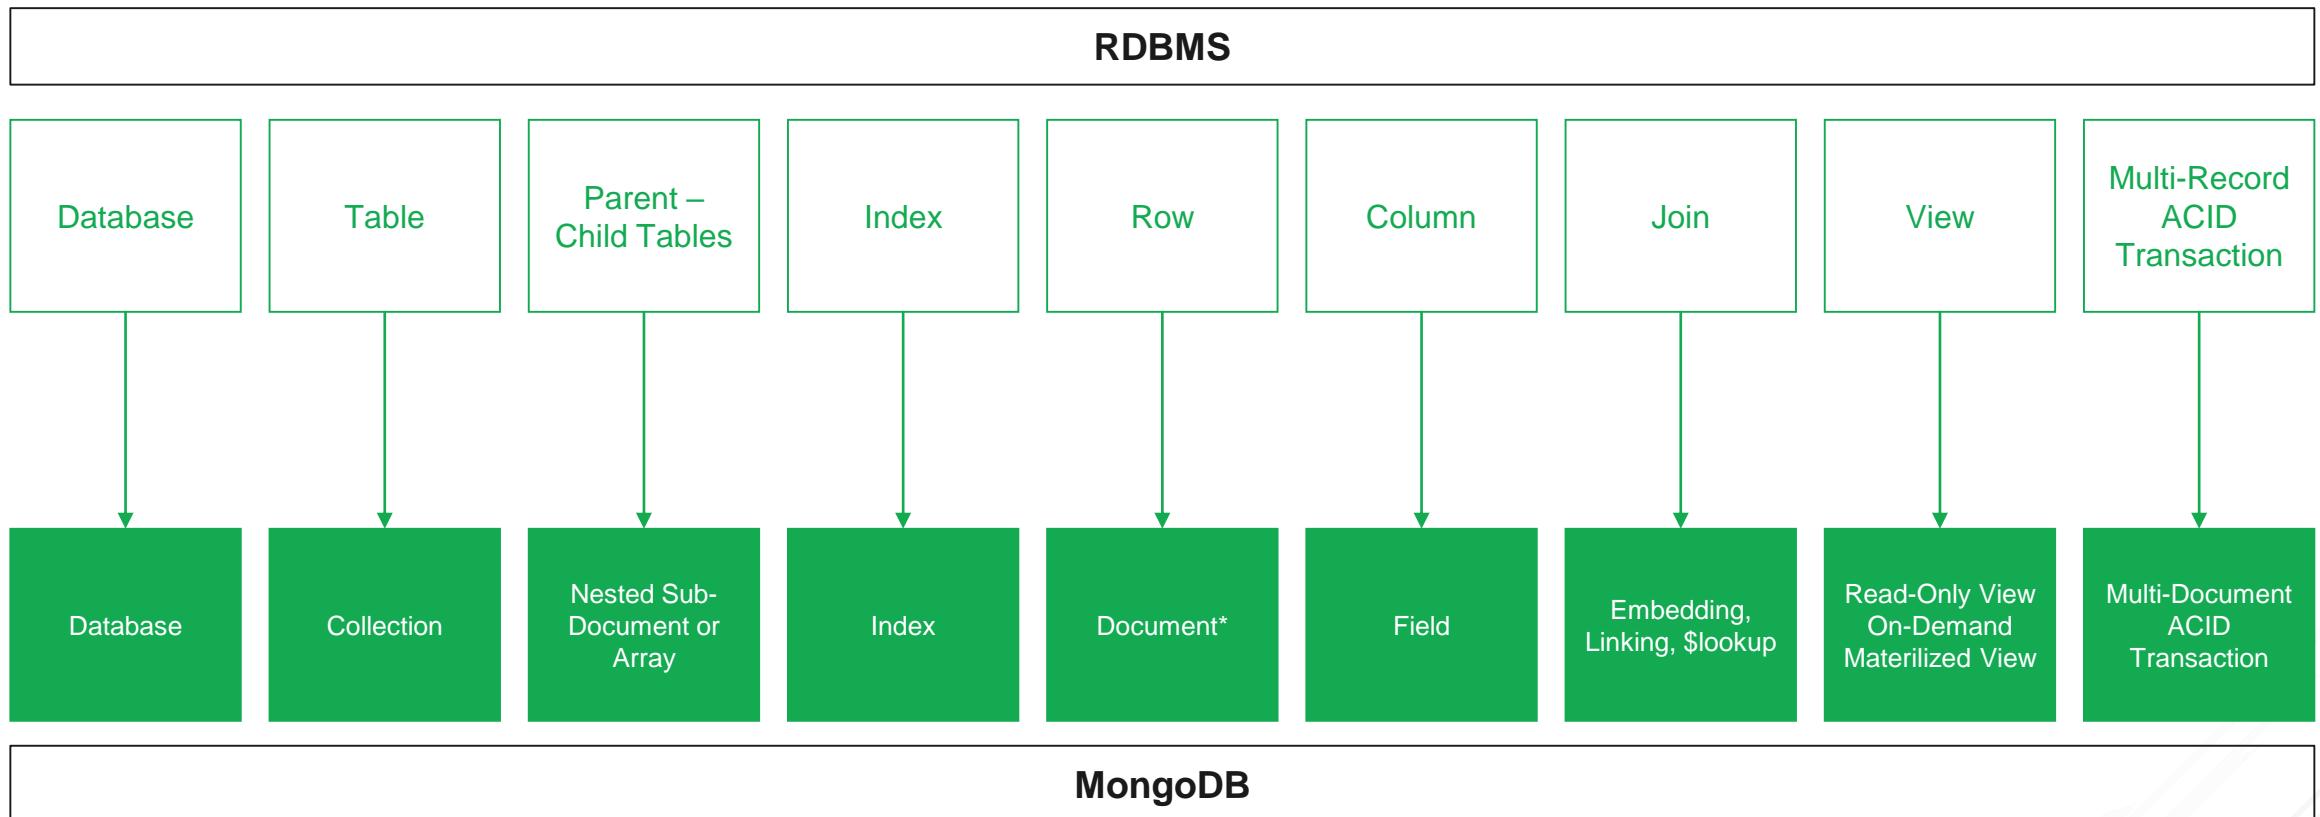
To this: store objects directly...





Some Terminology

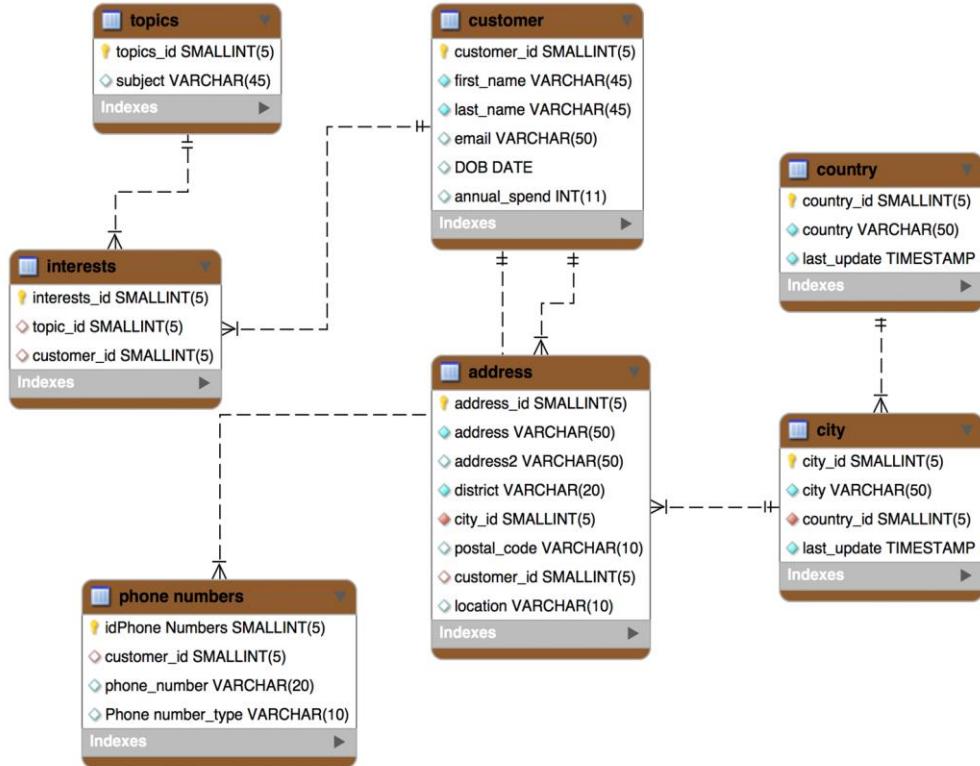
A comparison



* Proper document schema design yields more entity data per document than found in a relational database row



Easy: Contrasting data models



Tabular (Relational) Data Model

Related data split across multiple records and tables

```
{
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),
  "name" : {
    "first" : "John",
    "last" : "Doe" },
  "address" : [
    { "location" : "work",
      "address" : {
        "street" : "16 Hatfields",
        "city" : "London",
        "postal_code" : "SE1 8DJ" },
      "geo" : { "type" : "Point", "coord" : [
        51.5065752,-0.109081]}},
    +   { ... }
  ],
  "phone" : [
    { "location" : "work",
      "number" : "+44-1234567890"},
    +   { ... }
  ],
  "dob" : ISODate("1977-04-01T05:00:00Z"),
  "retirement_fund" : NumberDecimal("1292815.75")
}
```

Document Data Model

Related data contained in a single, rich document



Easy: Document data model

- Naturally maps to objects in code
- Represent data of any structure
- Strongly typed for ease of processing
 - Over 20 binary encoded JSON data types
- Access by idiomatic drivers in all major programming language

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    +   {...}  
  ],  
  "phone" : [  
    { "location" : "work",  
      "number" : "+44-1234567890"},  
    +   {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

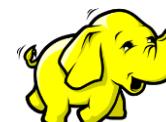


Easy: Drivers and Frameworks

Drivers



Frameworks





Easy: Transactional Data Guarantees

Single & Multi-Document ACID Transactions

For many apps,
single document
transactions meet the
majority of needs

```
_id: 12345678
> name: Object
> address: Array
> phone: Array
email: "john.doe.mongodb.com"
dob: 1966-07-30 01:00:00:000
▼ interests: Array
  0: "Cycling"
  1: "IoT"
```

Related data modeled in a single, rich document against
which ACID guarantees are applied



Easy: MongoDB Multi-Document ACID Transactions

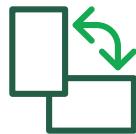


Just like relational transactions

- Multi-statement, familiar relational syntax
- Easy to add to any application
- Multiple documents in 1 or many collections and databases, across replica sets and sharded clusters

ACID guarantees

- Snapshot isolation, all or nothing execution
- No performance impact for non-transactional operations



Syntax

```
with client.start_session() as s:  
  
    s.start_transaction()  
  
    collection_one.insert_one(doc_one, session=s)  
  
    collection_two.insert_one(doc_two, session=s)  
  
    s.commit_transaction()
```

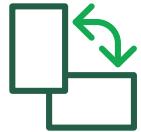
Natural for developers

- Idiomatic to the programming language
- Familiar to relational developers
- Simple



Syntax

```
try (ClientSession clientSession = client.startSession() ) {  
    clientSession.startTransaction();  
  
    collection.insertOne(clientSession, docOne);  
  
    collection.insertOne(clientSession, docTwo);  
  
    clientSession.commitTransaction();  
}
```



Comparing Syntax with Relational



```
db.start_transaction()  
cursor.execute(orderInsert, orderData)  
cursor.execute(stockUpdate, stockData)  
db.commit()
```

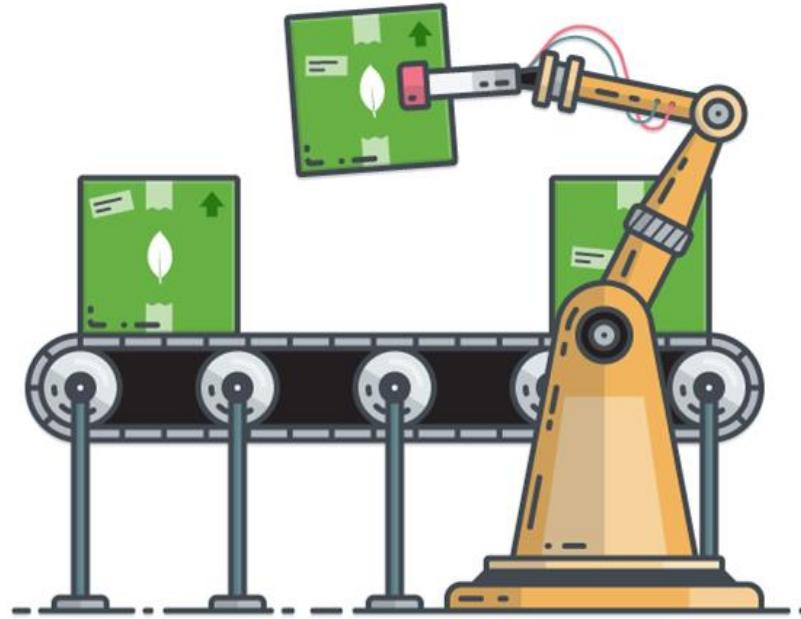


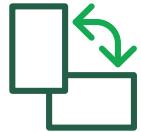
```
s.start_transaction()  
orders.insert_one(order, session=s)  
stock.update_one(item, stockUpdate, session=s)  
s.commit_transaction()
```



Transactions Best Practises

- Use MongoDB 4.2 drivers
- Transactions targeting a single shard will be faster than those spanning multiple shards
- Transactions automatically abort after 60 seconds – TUNABLE. Aborted transactions are fully rolled back by the database
- Best practice: No more than 1,000 documents modified in a single transaction
- No arbiters when running Distributed Transactions
- Chunk migrations will block behind running transactions





The Journey to Transactions

**Major engineering investment over 3+ years
touching every part of the server and drivers**

- Storage layer
- Replication consensus protocol
- Sharding architecture
- Consistency and durability guarantees
- Global logical clock
- Cluster metadata management
- Exposed to drivers through API enhancements





Major Engineering Projects

MongoDB 3.0	MongoDB 3.2	MongoDB 3.4	MongoDB 3.6	MongoDB 4.0	MongoDB 4.2
New Storage engine (WiredTiger)	Enhanced replication protocol: stricter consistency & durability	Shard membership awareness	Consistent secondary reads in sharded clusters	Replica Set Transactions	Global Transactions
	WiredTiger default storage engine		Logical sessions	Make catalog timestamp-aware	Oplog applier prepare support
	Config server manageability improvements		Retryable writes	Snapshot reads	Distributed commit protocol
	Read concern "majority"		Causal Consistency	Recoverable rollback via WT checkpoints	Global point-in-time reads
			Cluster-wide logical clock	Recover to a timestamp	More extensive WiredTiger repair
			Storage API to change to use timestamps	Sharded catalog improvements	Transaction manager
			Read concern majority feature always available		
			Collection catalog versioning		
			UUIDs in sharding		
			Fast in-place updates to large documents in WT		



Learn More

Engineering Chalk and Talks

Hear from the engineers who implemented transactions in MongoDB



WiredTiger timestamps: enforcing correctness in operation ordering across the distributed storage layer.
Hear from Dr. Michael Cahill, Director of Engineering for storage at MongoDB



Logical sessions: coordinating operations across a distributed cluster, presented by Jason Carey, Lead Engineer for platforms at MongoDB



Global logical clock: establishing a global snapshot of



Safe secondary reads: providing consistent reads against

Transactions Page

- Chalk and talks
- Code snippets & documentation
- Blogs



Flexible: Adapt to change

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "phone" : [  
    { "location" : "work",  
      "number" : "+44-1234567890"},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

Add new fields dynamically at runtime



Stepping through the code: Version 1

For each customer store:

- ID
- Name



Version 1: Initial efforts for both technologies

SQL

```
DDL: create table contact ( ... )  
  
init()  
{  
    contactInsertStmt = connection.prepareStatement  
        ("insert into contact ( id, name ) values ( ?,? )");  
    fetchStmt = connection.prepareStatement  
        ("select id, name from contact where id = ?");  
}  
  
save(Map m)  
{  
    contactInsertStmt.setString(1, m.get("id"));  
    contactInsertStmt.setString(2, m.get("name"));  
    contactInsertStmt.executeUpdate();  
}
```

```
Map fetch(String id)  
{  
    Map m = null;  
    fetchStmt.setString(1, id);  
    rs = fetchStmt.executeQuery();  
    if(rs.next()) {  
        m = new HashMap();  
        m.put("id", rs.getString(1));  
        m.put("name", rs.getString(2));  
    }  
    return m;  
}
```

MongoDB

```
DDL: none
```

```
save(Map m)
```

Let's assume for argument's sake that both approaches take the same amount of time

```
Map fetch(String id)  
{  
    Map m = null;  
    DDBObject dbo = new BasicDBObject();  
    dbo.put("id", id);  
    c = collection.find(dbo);  
    if(c.hasNext()) {  
        m = (Map) c.next();  
    }  
    return m;  
}
```



Version 1: Initial efforts for both technologies

SQL

```
DDL: create table contact ( ... )  
  
init()  
{  
    contactInsertStmt = connection.prepareStatement  
    ("insert into contact ( id, name ) values ( ?,? )");  
    fetchStmt = connection.prepareStatement  
    ("select id, name from contact where id = ?");  
}  
  
save(Map m)  
{  
    contactInsertStmt.setString(1, id);  
    contactInsertStmt.setString(2, name);  
    contactInsertStmt.executeUpdate();  
}  
  
Map fetch(String id)  
{  
    Map m = null;  
    fetchStmt.setString(1, id);  
    rs = fetchStmt.executeQuery();  
    if(rs.next()) {  
        m = new HashMap();  
        m.put("id", rs.getString(1));  
        m.put("name", rs.getString(2));  
    }  
    return m;  
}
```

MongoDB

```
DDL: none
```

```
save(Map m)
```

Let's assume for argument's sake that both approaches take the same amount of time

```
Map fetch(String id)  
{  
    Map m = null;  
    DDBObject dbo = new BasicDBObject();  
    dbo.put("id", id);  
    c = collection.find(dbo);  
    if(c.hasNext()) {  
        m = (Map) c.next();  
    }  
    return m;  
}
```



Version 2

New feature!

For each customer store:

- ID
- Name
- Title
- ContactDate

} **New Requirements**



Version 2: Add simple fields

```
m.put("name", "John");  
m.put("id", "K1");  
m.put("title", "Mr.");  
m.put("contactDate", new  
Date(2018, 05, 1));
```

- In the data access layer, the map is easy to change
- ...but now we have to change our persistence code

Brace yourself...



SQL Version 2 (changes in bold)

```
DDL: alter table contact add title varchar(8);
      alter table contact add contactDate date;

init()
{
    contactInsertStmt = connection.prepareStatement
    ("insert into contact ( id, name, title, contactdate ) values ( ?, ?, ?, ? )");
    fetchStmt = connection.prepareStatement
    ("select id, name, title, contactdate from contact where id = ?");
}

save(Map m)
{
    contactInsertStmt.setString(1, m.get("id"));
    contactInsertStmt.setString(2, m.get("name"));
    contactInsertStmt.setString(3, m.get("title"));
    contactInsertStmt.setDate(4, m.get("contactDate"));
    contactInsertStmt.execute();
}

Map fetch(String id)
{
    Map m = null;
    fetchStmt.setString(1, id);
    rs = fetchStmt.execute();
    if(rs.next()) {
        m = new HashMap();
        m.put("id", rs.getString(1));
        m.put("name", rs.getString(2));
        m.put("title", rs.getString(3));
        m.put("contactDate", rs.getDate(4));
    }
    return m;
}
```

Consequences:

1. Code release schedule linked to database upgrade (new code cannot run on old schema)
 2. DAL closely tied to RDBMS columns
 3. Issues with case sensitivity starting to creep in (many RDBMS are case insensitive for column names, but code is case sensitive)
 4. Changes require careful modifications in 4 places
1. Beginning of **technical debt**



SQL Version 2 (changes in bold)

```
DDL: alter table contact add title varchar(8);
      alter table contact add contactDate date;

init()
{
    contactInsertStmt = connection.prepareStatement
        ("insert into contact ( id, name, title, contactdate ) values (
        ?, ?, ?, ? )");
    fetchStmt = connection.prepareStatement
        ("select id, name, title, contactdate from contact where id = ?");
}

save(Map m)
{
    contactInsertStmt.setString(1, m.get("id"));
    contactInsertStmt.setString(2, m.get("name"));
    contactInsertStmt.setString(3, m.get("title"));
    contactInsertStmt.setDate(4, m.get("contactDate"));
    contactInsertStmt.execute();
}

Map fetch(String id)
{
    Map m = null;
    fetchStmt.setString(1, id);
    rs = fetchStmt.execute();
    if(rs.next()) {
        m = new HashMap();
        m.put("id", rs.getString(1));
        m.put("name", rs.getString(2));
        m.put("title", rs.getString(3));
        m.put("contactDate", rs.getDate(4));
    }
    return m;
}
```

Consequences:

1. Code release schedule linked to database upgrade (new code cannot run on old schema)
 2. DAL closely tied to RDBMS columns
 3. Issues with case sensitivity starting to creep in (many RDBMS are case insensitive for column names, but code is case sensitive)
 4. Changes require careful modifications in 4 places
1. Beginning of **technical debt**



MongoDB Version 2

```
save(Map m)
{
    collection.insert(m);
}

Map fetch(String id)
{
    Map m = null;
    DBObject dbo = new BasicDBObject();
    dbo.put("id", id);
    c = collection.find(dbo);
    if(c.hasNext())
        m = (Map) c.next();
    }
    return m;
}
```



NO CHANGE

Advantages:

1. Zero time and money spent on overhead code
2. Code and database not physically linked
3. New material with more fields can be added into existing collections; backfill is optional
4. Names of fields in database precisely match key names in code layer and directly match on name, not indirectly via positional offset
1. No technical debt is created



Version 3

New feature!

For each customer store:

- ID
- Name
- Title
- HireDate
- Zero to n phone numbers



New Requirement



Version 3: Add list of phone numbers

```
m.put("name", "John");
m.put("id", "K1");
m.put("title", "Mr.");
m.put("contactDate", new Date(2018, 05, 1));

n1.put("type", "work");
n1.put("number", "1-800-555-1212"));
list.add(n1);
n2.put("type", "home"));
n2.put("number", "1-866-444-3131"));
list.add(n2);
m.put("phones", list);
```

- It was still pretty easy to add this data to the map
- ... but meanwhile, in the persistence code ...

REALLY brace yourself...

🎥 SQL Version 3

```
DDL: create table phones ( ... )
```

```
init()  
{  
    contactInsertStmt = connection.prepareStatement  
        ("insert into contact ( id, name, title, hiredate ) values ( ?, ?, ?, ? )");  
    c2stmt = connection.prepareStatement("insert into phones (id, type,  
number) values (?, ?, ?);"  
    fetchStmt = connection.prepareStatement  
        ("select id, name, title, hiredate, type, number from contact, phones  
where phones.id = contact.id and contact.id = ?");  
}  
  
save(Map m)  
{  
    startTrans();  
    contactInsertStmt.setString(1, m.get("id"));  
    contactInsertStmt.setString(2, m.get("name"));  
    contactInsertStmt.setString(3, m.get("title"));  
    contactInsertStmt.setDate(4, m.get("contactDate"));  
  
    for(Map onePhone : m.get("phones")) {  
        c2stmt.setString(1, m.get("id"));  
        c2stmt.setString(2, onePhone.get("type"));  
        c2stmt.setString(3, onePhone.get("number"));  
        c2stmt.execute();  
    }  
    contactInsertStmt.execute();  
    endTrans();  
}
```

```
Map fetch(String id)  
{  
    Map m = null;  
    fetchStmt.setString(1, id);  
    rs = fetchStmt.execute();  
    int i = 0;  
    List list = new ArrayList();  
    while (rs.next()) {  
        if(i == 0) {  
            m = new HashMap();  
            m.put("id", rs.getString(1));  
            m.put("name", rs.getString(2));  
            m.put("title", rs.getString(3));  
            m.put("contactDate", rs.getDate(4));  
            m.put("phones", list);  
        }  
        Map onePhone = new HashMap();  
        onePhone.put("type", rs.getString(5));  
        onePhone.put("number", rs.getString(6));  
        list.add(onePhone);  
        i++;  
    }  
    return m;  
}
```



This took time and money



SQL Version 3: With Bugs Removed

```
init()
{
    contactInsertStmt = connection.prepareStatement
        ("insert into contact ( id, name, title, contactdate ) values (
    ?,?,?,? )");
    c2stmt = connection.prepareStatement("insert into phones (id,
    type, number) values (?, ?, ?)");
    fetchStmt = connection.prepareStatement
        ("select A.id, A.name, A.title, A.hiredate, B.type, B.number from
contact A left outer join phones B on (A.id = B. id) where A.id =
?");
}

while (rs.next()) {
    if(i == 0) {
        // ...
    }
    String s = rs.getString(5);
    if(s != null) {
        Map onePhone = new HashMap();
        onePhone.put("type", s);
        onePhone.put("number", rs.getString(6));
        list.add(onePhone);
    }
}
```

Some contacts that have no phone number → outer join.

But this ALSO means we have to change the unwind logic

This took more time and money!



MongoDB Version 3

```
save(Map m)
{
    collection.insert(m);
}

Map fetch(String id)
{
    Map m = null;
    DBObject dbo = new BasicDBObject();
    dbo.put("id", id);
    c = collection.find(dbo);
    if(c.hasNext())
        m = (Map) c.next();
    return m;
}
```



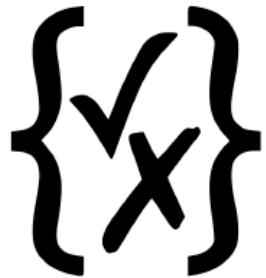
NO CHANGE

Advantages:

1. Zero time and money spent on overhead code
2. Code and database not physically linked
3. New material with more fields can be added into existing collections; backfill is optional
4. Names of fields in database precisely match key names in code layer and directly match on name, not indirectly via positional offset
1. No technical debt is created



Flexible: Govern



JSON Schema

Enforces strict schema structure over a complete collection for data governance & quality

- Builds on document validation introduced by restricting new content that can be added to a document
- Enforces presence, type, and values for document content, including nested array
- Simplifies application logic

Tunable: enforce document structure, log warnings, or allow complete schema flexibility

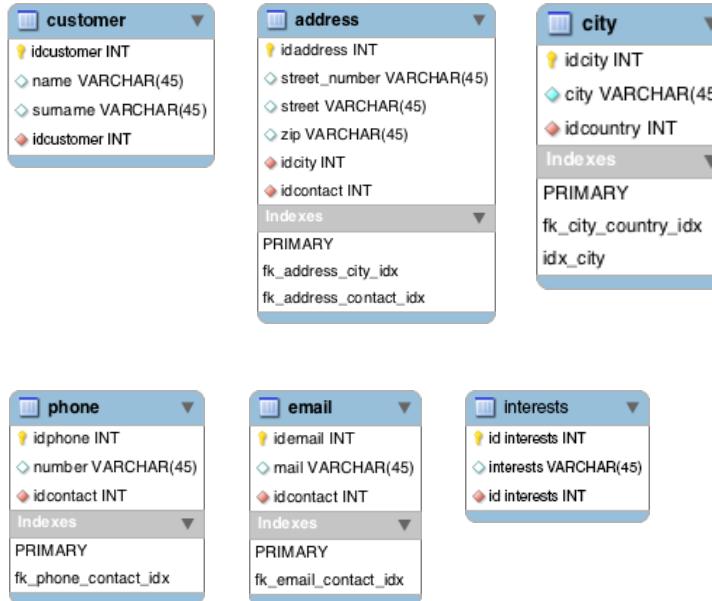
Queryable: identify all existing documents that do not comply



Fast: To work with data

Compared to storing data across multiple tables, a single document data structure:

- Presents a single place for the database to read and write data
- Denormalized data eliminates JOINs for most operational queries
- Simplifies query development and optimization



```
_id: 12345678
> name: Object
> address: Array
> phone: Array
  email: "john.doe@mongodb.com"
  dob: 1966-07-30 01:00:00:000
  < interests: Array
    0: "Cycling"
    1: "IoT"
```



SQL needed to insert a customer

```
import mysql.connector
from mysql.connector import errorcode

def addUser(connection, user):
    cursor = connection.cursor()

    customerInsert = (
        "INSERT INTO customer (first_name, last_name,
                           email,
                           DOB, annual_spend) VALUES "
        "('%(first)s', '%(last)s', %(email)s, %(dob)s, %(spend)
         s)")

    customerData = {
        'first': user['name']['first'],
        'last': user['name']['second'],
        'email': user['email'],
        'dob': user['dob'],
        'spend': user['annualSpend']
    }

    cursor.execute(customerInsert, customerData)
    customerId = cursor.lastrowid

    cityQuery = ("SELECT city_id FROM city WHERE city = %(city)s")
    for address in user['address']:
        cursor.execute(cityQuery, {'city': address['city']})
        city_id = cursor.fetchone()[0]

        addressInsert = (
            "INSERT INTO address (address, address2,
                               district,
                               city_id, postal_code, customer_id, location) "
            "VALUES (%(add)s, %(add2)s, %(dist)s, %(city)s
                  , %(post)s, %(cust)s, %(loc)s)"
```

```
addressData = {
    'add': address['number'],
    'add2': address['street'],
    'dist': address['state'],
    'city': city_id,
    'post': address['postalCode'],
    'cust': customerId,
    'loc': address['location']
}

cursor.execute(addressInsert, addressData)

topicQuery = ("SELECT topics_id FROM topics WHERE
               subject = %(subj)s")
interestInsert = (
    "INSERT into interests (topic_id, customer_id) "
    "VALUES (%(topic)s, %(cust)s)")
```

```
for interest in user['interests']:
    topicId = 0
    topicData = {
        'subj': interest['interest']
    }

    cursor.execute(topicQuery, topicData)
    row = cursor.fetchone()
    if row is None:
        topicInsert = ("INSERT INTO topics (subject)
                       VALUES (%(subj)s)")
        cursor.execute(topicInsert, topicData)
        topicId = cursor.lastrowid
    else:
        topicId = row[0]

    interestData = {
        'topic': topicId,
        'cust': customerId
    }
    cursor.execute(interestInsert, interestData)

phoneInsert = (
    "INSERT INTO `phone numbers` (customer_id,
                                  phone_number, `Phone number_type`)"
    "VALUES (%(cust)s, %(num)s, %(type)s)")
for phoneNumber in user['phone']:
    phoneData = {
        'cust': customerId,
        'num': phoneNumber['number'],
        'type': phoneNumber['location']
    }
    cursor.execute(phoneInsert, phoneData)

connection.commit()
cursor.close()
return customerId
```

Code snippet:
<https://git.io/vpxnx>



Fast: MongoDB requires just 2 lines of code

```
import mysql.connector
from mysql.connector import errorcode

def addUser(connection, user):
    cursor = connection.cursor()

    customerInsert = (
        "INSERT INTO customer (first_name, last_name,
                           email,
                           DOB, annual_spend) VALUES "
        "({first}s, {last}s, {email}s, {dob}s, {spend}
         s)")

    customerData = {
        'first': user['name']['first'],
        'last': user['name']['second'],
        'email': user['email'],
        'dob': user['dob'],
        'spend': user['annualSpend']
    }

    cursor.execute(c
customerId = cur
cityQuery = ("SE
    city)s")
for address in u
cursor.execu
)
city_id = cursor.fetchone()['
addressInsert = (
    "INSERT INTO address (address, address2,
                           district,
                           city_id, postal_code, customer_id, location)"
    "VALUES ({add1}s, {add2}s, {dist}s, {city}s
             , {post}s, {cust}s, {loc}s)")
addressData = {
    'add': address['number'],
    'add2': address['street'],
    'dist': address['state'],
    'city': city_id,
    'post': address['postalCode'],
    'cust': customerId,
    'loc': address['location']
}
cursor.execute(addressInsert, addressData)

topicQuery = ("SELECT topics_id FROM topics WHERE
               subject = %(subj)s")
interestInsert = (
    "INSERT into interests (topic_id, customer_id) "
    "VALUES (%(topic)s, %(cust)s)")

for interest in user['interests']:
    topicId = 0
    topicData = {
        'subj': interest['interest']
    }

    cursor.execute(topicQuery, topicData)
    row = cursor.fetchone()
    if row is None:
        topicInsert = ("INSERT INTO topics (subject)
                       VALUES (%(subj)s)")
        cursor.execute(topicInsert, topicData)
        topicId = cursor.lastrowid
    else:
        topicId = row[0]

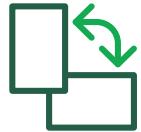
    interestData = {
        'topic': topicId,
        'cust': customerId
    }
    cursor.execute(interestInsert, interestData)

    t, interestData)

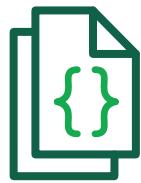
    (customer_id,
     per_type) "
     %(type)s"))
    :
    er'],
    type * phone_number + location')
}
cursor.execute(phoneInsert, phoneData)

connection.commit()
cursor.close()
return customerId
```

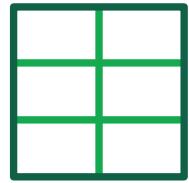
Code snippet: <https://git.io/vpnpG>



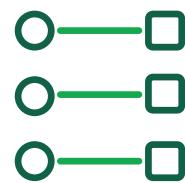
Versatile: Multiple data models, rich query functionality



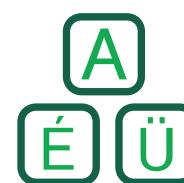
JSON Documents



Tabular



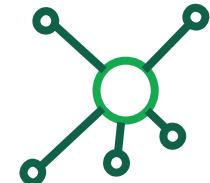
Key-Value



Text



Geospatial



Graph

Rich Queries

Point | Range | Geospatial | Faceted Search | Aggregations | JOINs | Graph Traversals



Versatile: Rich query functionality

Expressive Queries	<ul style="list-style-type: none">Find anyone with phone # “1-212...”Check if the person with number “555...” is on the “do not call” list
Geospatial	<ul style="list-style-type: none">Find the best offer for the customer at geo coordinates of 42nd St. and 6th Ave
Text Search	<ul style="list-style-type: none">Find all tweets that mention the firm within the last 2 days
Aggregation	<ul style="list-style-type: none">Count and sort number of customers by city, compute min, max, and average spend
Native Binary JSON Support	<ul style="list-style-type: none">Add an additional phone number to Mark Smith’s record without rewriting the documentUpdate just 2 phone numbers out of 10Sort on the modified date
JOIN (\$lookup)	<ul style="list-style-type: none">Query for all San Francisco residences, lookup their transactions, and sum the amount by person
Graph Queries (\$graphLookup)	<ul style="list-style-type: none">Query for all people within 3 degrees of separation from Mark

MongoDB

```
{   customer_id : 1,  
    first_name : "Mark",  
    last_name : "Smith",  
    city : "San Francisco",  
    phones: [      {  
        number : "1-212-777-1212",  
        type : "work"  
    },  
    {  
        number : "1-212-777-1213",  
        type : "cell"  
    }]  
.... . . . }
```



Fully Indexable

Fully featured secondary indexes

Index Types

- Primary Index
 - Every Collection has a primary key index
- Compound Index
 - Index against multiple keys in the document
- MultiKey Index
 - Index into arrays
- Wildcard Index
 - Auto-index all matching fields, sub-documents & arrays
- Text Indexes
 - Support for text searches
- GeoSpatial Indexes
 - 2d & 2dSphere indexes for spatial geometries
- Hashed Indexes
 - Hashed based values for sharding

Index Features

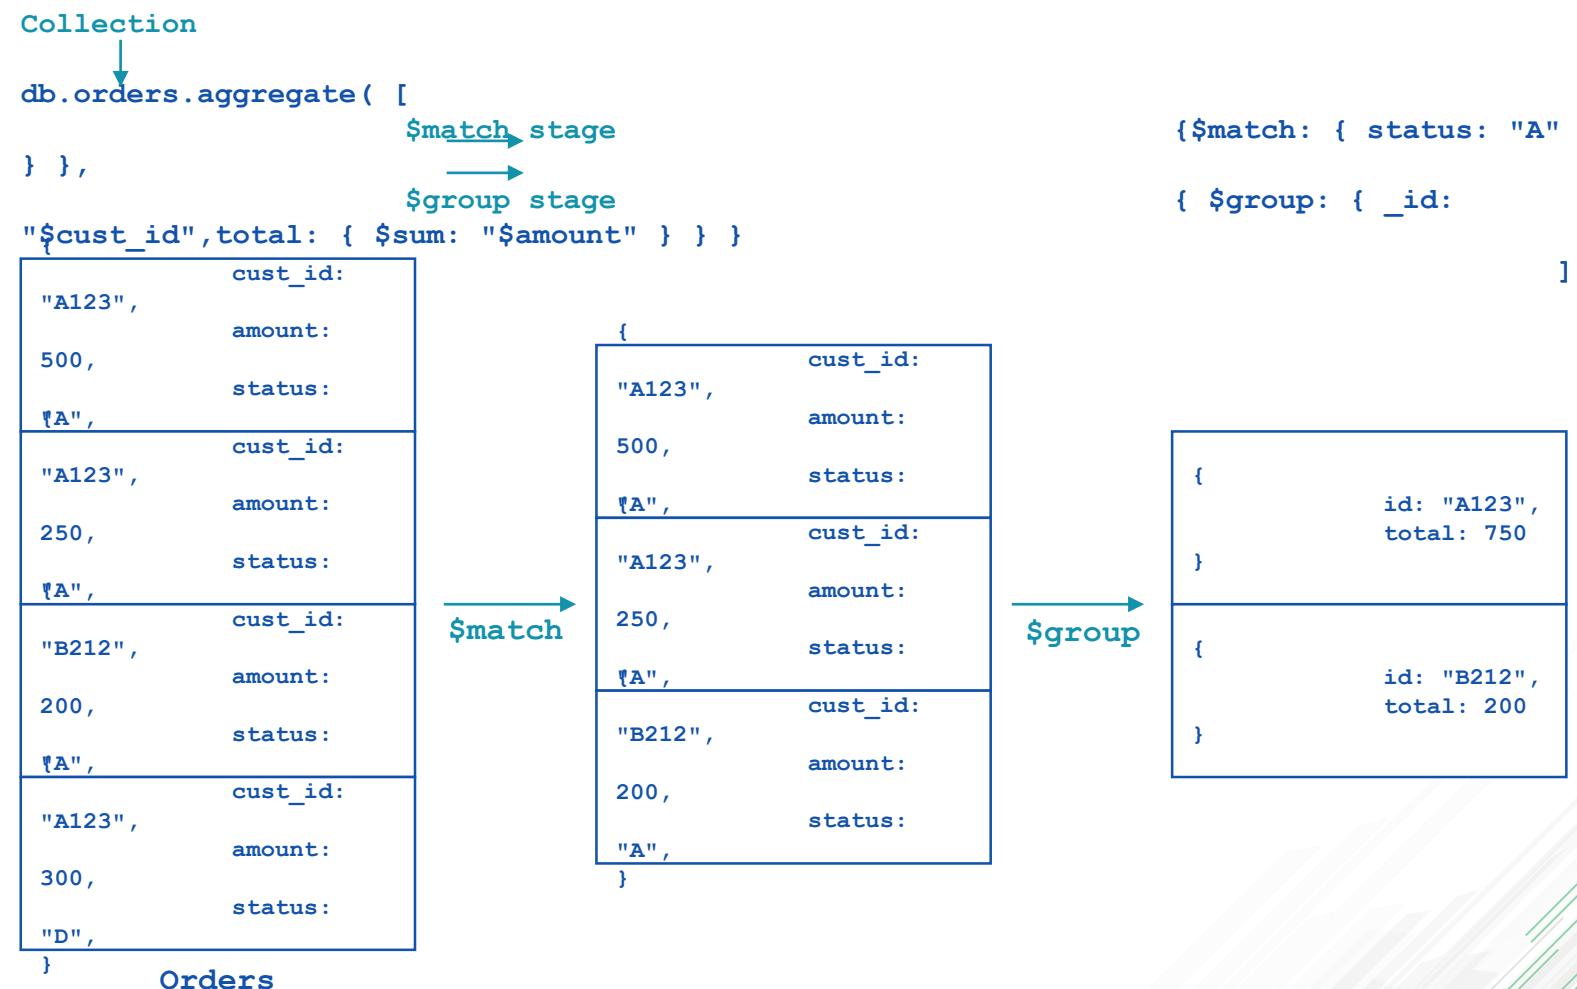
- TTL Indexes
 - Single Field indexes, when expired delete the document
- Unique Indexes
 - Ensures value is not duplicated
- Partial Indexes
 - Expression based indexes, allowing indexes on subsets of data
- Case Insensitive Indexes
 - supports text search using case insensitive search
- Sparse Indexes
 - Only index documents which have the given field



Aggregations

Advanced data processing pipeline for transformations and analytics

- Multiple stages
- Similar to a unix pipe
 - Build complex pipeline by chaining commands together
- Rich Expressions





Aggregation Features

A feature rich analytical framework

Pipeline Stages

- \$match
- \$group
- \$facet
- \$geoNear
- \$graphLookup
- \$lookup
- \$merge
- \$project
- \$sort
- \$unwind
- ...and more

Operators

- Mathematical
 - \$add, \$abs, \$subtract, \$multiply, \$divide, \$log, \$log10, \$stdDevPop, \$stdDevSam, \$avg, \$sqrt, \$pow, \$sum, \$zip, \$convert, \$round, etc
- Array
 - \$push, \$reduce, \$reverseArray, \$addToSet, \$arrayElemAt, \$slice, etc.

- Conditionals
 - \$and, \$or, \$eq, \$lt, \$lte, \$gt, \$gte, \$cmp, \$cond, \$switch, \$in, etc
- Date
 - \$dateFromParts, \$dateToParts, \$dateFromString, \$dateToString, \$dayOfMonth, \$isoWeek, \$minute, \$month, \$year, etc.
- String
 - \$toUpperCase, \$toLowerCase, \$substr, \$strcasecmp, \$concat, \$split, etc.
- Literals
 - \$exp, \$let, \$literal, \$map, \$type, etc
- Regex
 - \$regexxFnd, \$regexMatch, etc
- Trigonometry
 - \$sin, \$cos, \$degreesToRadians, etc



Compared to SQL JOINs and aggregation

```
SELECT
    city,
    SUM(annual_spend) Total_Spend,
    AVG(annual_spend) Average_Spend,
    MAX(annual_spend) Max_Spend,
    COUNT(annual_spend) customers
FROM (
    SELECT t1.city, customer.annual_spend
    FROM customer
    LEFT JOIN (
        SELECT address.address_id, city.city,
               address.customer_id, address.location
        FROM address LEFT JOIN city
        ON address.city_id = city.city_id
    ) AS t1
    ON
        (customer.customer_id = t1.customer_id AND
         t1.location = "home")
) AS t2
GROUP BY city;
```

SQL queries have a nested structure

Understanding the outer layers requires understanding the inner ones

So SQL has to be read “inside-out”



Versatile: Complex queries fast to create, optimize, & maintain

```
db.customers.aggregate ([  
    {  
        $unwind: "$address",  
    },  
    {  
        $match: {"address.location": "home"}  
    },  
    {  
        $group: {  
            _id: "$address.city",  
            totalSpend: {$sum: "$annualSpend"},  
            averageSpend: {$avg: "$annualSpend"},  
            maximumSpend: {$max: "$annualSpend"},  
            customers: {$sum: 1}  
        }  
    }  
])
```



These “phases” are distinct and easy to understand

They can be thought about in order... no nesting.

MongoDB’s aggregation framework has the flexibility you need to get value from your data, but without the complexity and fragility of SQL



Versatile: On-Demand Materialized Views



- Faster insights on your data: pre-compute and store results of common analytics queries
- With \$merge stage aggregation pipeline outputs with existing result sets to increment and enrich views
 - Updated each time the pipeline is run
 - Output to sharded and unsharded collections
 - Define indexes on each view
- With uniqueKey, control how documents are added to the view: Insert, Replace, Merge

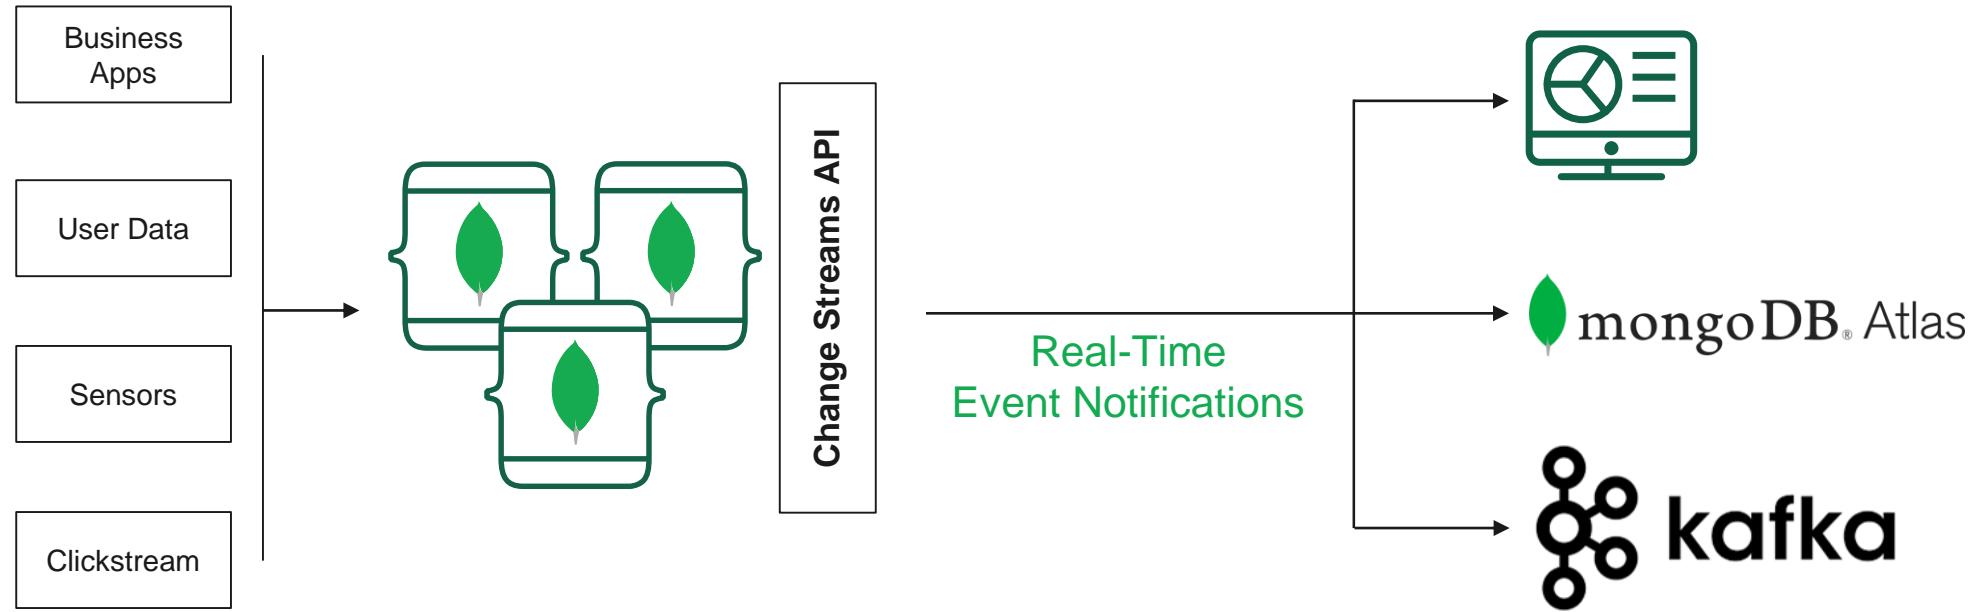


\$merge Syntax

```
{ $merge: {  
  to: "<output-collection>",  
  on: <field1> | [ <field1>,<field2>,... ],  
  whenNotMatched: <"insert" | "discard" | "fail">,  
  whenMatched: <"merge" | "replace" | "keepExisting" | "fail" | [...] >  
} }
```

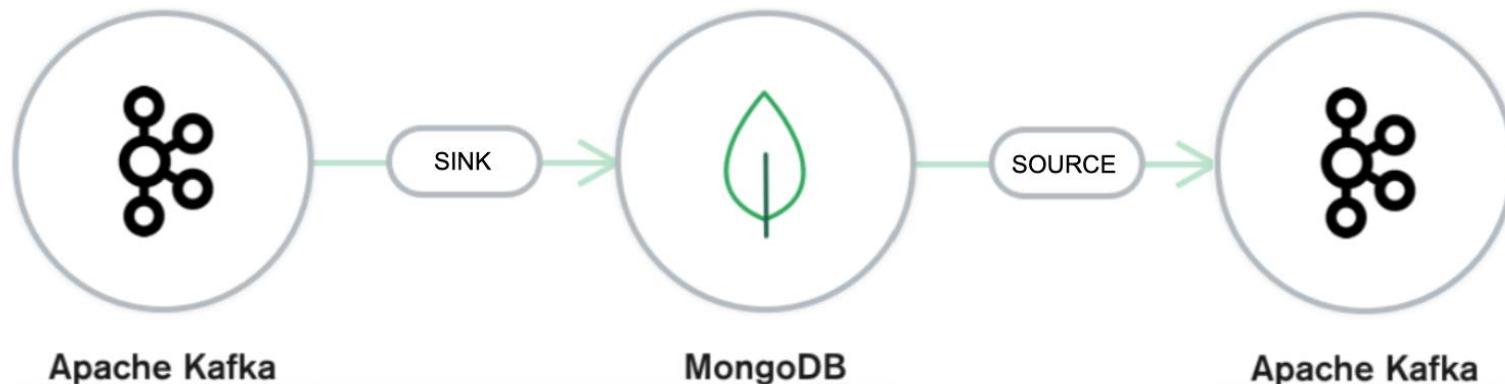


Versatile: MongoDB Change Streams



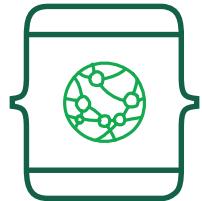
Enabling developers to build
reactive, real-time services

MongoDB Connector for Apache Kafka (Beta)



- Build robust data pipelines for microservices and Event Driven Architectures
- Developed with the community and supported by MongoDB engineers, verified by Confluent
- Supports MongoDB as a sink and a source for Kafka
- Integrate with Change Streams and Atlas triggers to create fully reactive, event driven pipelines

Intelligently put data where you need it



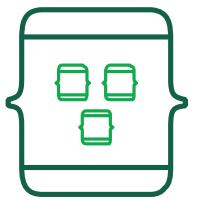
Highly Available

Built-in **multi-region** high availability, replication & automated failover



Workload Isolation

Ability to run both **operational & analytics workloads** on same cluster, for timely insight and lower cost



Scalability

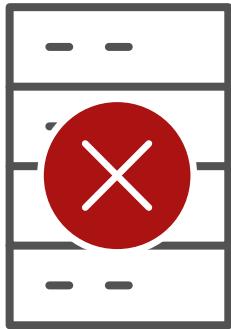
Elastic horizontal scalability – add/remove capacity dynamically **without downtime**



Locality

Declare **data locality rules** for governance (e.g. data sovereignty), class of service & local low latency access

Challenges of scale & HA with tabular databases



Single Node, Vertically Scaled

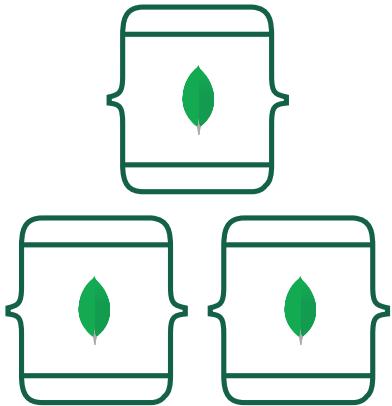
Scale-Out

- **Developer pain:** complex application-level sharding or external sharding frameworks
- **Static:** Limited elasticity
- **Trade-Offs:** Sacrifice key RDBMS functionality: cross-shard transactions, JOINs, referential integrity

Availability

- **Complex:** requires integration of external cluster managers for failure detection and recovery
- **Downtime:** extended outages with multi-minute recovery times
- **Developer pain:** write complex exception-handling code to handle multi-minute node failovers

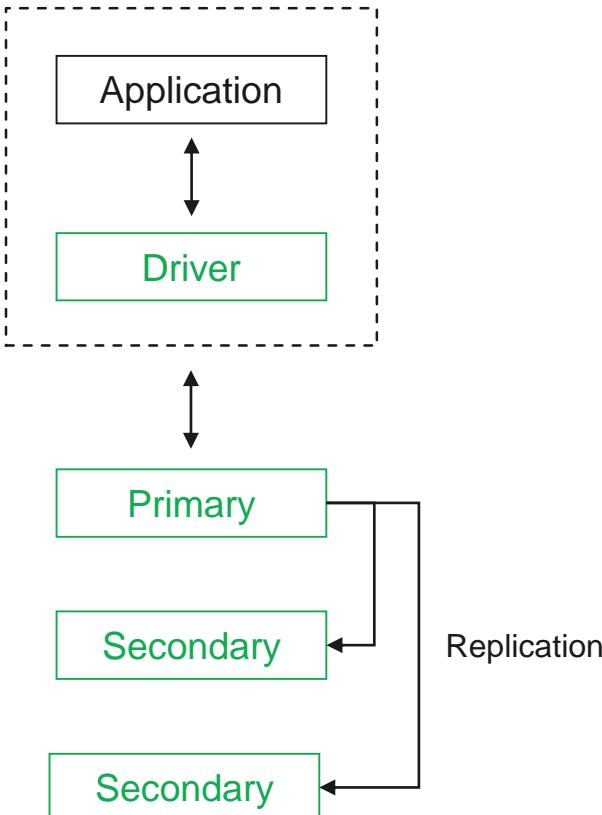
Put data where you need it: Availability



Replica Sets

- Up to 50 replicas, distributed across racks, data centers, and regions
- Self-healing using RAFT-based consensus protocol for automated failover & recovery
- Tunable durability and consistency controls
- Resilience with less code: retryable reads and writes

MongoDB Replica Sets



Replica Set – 2 to 50 copies

Self-healing

Data Center Aware

Addresses availability considerations:

- High Availability
- Disaster Recovery
- Maintenance

Workload Isolation: operational & analytics

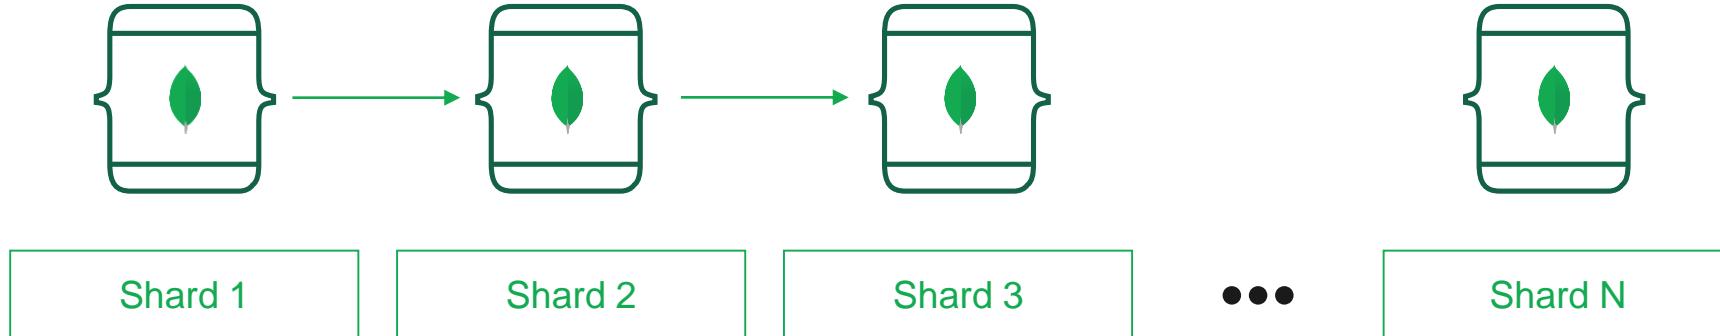
Serving global audiences with relational databases



MongoDB's native replica sets puts your entire database right next to users



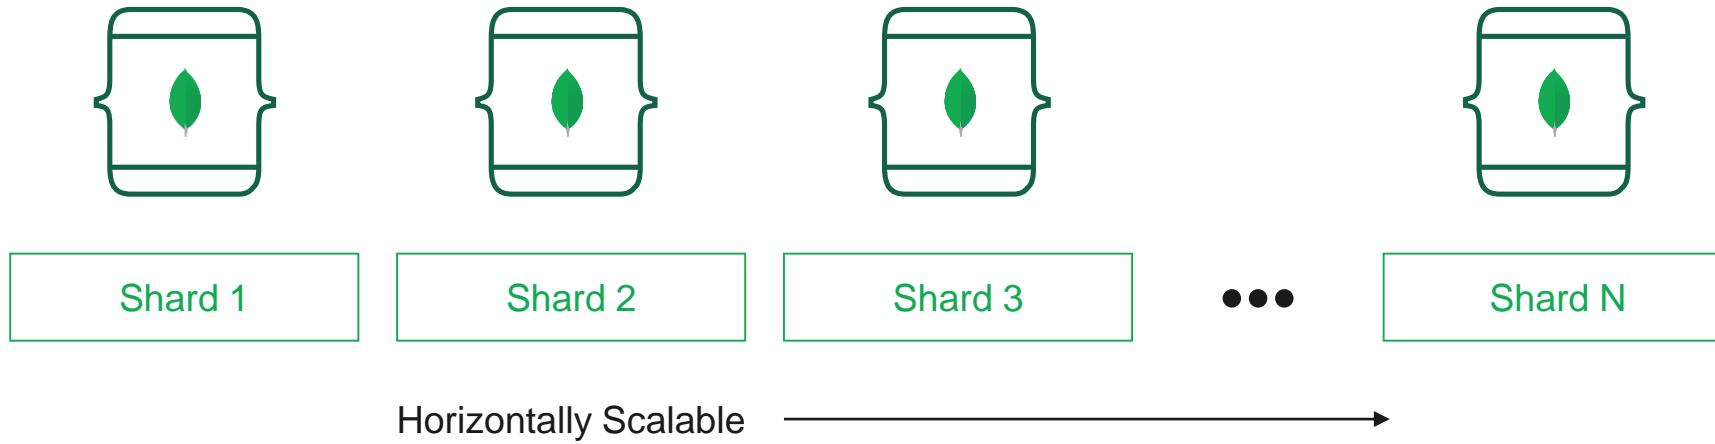
Put data where you need it: Scalability



Auto-Sharding

- Automatically scale beyond the constraints of a single node
- Application transparent
- Scale and rebalance incrementally, in real time
- Unlike NoSQL systems that randomly spray data across a cluster, MongoDB exposes multiple data distribution policies to optimize for query patterns and locality

Scaling MongoDB: Automatic Sharding

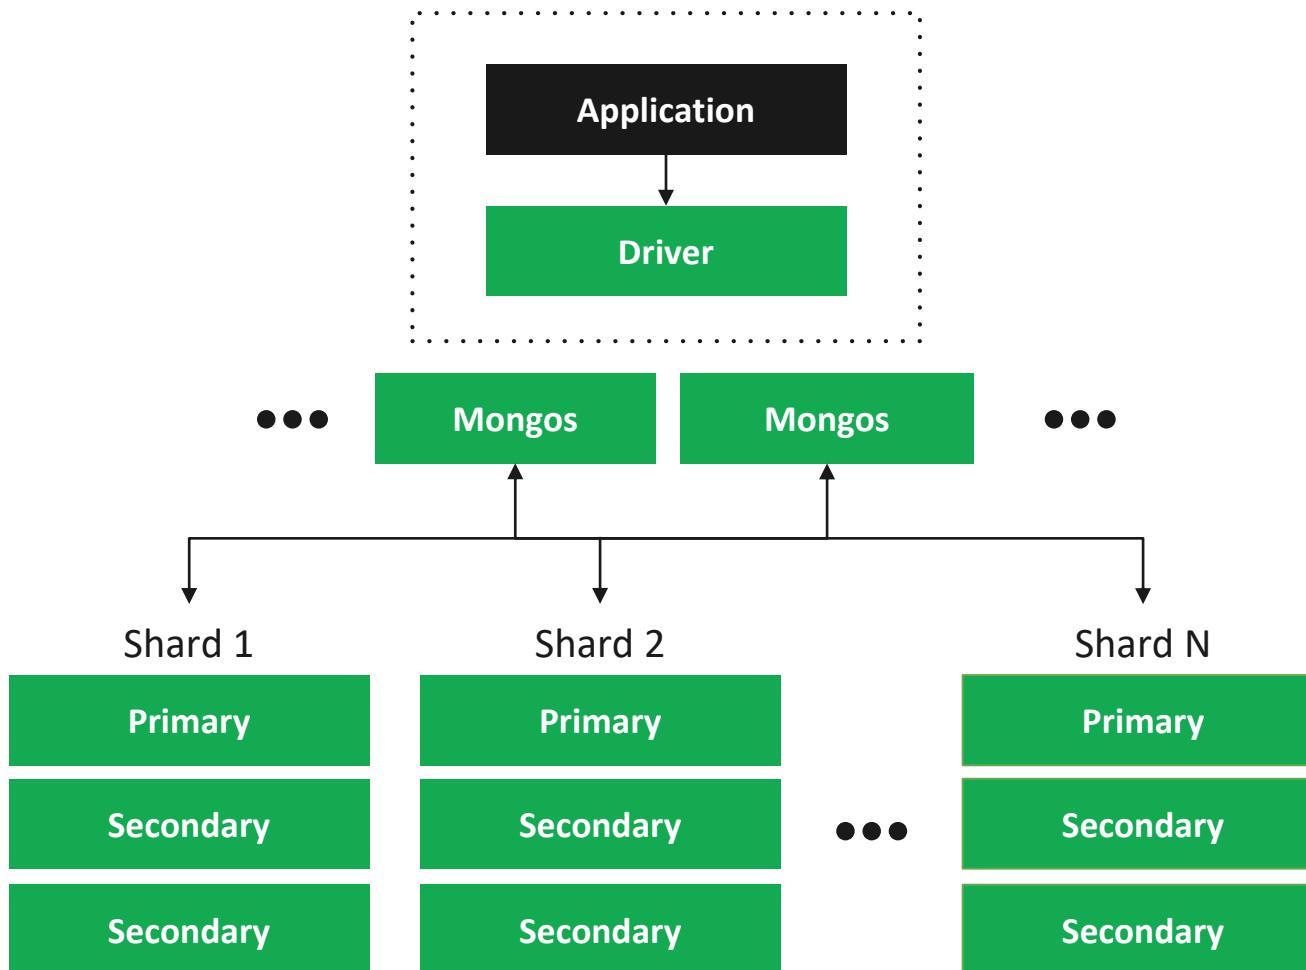


Multiple sharding policies: hashed, ranged, zoned

Increase or decrease capacity as you go

Automatic balancing for elasticity

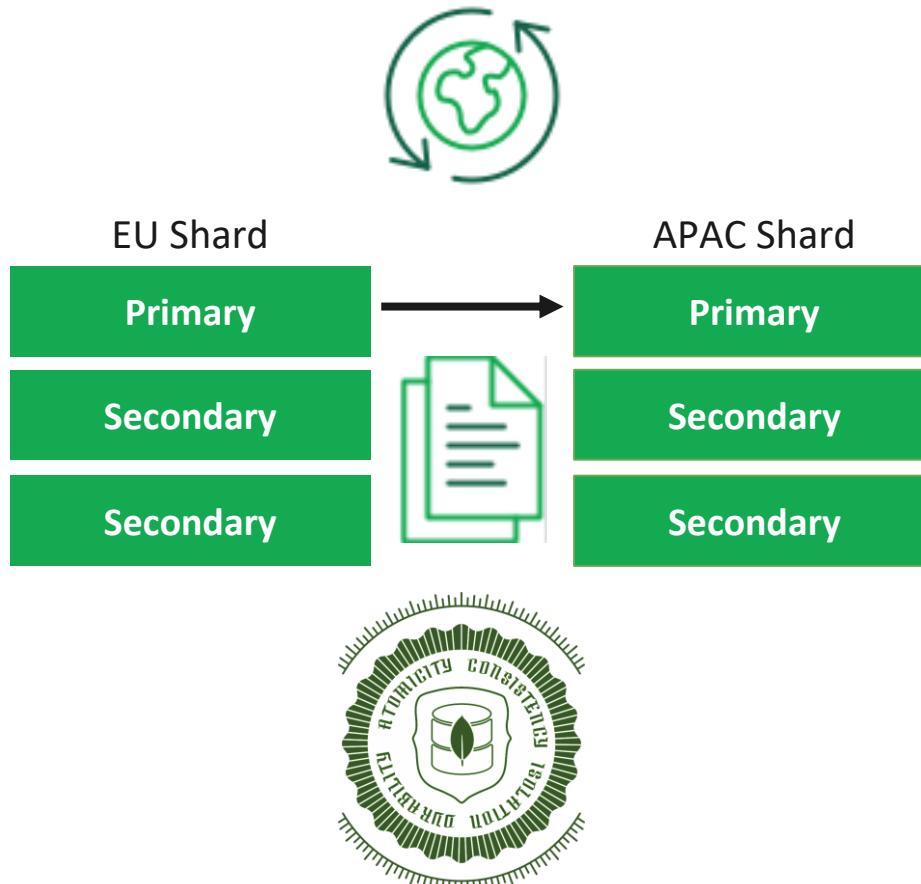
Sharding Architecture



High availability
- Replica sets

Horizontal scalability
- Sharding

Mutable Shard Key Values with Distributed Transactions

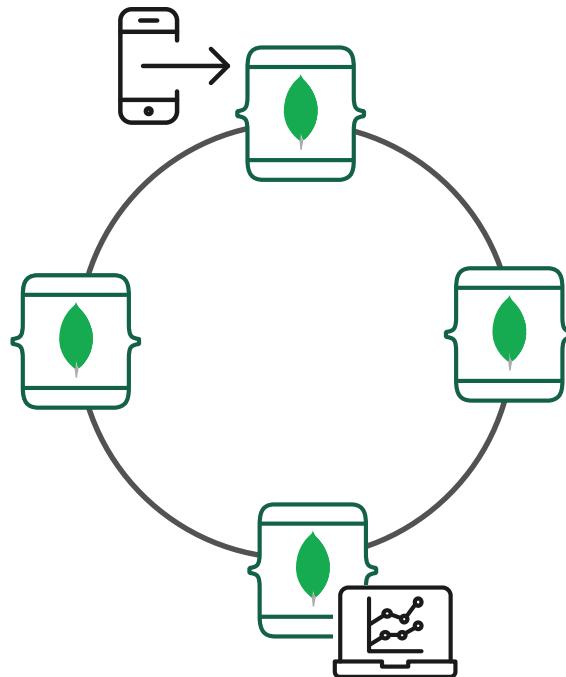


Default 60-second run time limit



- Improved sharding flexibility
- Modifying a shard key value will migrate a document using a **Distributed Transaction**
 - **Global-redistribution:** Rehoming documents to a new region
 - **Tiered storage:** Aging out older documents to low cost storage shard
- No need to delete then re-insert document with the new shard key value, and control atomicity in the app

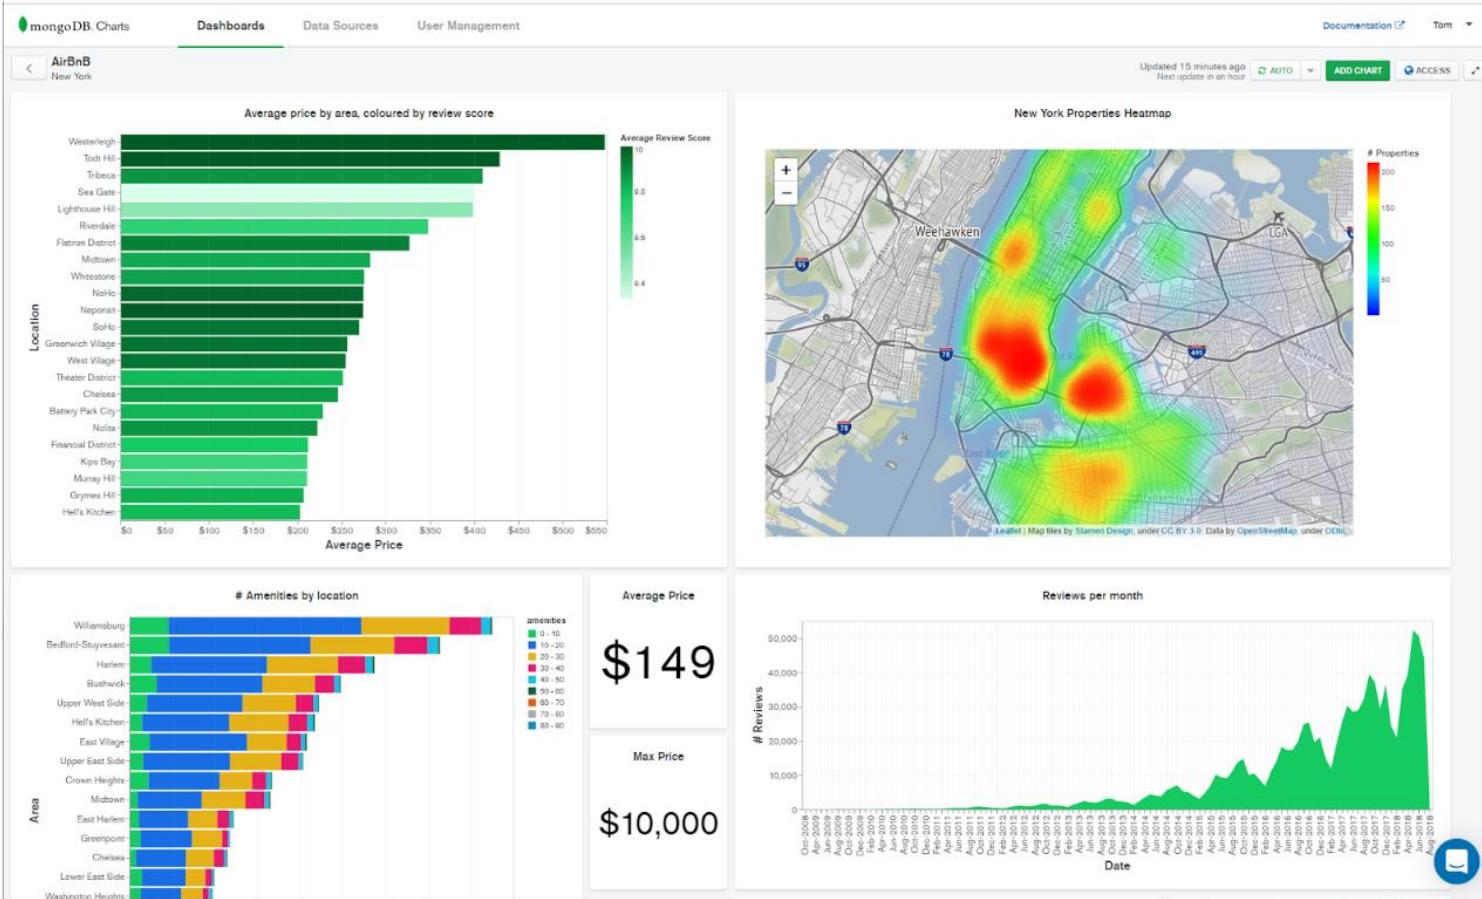
Put data where you need it: Workload Isolation



Enable different workloads on the same data

- Combine operational and analytical workloads on a single data platform
- Extract live insights from real-time data to enrich applications
- One set of nodes serving operational apps, replicating to dedicated nodes serving analytics: up to 50 nodes in a single replica set
- ETL-free

Sophisticated Analytics & Visualizations Of Data In Place



- Rich MongoDB query language & idiomatic drivers
- Charts
- Connector for BI
- Connector for Spark

MongoDB Charts: Create, Visualize, Share

```
{  
  "_id": {"$oid": "5afb2c3dc09c8d2dd5852cf2"},  
  "saleDate": {"$date": "2017-11-08T19:06:53.449Z"},  
  "items": [  
    {  
      "name": "envelope",  
      "tags": ["stationary", "office", "general"],  
      "price": {"$numberDecimal": "9.83"},  
      "quantity": 10  
    },  
    {  
      "name": "pens",  
      "tags": ["office", "writing", "school", "stationary"],  
      "price": {"$numberDecimal": "73.62"},  
      "quantity": 2  
    },  
    {  
      "name": "laptop",  
      "tags": ["office", "school", "electronics"],  
      "price": {"$numberDecimal": "595.72"},  
      "quantity": 4  
    },  
    {  
      "name": "notepad",  
      "tags": ["office", "writing", "school"],  
      "price": {"$numberDecimal": "34.65"},  
      "quantity": 3  
    }  
  ],  
  "storeLocation": "Seattle",  
  "customer": {  
    "gender": "M",  
    "age": 45,  
    "email": "uga@e.son",  
    "satisfaction": 4  
  },  
  "couponUsed": false,  
  "purchaseMethod": "Online"  
}
```



Save and Close

Data Source: test.supplySales

Filters: { storeLocation: { \$in: ["Denver", "New York"] }, 'items' }

FIELDS: filter, _id, couponUsed, customer, items, purchaseMethod, saleDate, storeLocation

Chart Type: Donut

Label: tags

Arc: # tags

Enter a title for your chart

Items.tags: writing, office, stationary, school, electronics

Share icon



Work with complex data

Connect to data sources securely.
Filter. Sample. Visualize. Share

Create standalone
dashboards or embed
directly into web apps

Co-locating operational and analytical workloads

Transactional Nodes



Primary

Secondary

Secondary

Secondary
{use = analytics}

Secondary
{use = analytics}

Analytics Nodes



Predictive Analytics & Data Science

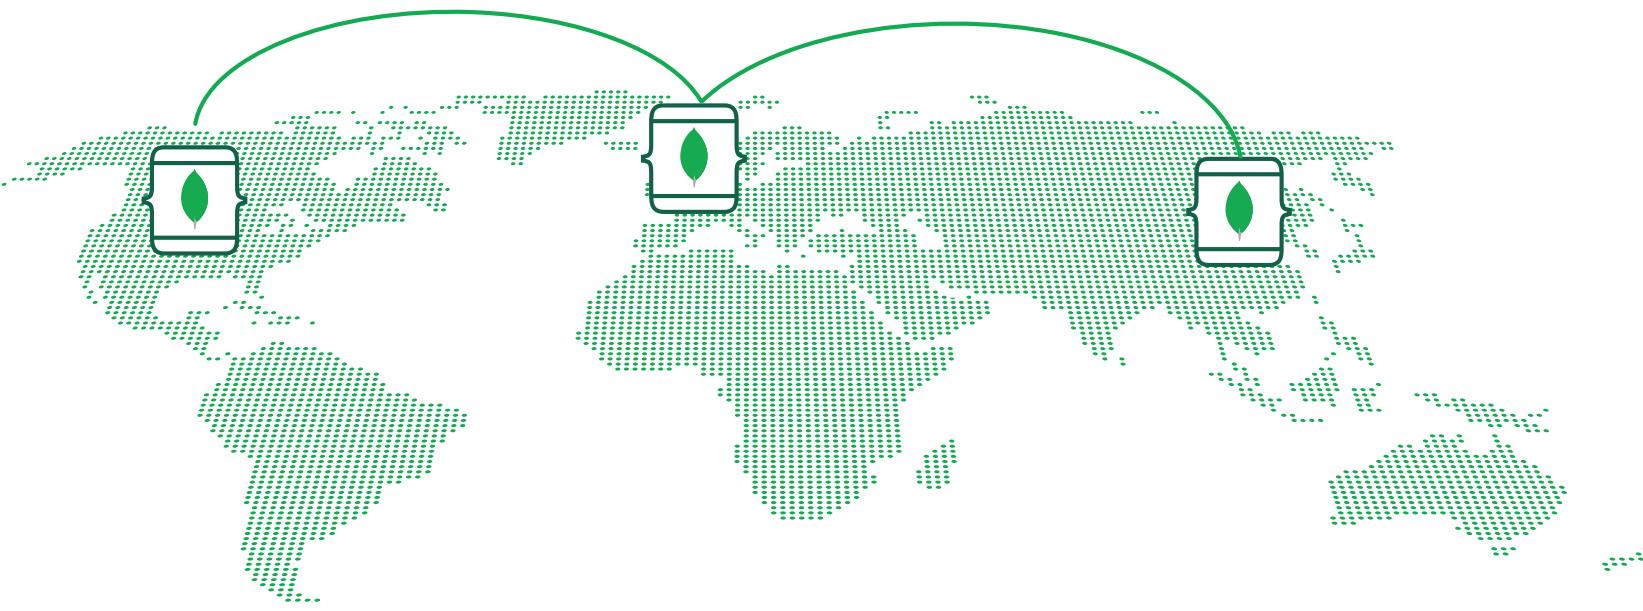


Technical Implementation (automated in Atlas)

```
// Set up replica set
rs.initiate(
{
  "_id": "rsl",
  "version": 1,
  "members": [
    {"priority": 1, "host": "localhost:27001", "_id": 0, "tags": {"use": "op"}},
    {"priority": 1, "host": "localhost:27002", "_id": 1, "tags": {"use": "op"}},
    {"priority": 1, "host": "localhost:27003", "_id": 2, "tags": {"use": "op"}},
    {"priority": 0, "host": "localhost:27004", "_id": 3, "tags": {"use": "analytics"}},
    {"priority": 0, "host": "localhost:27005", "_id": 4, "tags": {"use": "analytics"}}
  ]
})

// Force reads to analytics servers
read_client = MongoClient(host = 'localhost:27001', replicaset = 'rsl',
                          readPreference='secondary', readPreferenceTags = ["use:analytics"])
```

Put data where you need it: Locality



Governance

Class of Service

Latency

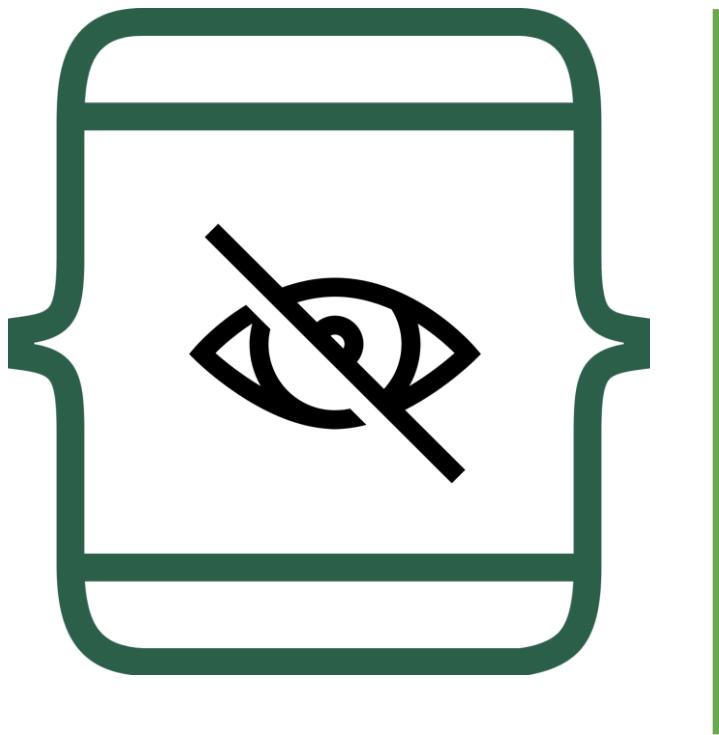
Intelligent Distribution via Zoned Sharding

- Policies to define data placement
- Name a server by region, tag your documents by region and MongoDB does the rest
- Documents automatically migrated as shard key ranges are modified
- Global queries across all data in all zones
- Fully managed as Global Clusters in Atlas

Secure: Wherever you put your data

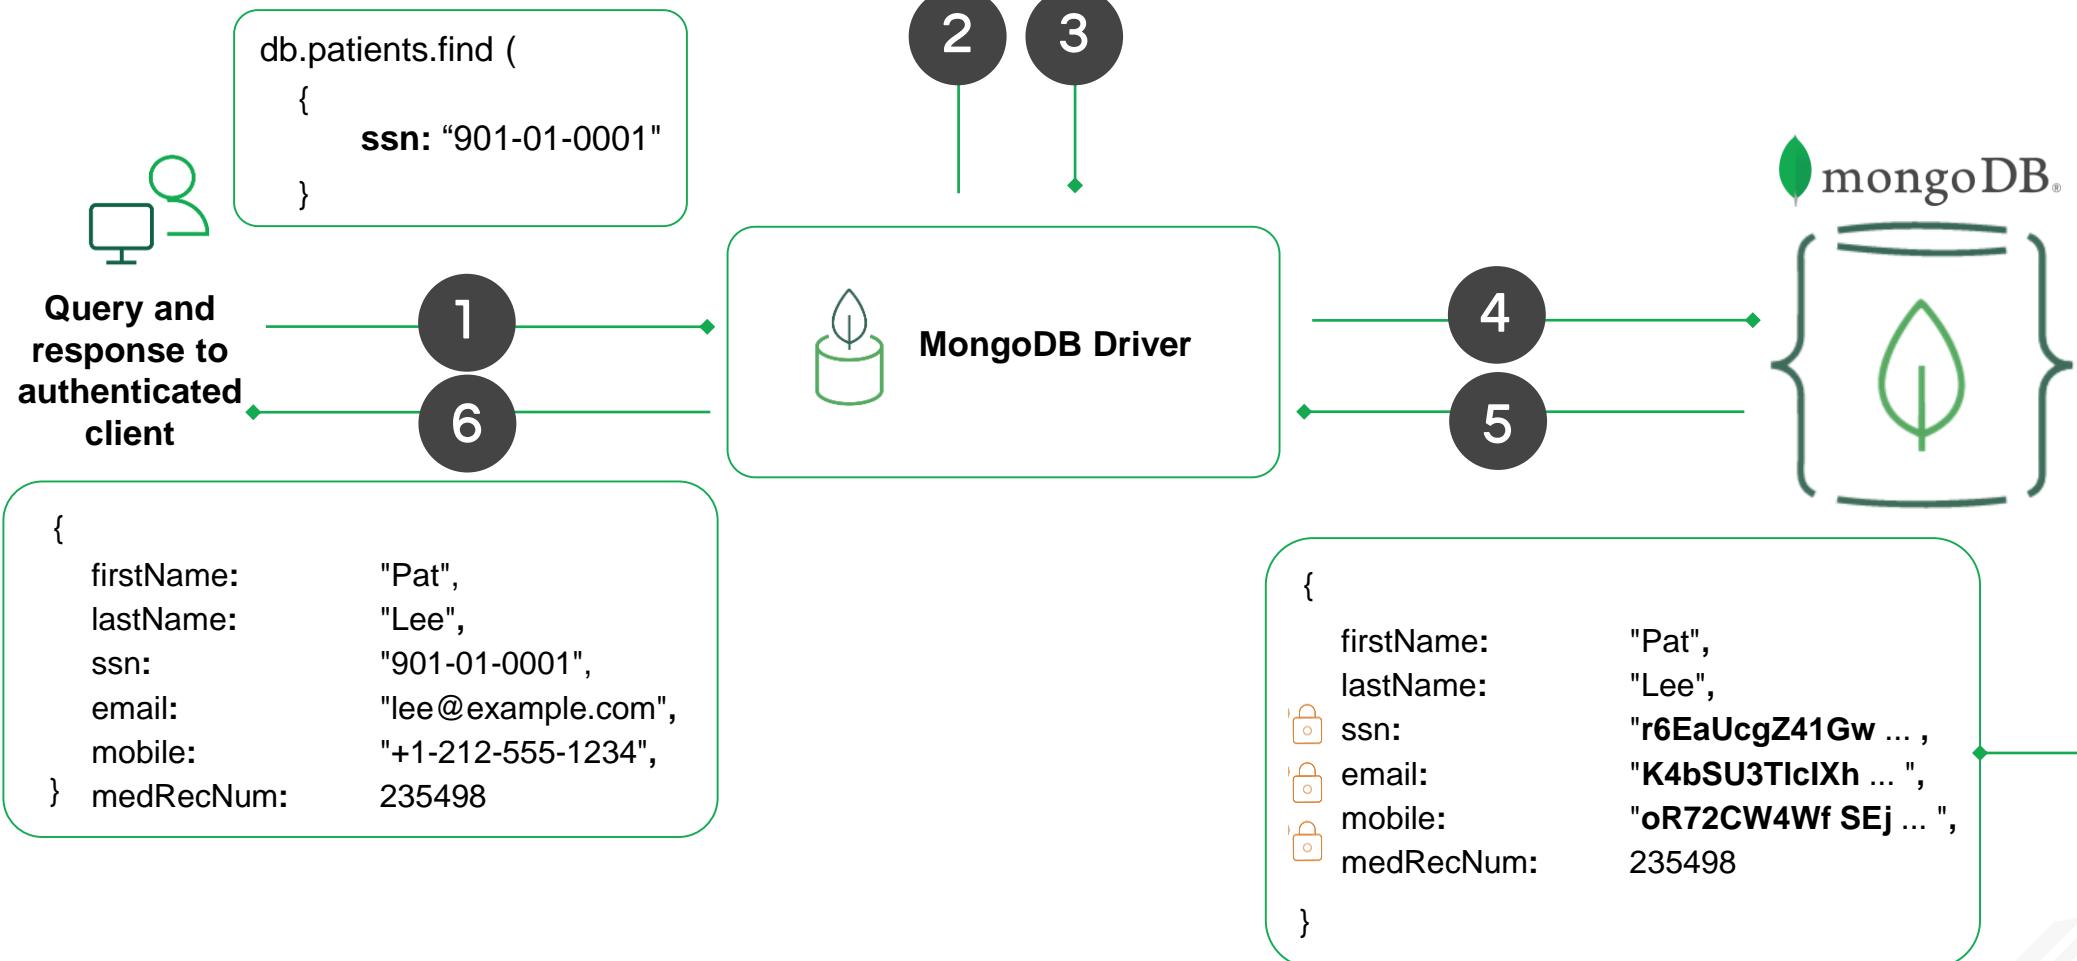
Business Needs	Security Features
Authentication	SHA-2, LDAP, Kerberos, x.509 Certificates
Authorization	RBAC, Read-Only Views, Field-Level Redaction
Auditing	Admin, DML, DDL, Role-based
Encryption	Network: TLS (with FIPS 140-2) In-Use: Client-Side Field Level Encryption (Beta) Disk: Encrypted Storage Engine or Partner Solutions

New: Client-Side Field Level Encryption (Beta)



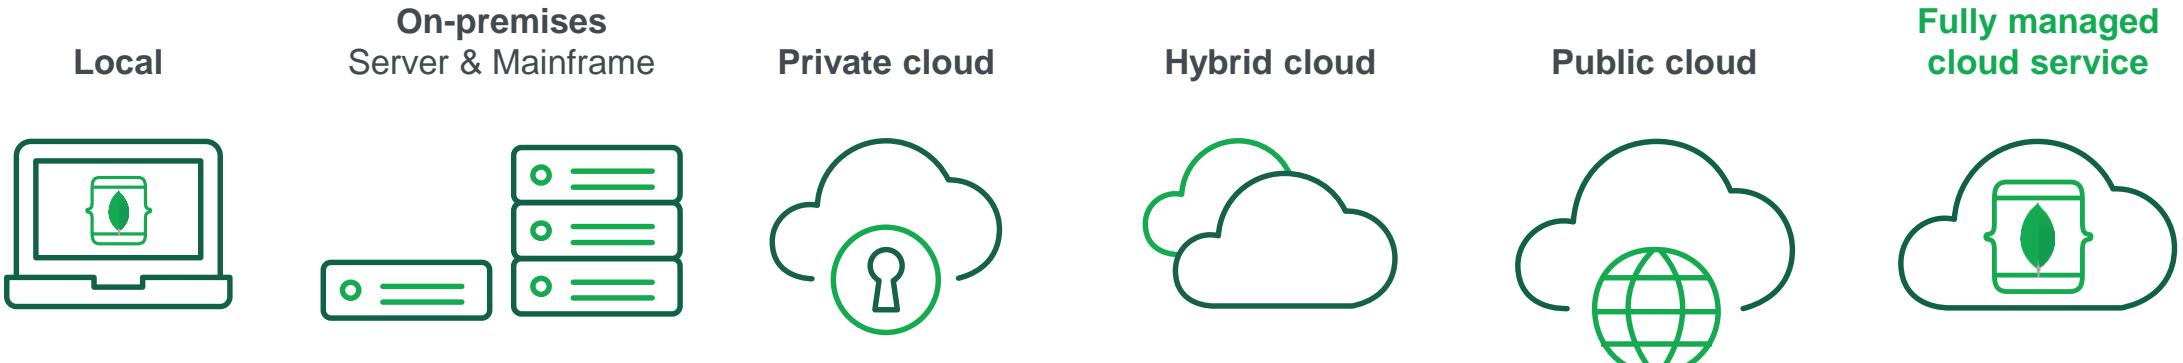
- Individual document fields encrypted by own key
- Database only sees ciphertext
- Many Advantages
 - Easy: Automatic and Transparent
 - Separation of Duties: (simplifies move to service-based systems as no service engineers ever see plaintext)
 - Compliant: Regulatory “right to be forgotten”
 - Fast: Minimal performance penalty

FLE Query Flow



Encrypted fields always stored, transmitted, and retrieved as ciphertext

Freedom to run anywhere



- Database that runs the same everywhere
- Leverage the benefits of a multi-cloud strategy
- Global coverage
- Avoid lock-in

Convenience: same codebase, same APIs, same tools, wherever you run

Manage & Run MongoDB Yourself Or Use DBaaS

Laptop



Mainframe



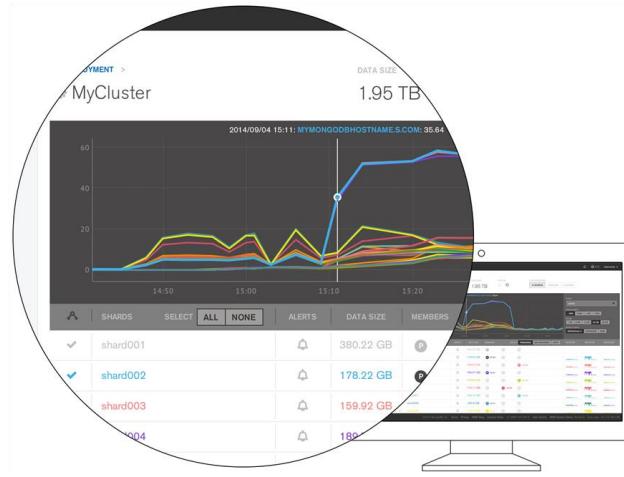
Server



Self-managed
in the cloud



Database as a Service

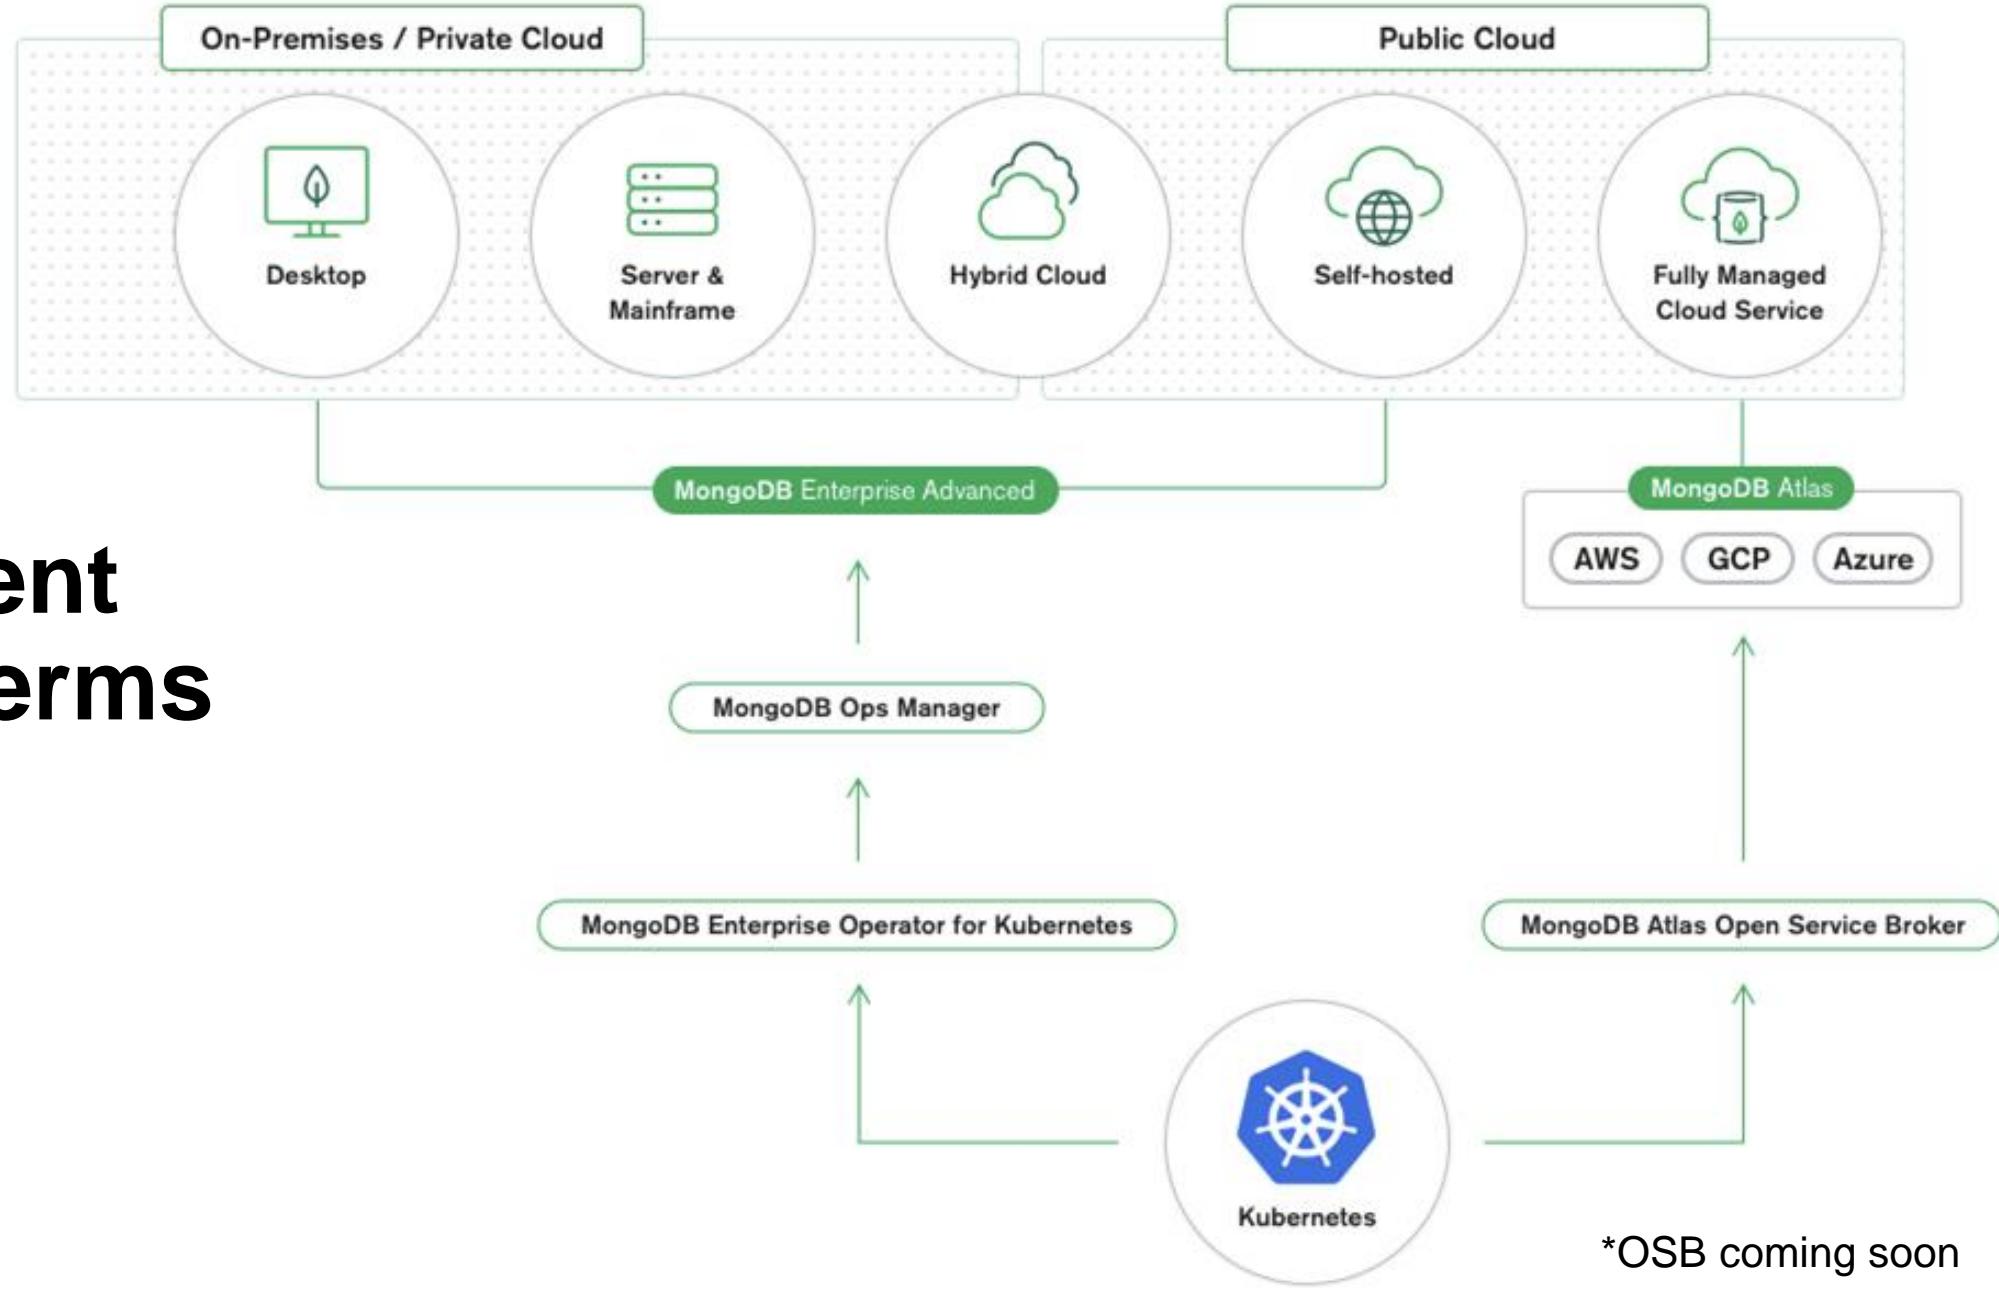


Self-Managed: Ops Mgr

- Automation
- Monitoring Alerting
- Backup/Restore
- Patching
- Scaling
- Performance Advice
- Kubernetes Operator

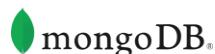
Database as a Service

- Self-service & fully automated
- Global, elastic clusters
- Highly available and secure out of the box
- Operational tooling built in: monitoring, backups
- **On AWS, Azure, GCP, 60+ regions total**



Deployment on your terms

Don't I get this freedom with relational databases?

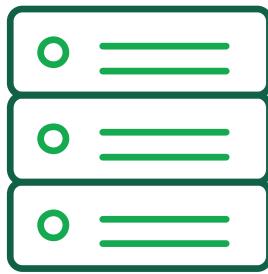


Fragmentation across different cloud platforms

- Different versions
- Different database features and options
- Different scaling characteristics
- Different cross-region options
- Different monitoring, alerting, backup, and recovery
- Different security controls

Architectural mis-alignment

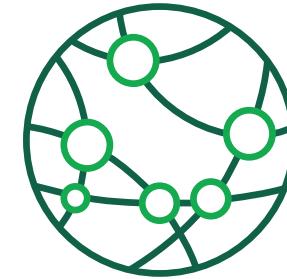
Choosing the right data platform for the cloud



**Exploit scale-out,
commodity platforms**

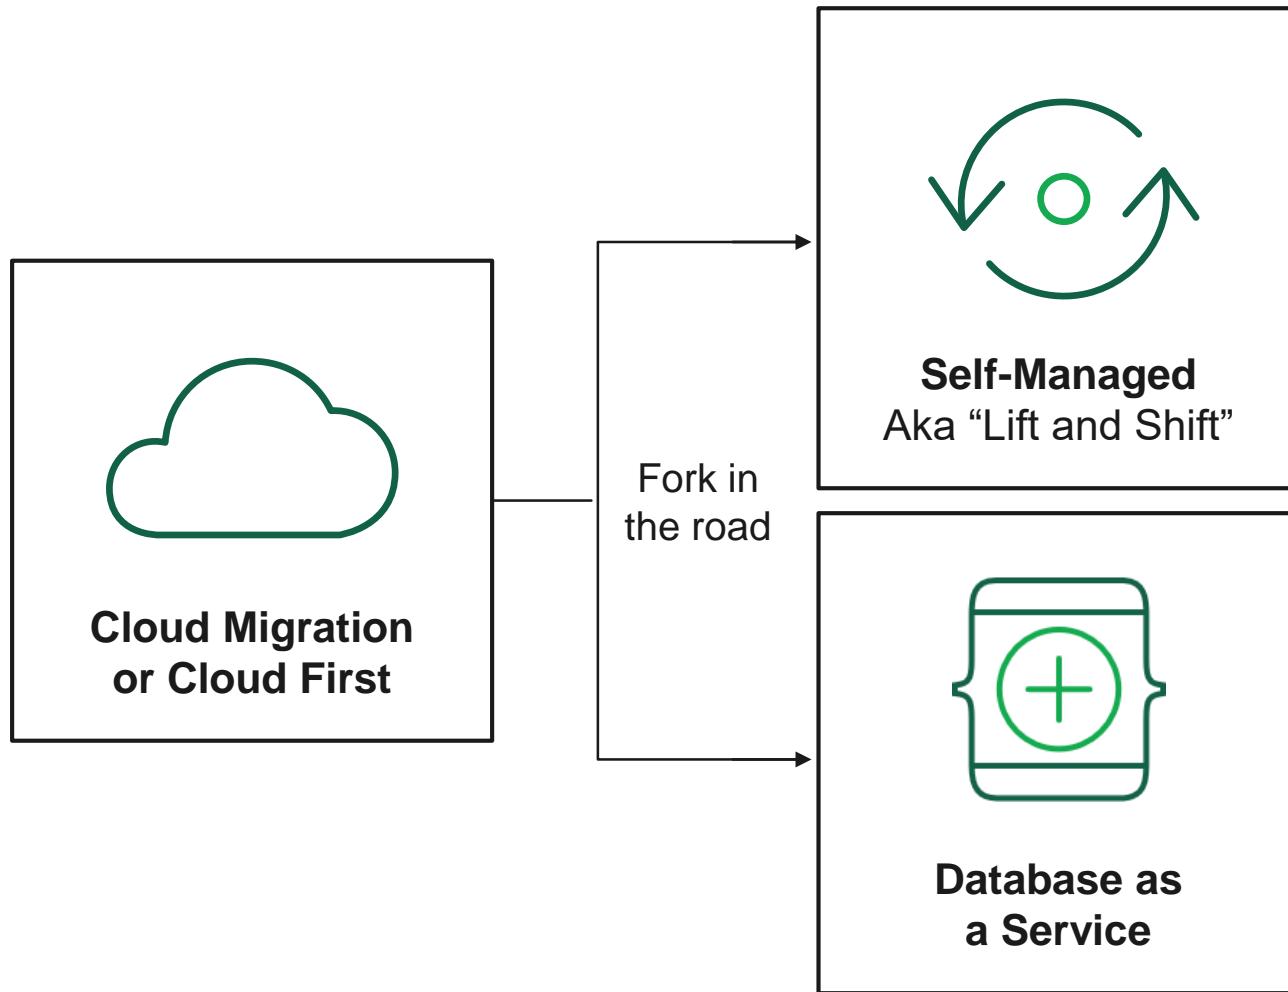


**Support velocity of
modern app development**

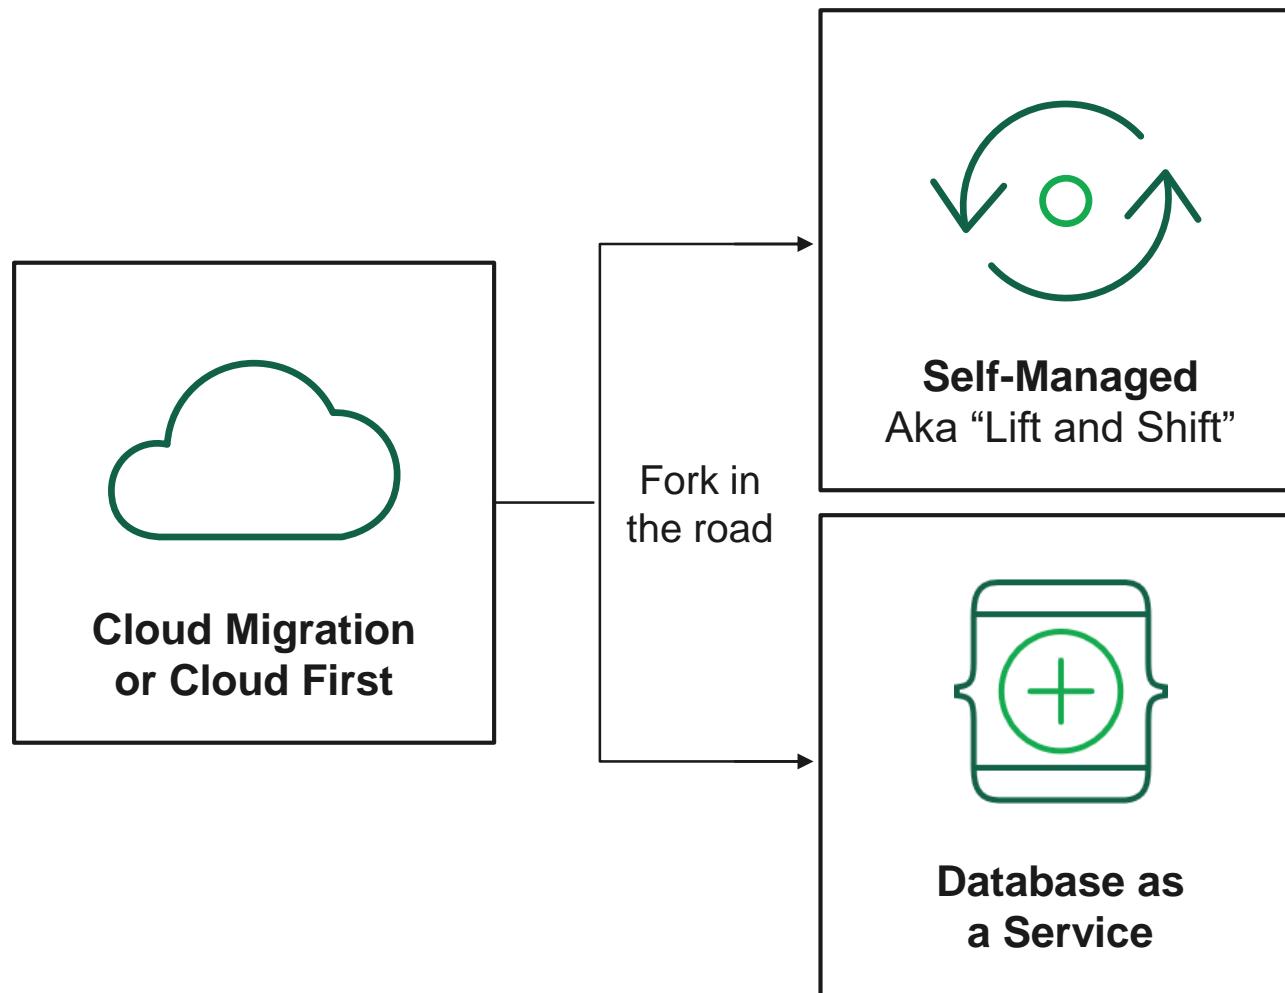


**Deploy globally,
on-demand**

2 choices as you move to the cloud: Self-Managed or DBaaS

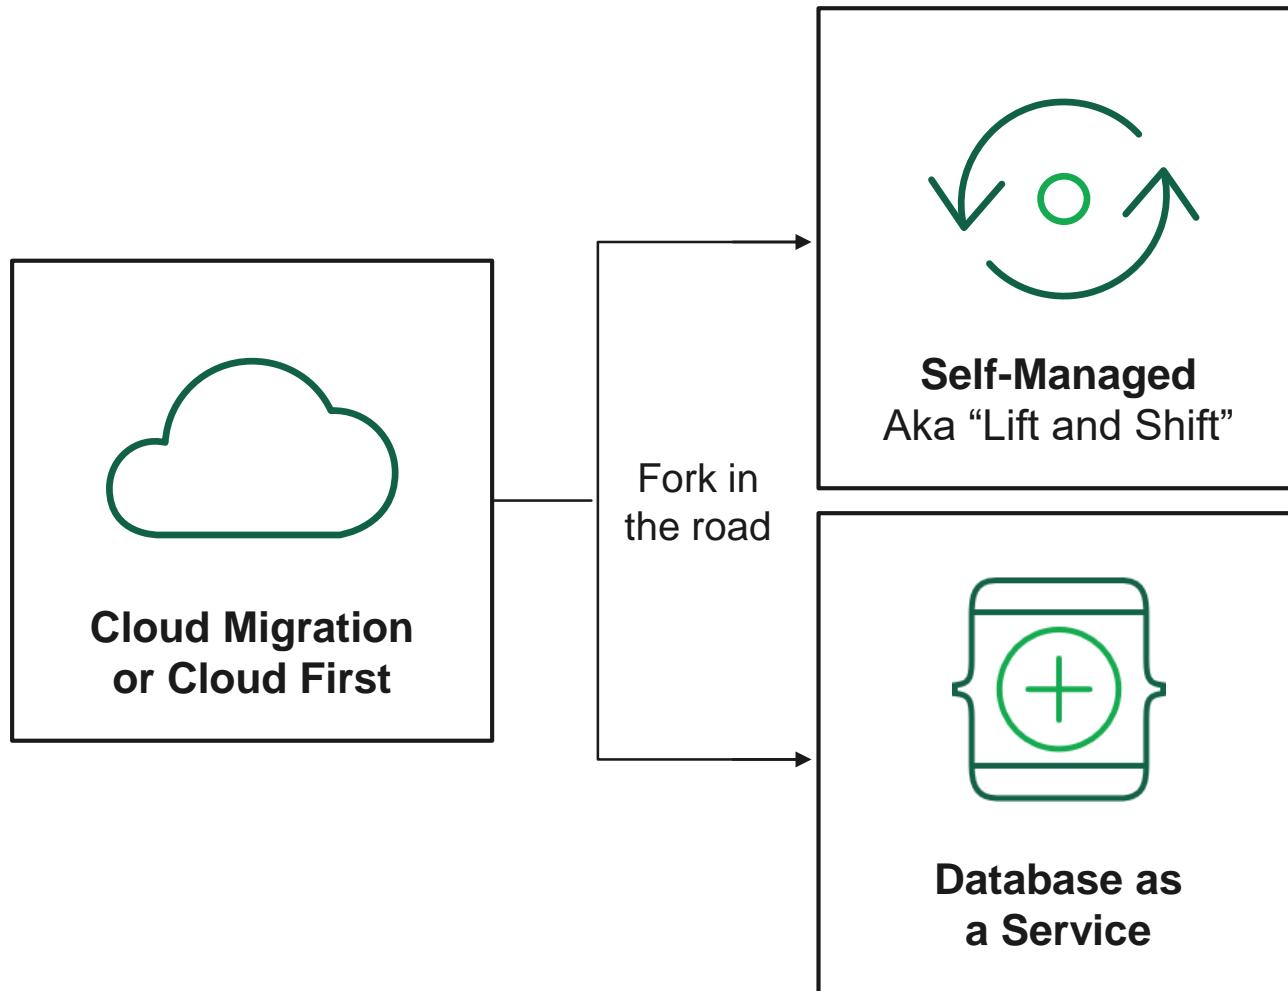


2 choices as you move to the cloud: Self-Managed or DBaaS



1. Provision instances and storage
2. Configure HA
3. Configure security
4. Configure backup/restore
5. Monitoring & alerting
6. Ongoing upgrades & maintenance

2 choices as you move to the cloud: Self-Managed or DBaaS



1. Provision instances and storage
2. Configure HA
3. Configure security
4. Configure backup/restore
5. Monitoring & alerting
6. Ongoing upgrades & maintenance

**Choose instance, hit deploy,
available in a couple of
minutes**

Atlas

unlocks **agility**
and **reduces**
cost



Self-service and
elastic



Global and highly
available



Secure by default



Comprehensive
monitoring



Managed backup



Sync & Serverless

MongoDB Atlas — Global Cloud Database

<p>Self-service & elastic</p> <p>Deploy, modify, and upgrade on demand with best-in-class operational automation</p> <p>Automated database maintenance</p> <p>Database and infrastructure resources as code</p> <p>Scale up, out, or down in a few clicks or API calls</p>	<p>Global & cloud-agnostic</p> <p>Available in 60+ regions across AWS, Azure, GCP</p> <p>Global clusters for read/write anywhere deployments and multi-region fault tolerance</p> <p>Easy migrations with a consistent experience across cloud providers</p>	<p>Enterprise-grade security & SLAs</p> <p>Network isolation, VPC peering, end-to-end encryption, and role-based access controls</p> <p>Encryption key management, LDAP integration, granular database auditing</p> <p>SOC 2 / Privacy Shield / HIPAA</p> <p>Guaranteed reliability with SLAs</p>
<p>Comprehensive monitoring</p> <p>Deep visibility into 100+ KPIs with proactive alerting</p> <p>Real-time performance tracking and Performance Advisor</p> <p>APIs to integrate with monitoring dashboards</p>	<p>Managed backup</p> <p>Point-in-time data recovery</p> <p>Queryable backup snapshots</p> <p>Consistent snapshots of sharded deployments</p> <p>Cloud data mobility</p>	<p>Sync and Serverless</p> <p>Simple, serverless functions for backend logic, service integrations, and APIs</p> <p>Database access from your frontend secured by straightforward, field-level access rules</p> <p>Database and authentication triggers to react to changes in real time</p>

Powering Mission-Critical Workloads

Queries per day

1700x in 2.5 years

0B

7B

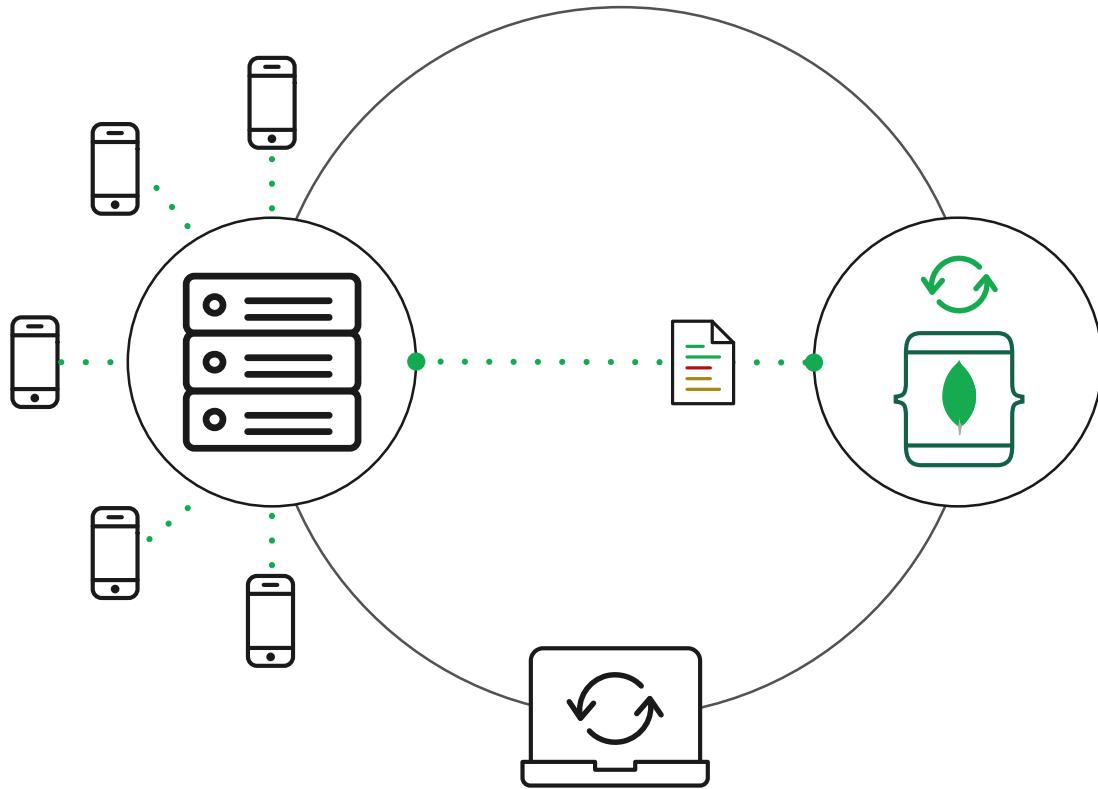
112B

January 2017

January 2018

June 2019

Freedom to run Anywhere: Cloud Migration



Migrate existing deployments running anywhere into MongoDB Atlas with minimal impact to your application.

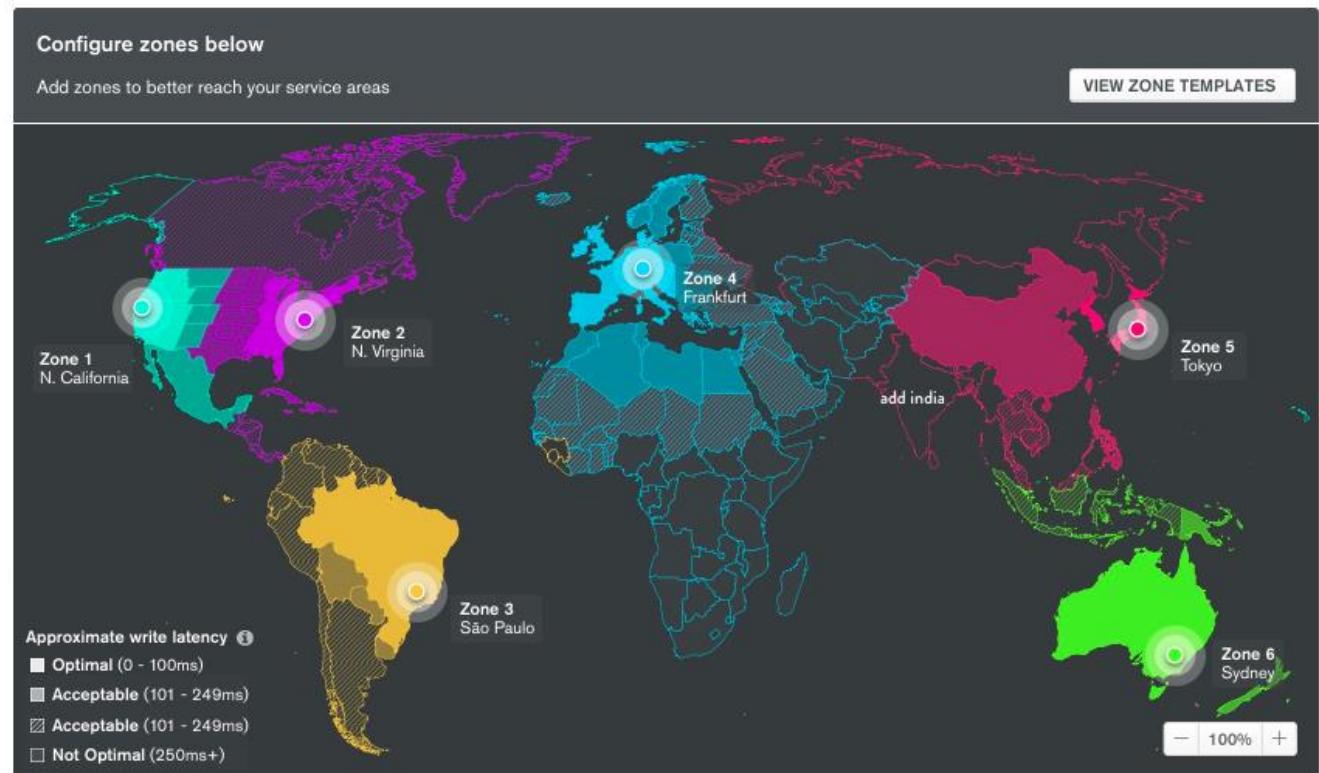
Live migration works by:

- Performing a sync between your source database and a target database hosted in MongoDB Atlas
- Syncing live data between your source database and the target database by tailing the oplog
- Notifying you when it's time to cut over to the MongoDB Atlas cluster

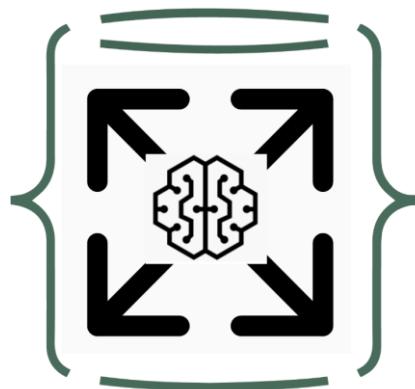
Geo-Distributed apps: MongoDB Atlas Global Clusters

Distribute your fully automated database across multiple geographically distributed zones made up of one or more cloud regions

- Read and write locally to provide single-digit millisecond latency for your distributed applications
- Ensure that data from certain geographies lives in predefined zones
- Easily deploy using prebuilt zone templates or build your own zones by choosing cloud regions in an easy-to-use, visual interface

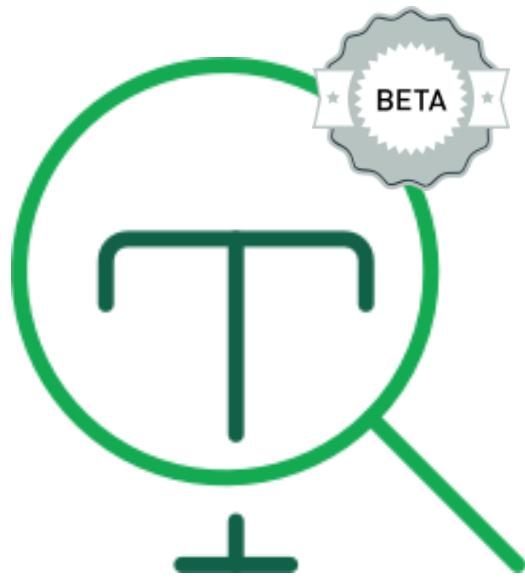


New: MongoDB Atlas Auto Scaling (late-Summer)



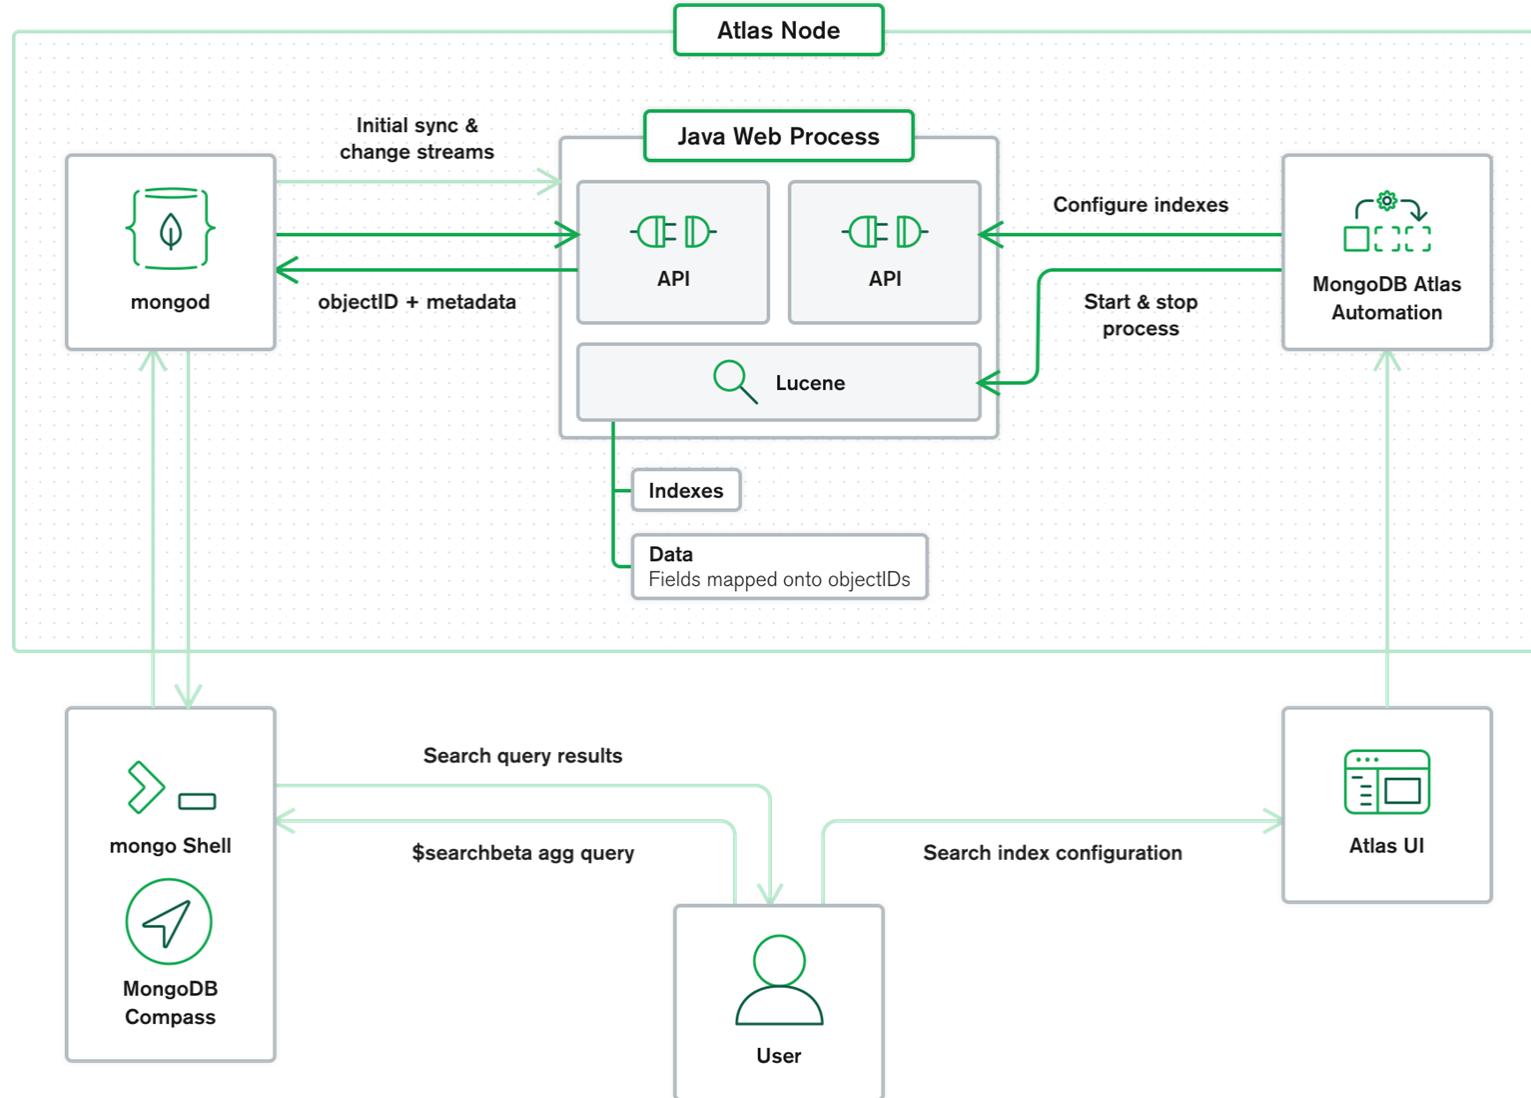
- Elastic scaling of instance sizes so your provisioned capacity responds to demand
 - Monitors key resource utilization metrics
 - Toggle on and off via UI or API
 - Cap peak instance sizes to control costs
- Rolling restarts across replica sets to maintain app availability
- Auto-storage scaling available since 2018

New: MongoDB Atlas Full Text Search (Beta)



- Adds Full Text Search as a fully managed service to your Atlas cluster
 - Power of Lucene 8, without provisioning and running a separate search platform
 - Integrated with MongoDB Query Language, so no separate APIs to learn
 - Dynamic and static indexing supporting fuzzy & wildcard search, Boolean & compound queries, language analyzers, scoring and snippets
 - Configured via Atlas Data Explorer or API

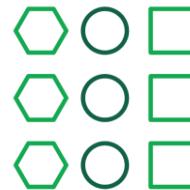
MongoDB Atlas Full Text Search Architecture



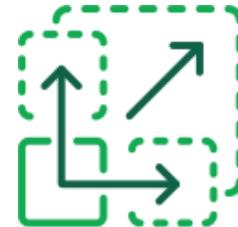
New: MongoDB Atlas Data Lake (Beta)

Analyze data in any format on S3 using the MongoDB Query Language

**Multiple Formats,
No Schema**



**Auto-Scale,
At Any Scale**



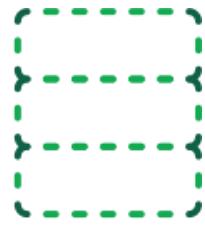
**Best Tools,
High Productivity**



**Integrated with Atlas, Single UI,
Billing, Permissioning**



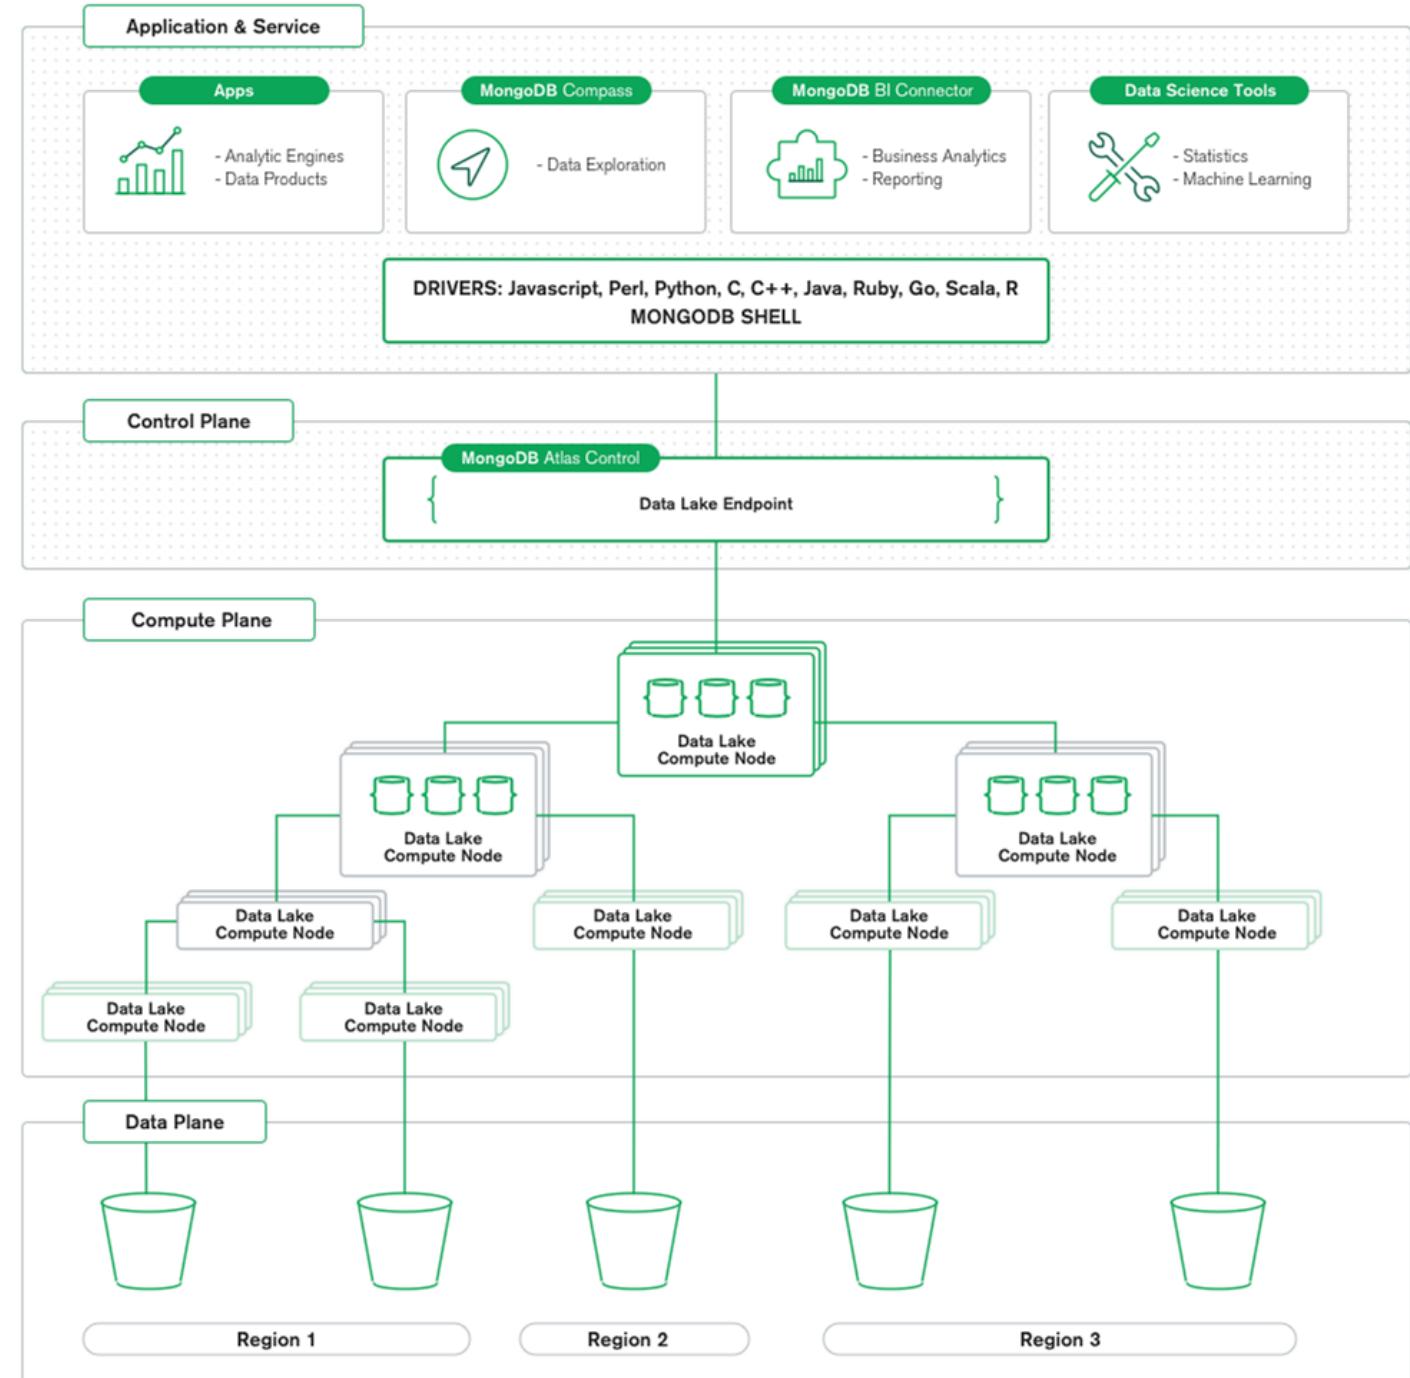
**Serverless, No
Infrastructure Management**



**On-Demand, Usage-
Based Pricing**



MongoDB Atlas Data Lake Architecture



MongoDB Atlas Data Lake UI

mongoDB.Atlas OK All Clusters

CONTEXT Query Engine MONGODB ENTERPRISE TOOLS > QUERY ENGINE

Please set your time zone Usage This Month:\$0.00 [details](#) Admin Jane

Data Lake Configure a New Data Lake

Atlas-DATA-Lake-Testing 4 DBS 51 COLLECTIONS

CONNECT CONFIGURATION ...

Queries Executed: 2463 (2463 SUCCESSFUL, 59 FAILED)

Total Data Scanned: 93.06 MEGABYTES

Total Data Returned: 42.28 MEGABYTES

Average Execution Time: 14.13 SECONDS

eliot 1 DBS 1 COLLECTIONS

Queries Executed: 144 (144 SUCCESSFUL, 6 FAILED)

Total Data Scanned: 264.17 MEGABYTES

Total Data Returned: 15.96 KILOBYTES

System Status: All Good Last Login: 76.218.106.48 Version: b1bfe792f5@master

Atlas Plan: NDS Effective Plan: NDS Plan Start Date: 2019-05-03T19:19:49Z Central URL: https://cloud-dev.mongodb.com Organization Name: Mong

©2019 MongoDB, Inc. [Status](#) [Terms](#) [Privacy](#) [Atlas Blog](#) [Contact Sales](#)

DATA LAKE BETA

Clusters

Data Lake BETA

Database Access

Network Access

Advanced

Access Management

Activity Feed

Alerts 1

Settings

Charts

Stitch

Triggers

Docs

Support

System Status: All Good Last Login: 76.218.106.48 Version: b1bfe792f5@master

Atlas Plan: NDS Effective Plan: NDS Plan Start Date: 2019-05-03T19:19:49Z Central URL: https://cloud-dev.mongodb.com Organization Name: Mong

©2019 MongoDB, Inc. [Status](#) [Terms](#) [Privacy](#) [Atlas Blog](#) [Contact Sales](#)

Configure a New Data Lake

Overview Name List Buckets Role & Policy Validate

Analyze your data in S3 with MQL

We'll guide you through the set up process

1. You'll create an AWS IAM role for Atlas to use.

2. Then you'll give the Atlas user read-only access to the S3 buckets you want to query.

3. You'll connect to your data lake with the Mongo Shell and map your S3 directories to databases and collections

4. Then you can build and run your queries in Mongo Shell, MongoDB Compass, or any MongoDB Driver.

View Data Lake docs

Cancel Configure a New Data Lake

Use Cases: Atlas Data Lake

Data Lake Analytics

- explore all of your rich data naturally
- get to data as it lands via streams or microservices
- democratize access across diverse user groups



Data Products and Services

- monetize data
- market research, data- and insight-as-a-service
- snapshots, time series analysis, predictive analytics to innovate faster



Active Archives

- historical analysis against data assets retained in long term cold storage
- cost-effective data strategy



mongoDB® Atlas : The only *true* multi-cloud database as a service



Google Cloud Platform



Database that runs the same everywhere

Consistent experience
across AWS, Azure, and
GCP

Coverage in any geography

Deploy in 59 regions
worldwide

Create globally distributed
databases with a few clicks

Leverage the benefits of a multi-cloud strategy

Exploit the benefits of AWS,
Azure, or GCP services on
your data

Avoid lock-in

Easily migrate data
between cloud providers

Beyond the Server

The Broader IODP Accelerates Everything

3-5x increase in productivity by leveraging the Intelligent Operational Data Platform

MongoDB Mobile The power of MongoDB in your device

MongoDB Stitch

Serverless platform that allows developers to focus on innovation rather than plumbing, services orchestration, and boilerplate code

MongoDB Atlas

Rapidly deploy, dynamically scale, and distribute databases across regions and cloud providers

Client Application or Service

Application Logic

Data

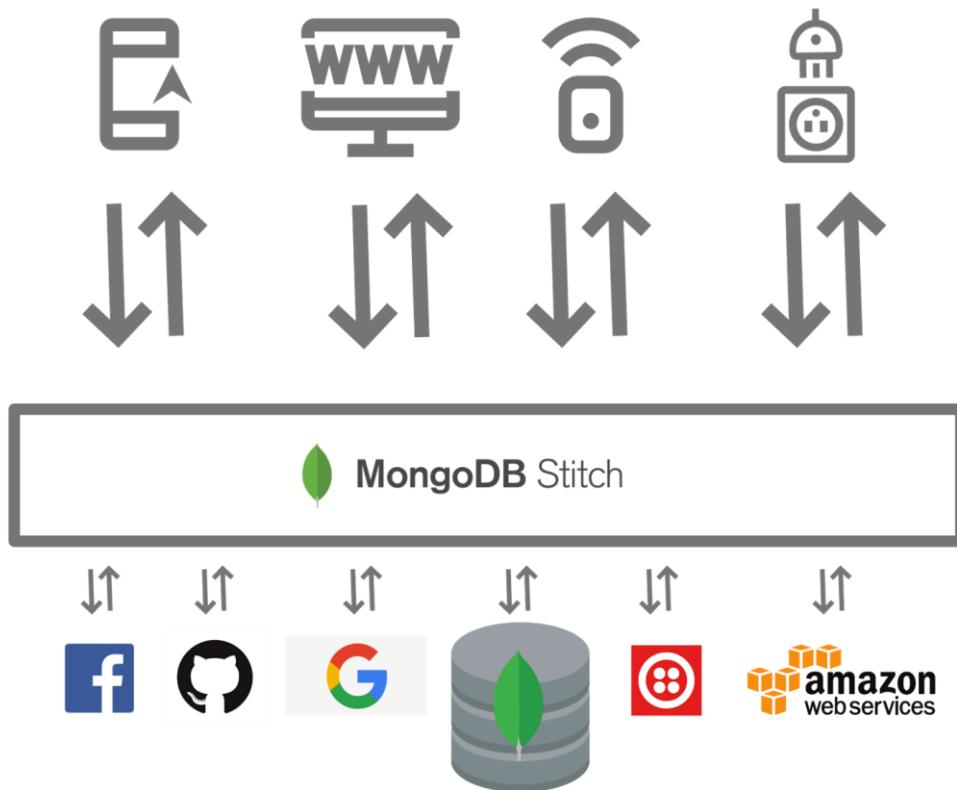
Application Logic

Services and APIs



Cloud Infrastructure

MongoDB Stitch Serverless Platform



Streamlines app development with simple, secure access to data and services from the client with thousands of lines less code to write and no infrastructure to manage.

Getting your apps to market faster while reducing operational costs.

MongoDB Stitch Serverless Platform – Services



Stitch QueryAnywhere

Brings MongoDB's rich query language safely to the edge

iOS, Android, Web, IoT



Stitch Functions

Integrate microservices + server-side logic + cloud services

Build full apps, or Data as a Service through custom APIs



Stitch Triggers

Real-time notifications let your application functions react in response to database changes, as they happen



Stitch Mobile Sync

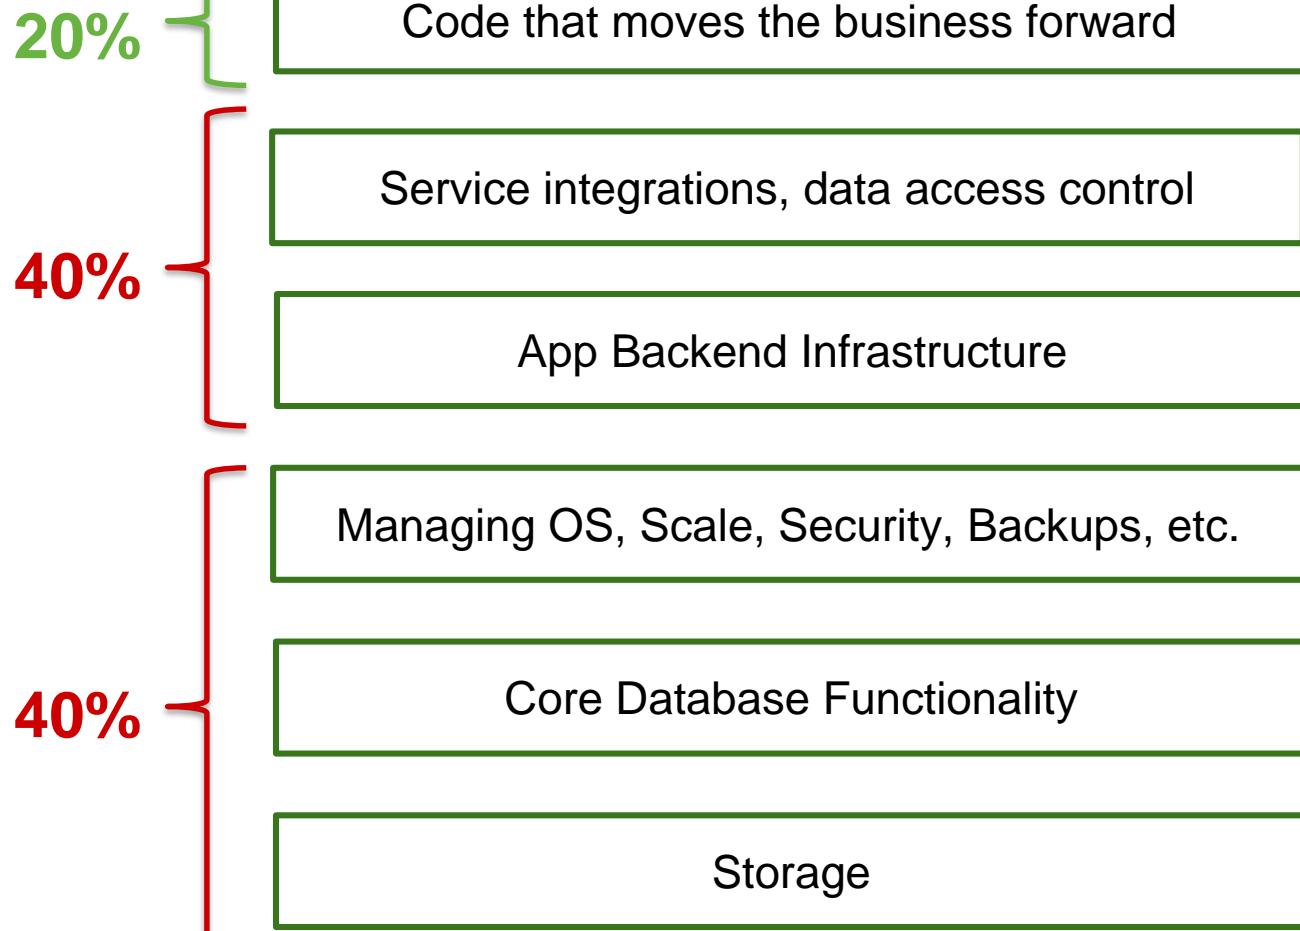
Automatically synchronizes data between documents held locally in MongoDB Mobile and your backend database

(beta)

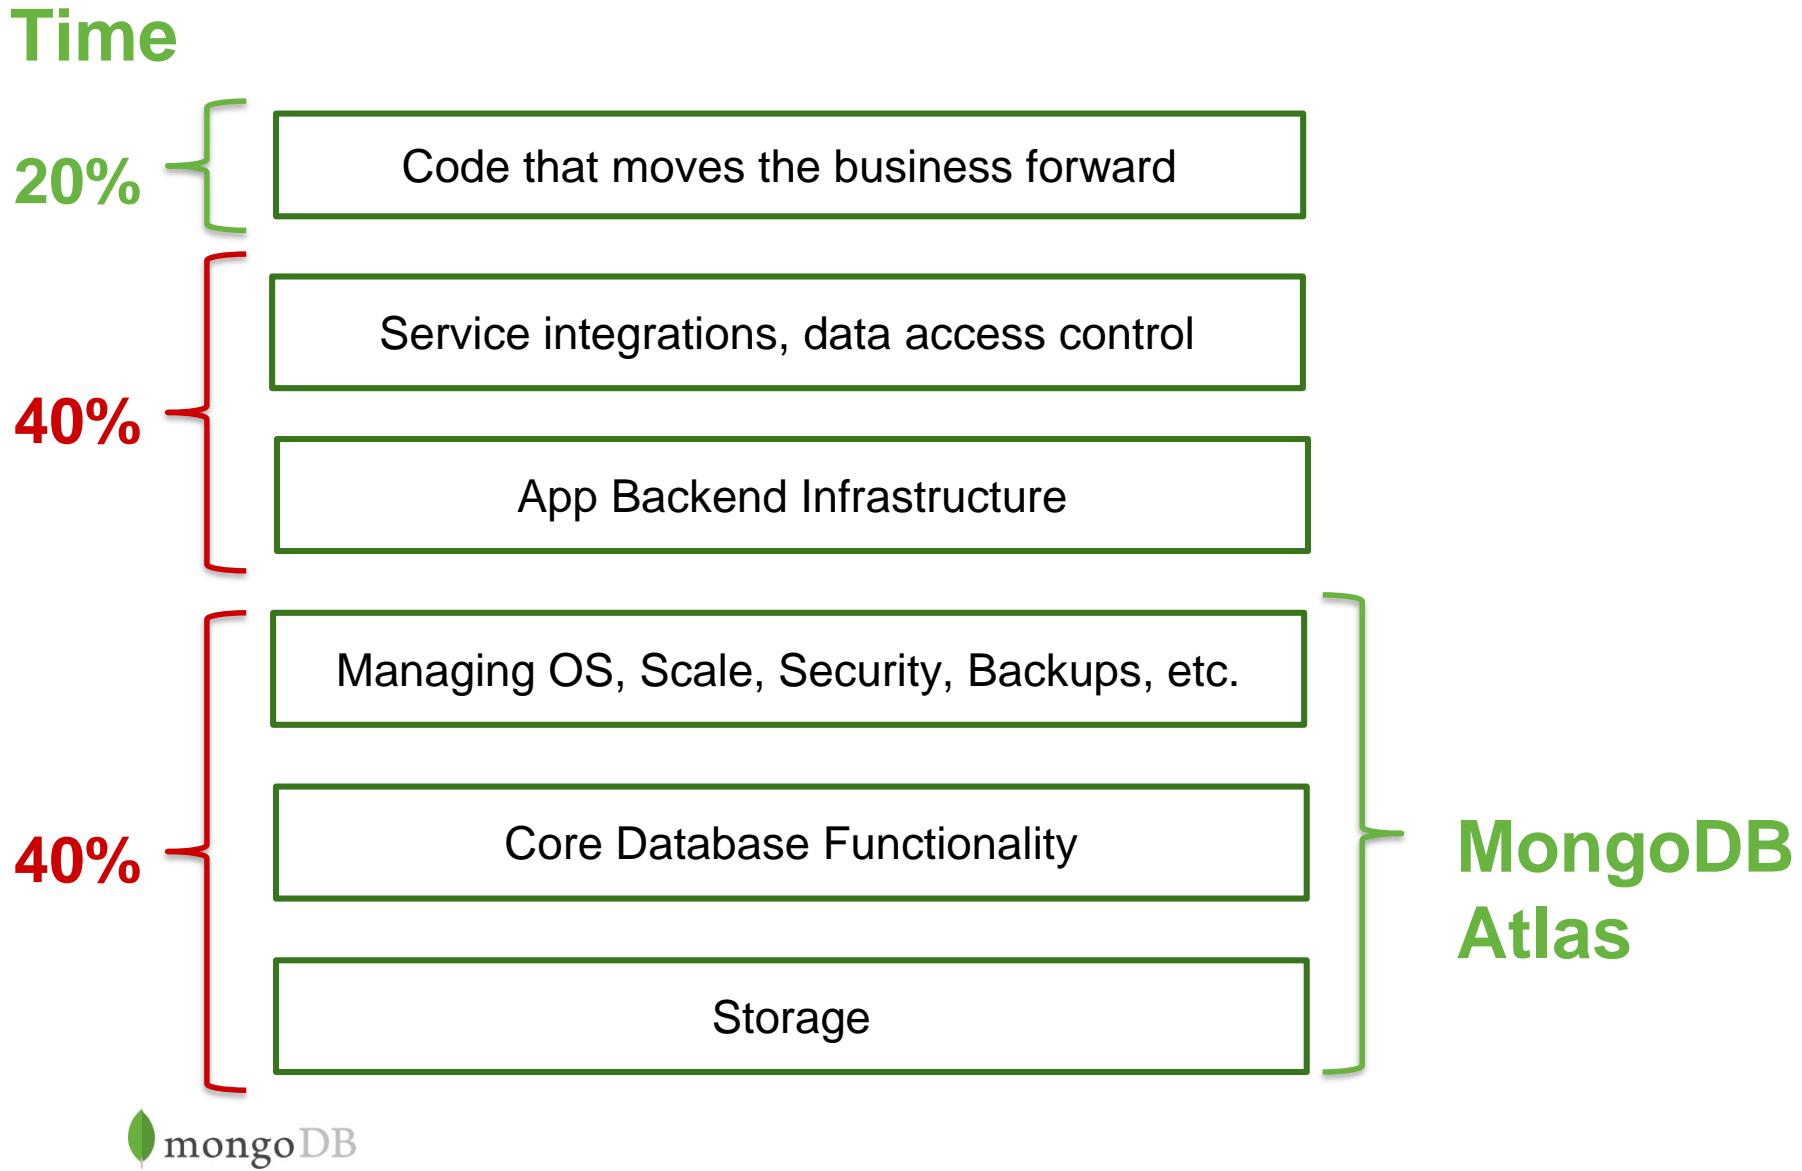
Streamlines app development with simple, secure access to data and services from the client with thousands of lines less code to write and no infrastructure to manage – getting your apps to market faster while reducing operational costs.

Focus Your Energy Where You Can Make a Difference

Time



Focus Your Energy Where You Can Make a Difference



Focus Your Energy Where You Can Make a Difference

Time

20%

Code that moves the business forward

40%

Service integrations, data access control

App Backend Infrastructure

40%

Managing OS, Scale, Security, Backups, etc.

Core Database Functionality

Storage

**MongoDB
Stitch**

**MongoDB
Atlas**

Fully managed
Elastic scale
Highly Available
Secure

Focus Your Energy Where You Can Make a Difference

Time

20%

Code that moves the business forward

You can focus here

40%

Service integrations, data access control

MongoDB
Stitch

App Backend Infrastructure

Fully managed
Elastic scale
Highly Available
Secure

40%

Managing OS, Scale, Security, Backups, etc.

MongoDB
Atlas

Core Database Functionality

Storage

Save weeks of development and thousands of lines of code

Backend

- Without Stitch
- Provision backend server
 - Install runtime environment
 - Add code to make backend HA
 - Add code to scale backend
 - Monitor & manage backend infrastructure
 - Code REST API for frontend to use backend
 - Code backend application logic

Data Access

- Code user authentication
- Code data access controls

Frontend

- Code application frontend
- Code against each external service API
- Continuously poll database for changes

With Stitch

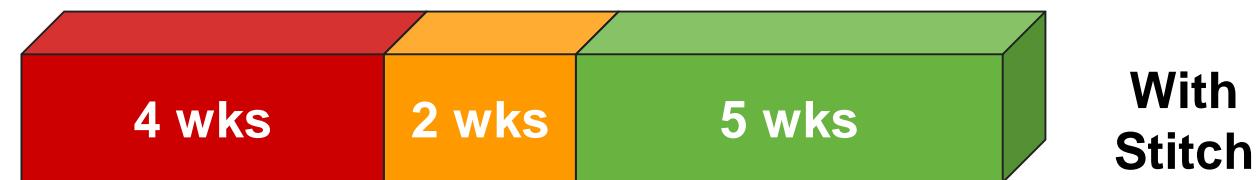
Handled automatically by Stitch and Atlas

Provide JS code for Stitch Functions

Simple JSON Config

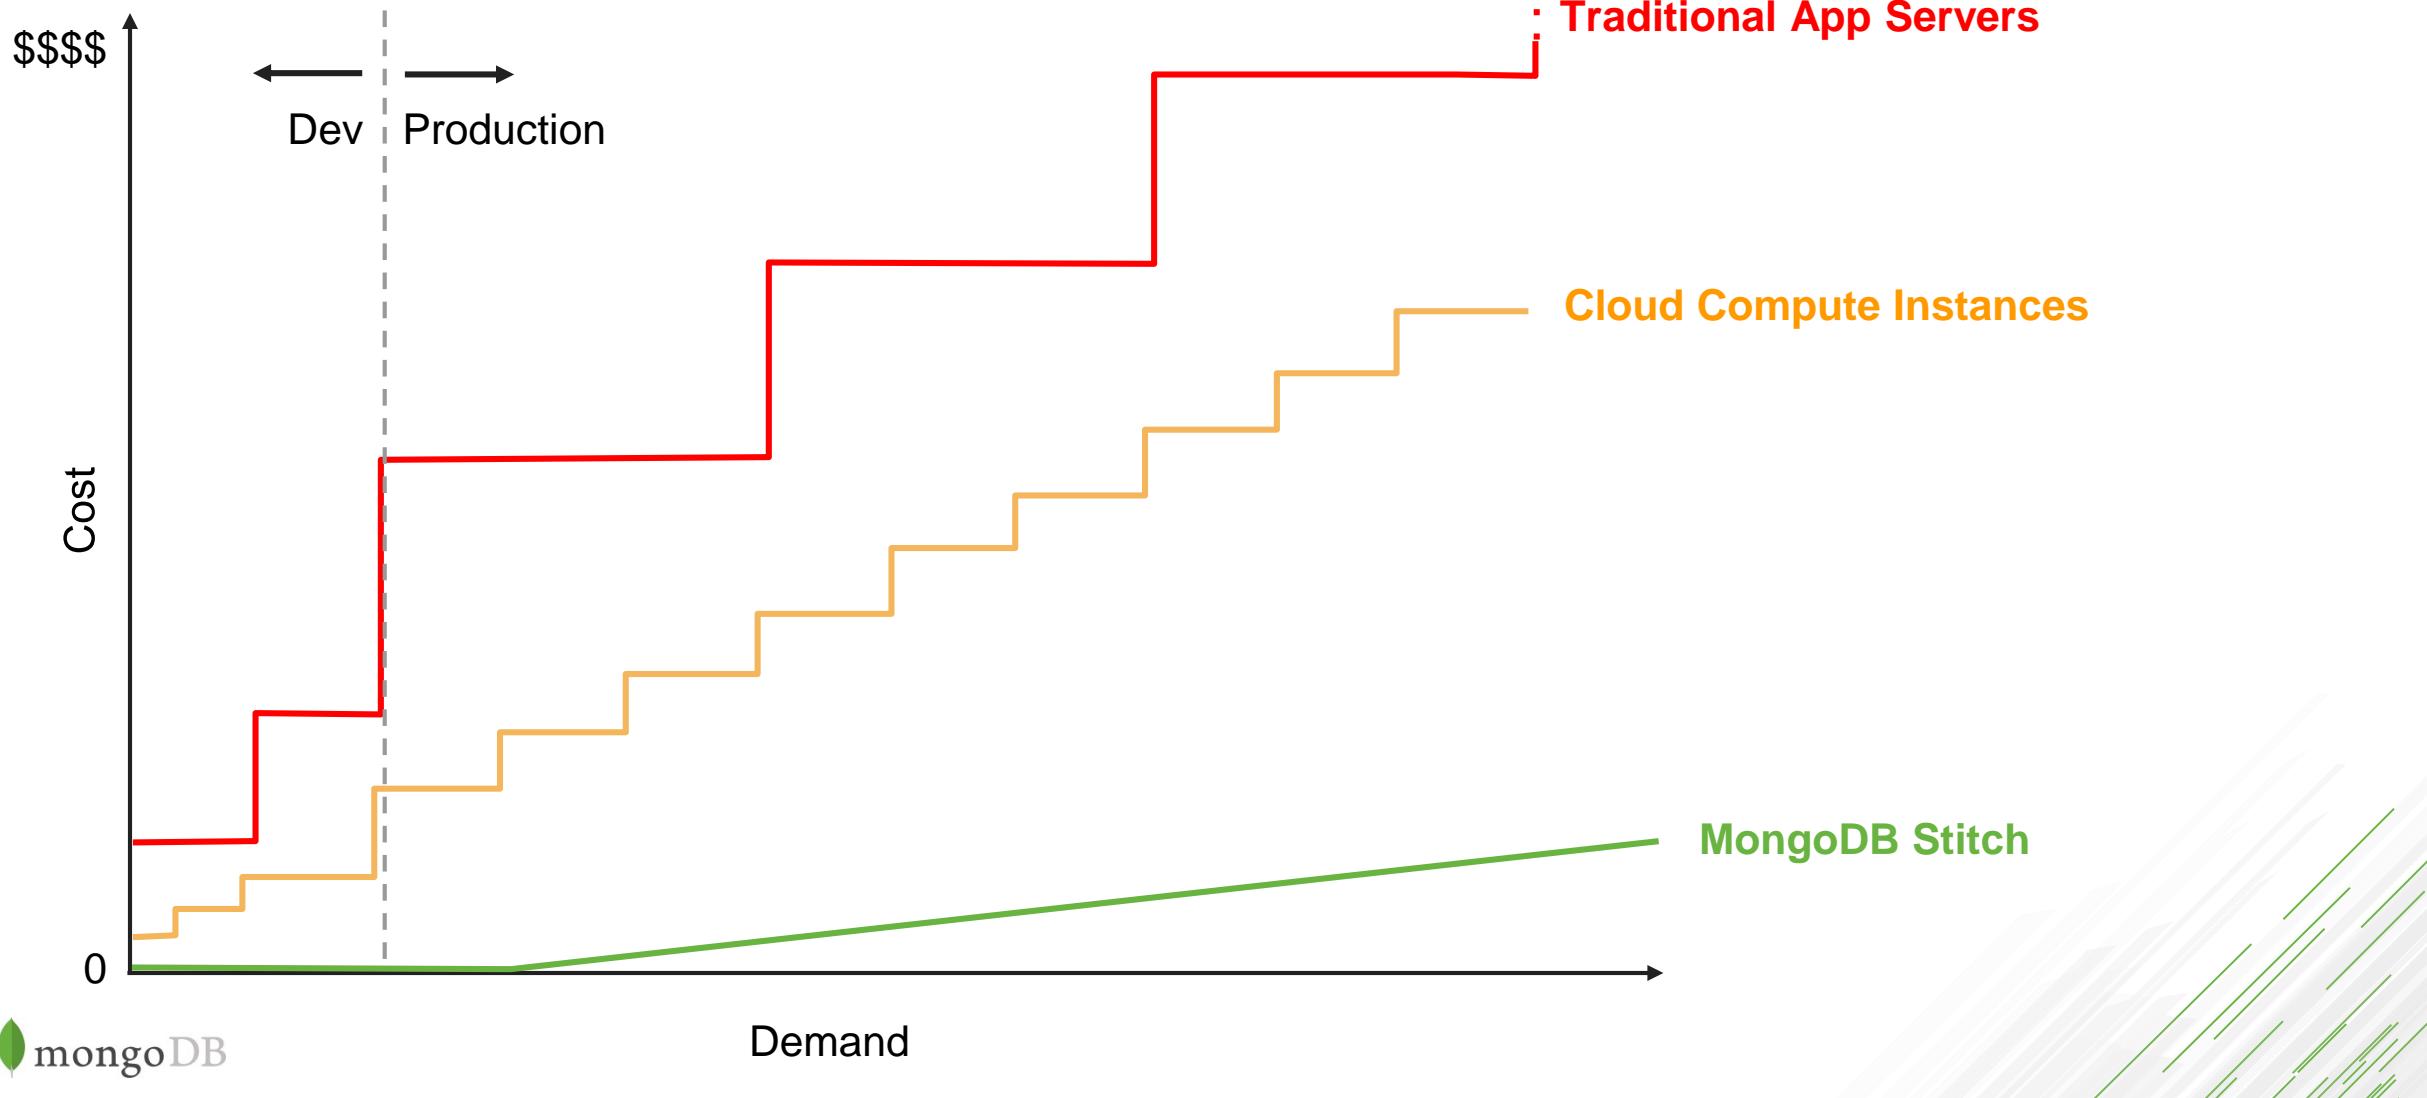
Code frontend using single SDK/API to access backend services

Development Time With/Without Stitch



Cut development time in half

MongoDB Stitch Reduces Infrastructure Costs



Summary: MongoDB **Uniquely** Delivers...

ACID
transactional
guarantees
of relational
databases

Documents,
the best way
to work with
data

Freedom
to
Run Anywhere

Intelligent
distributed
systems design

APPENDIX

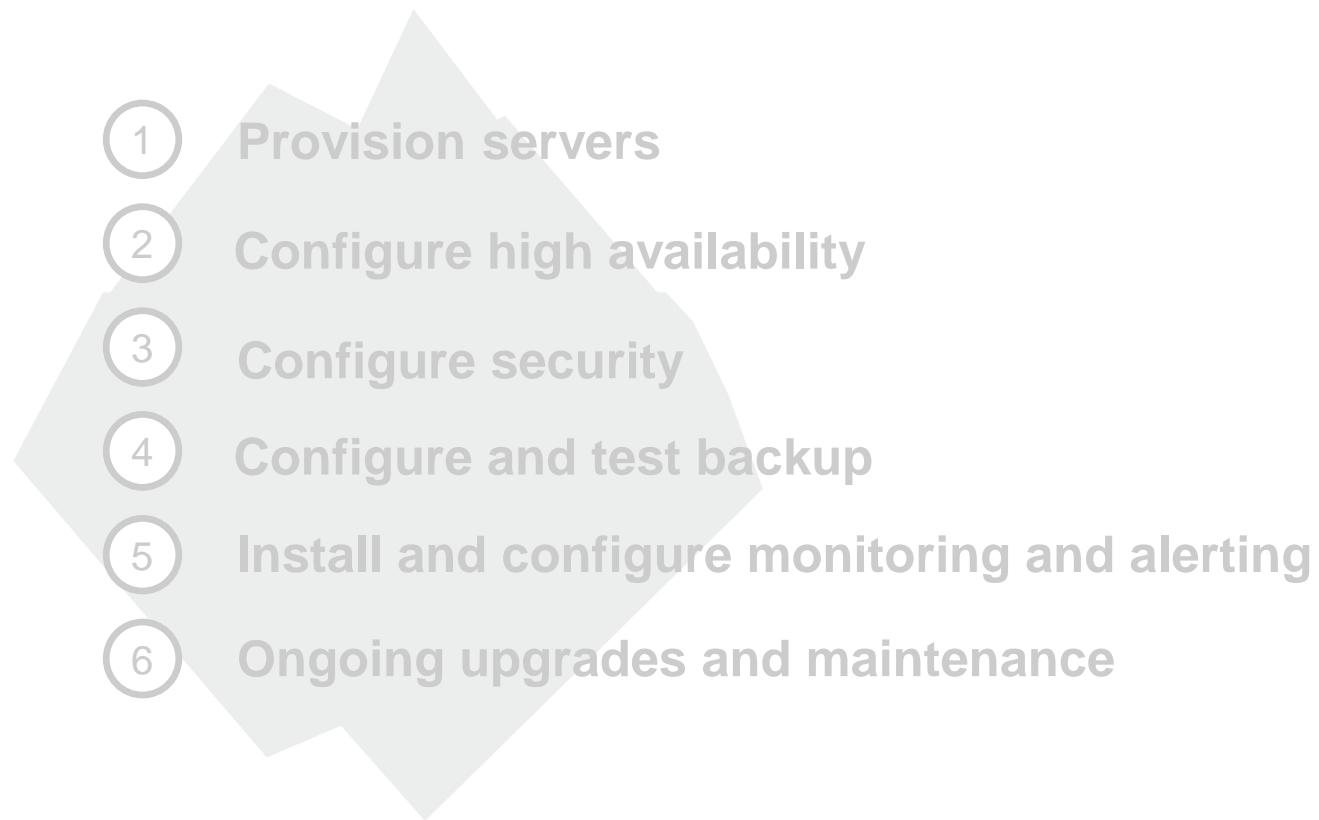


Hidden Operational Costs

Complicated
Operational Process

- 
- 1 Provision servers
 - 2 Configure high availability
 - 3 Configure security
 - 4 Configure and test backup
 - 5 Install and configure monitoring and alerting
 - 6 Ongoing upgrades and maintenance

The HIDDEN COST of resources

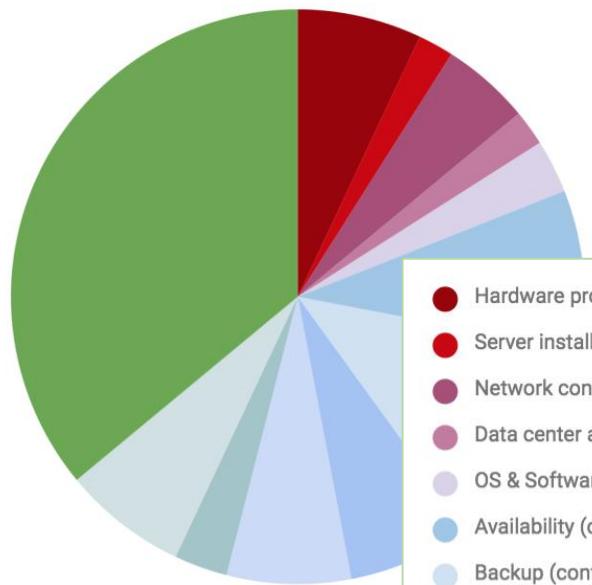


PEOPLE COSTS

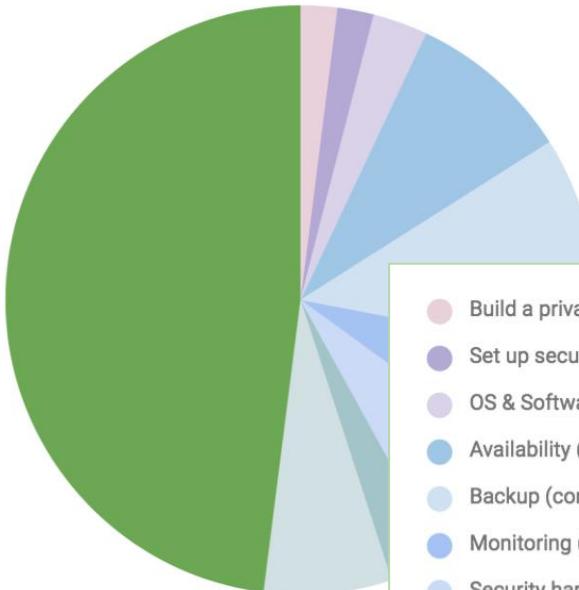
-  **DBAs**
~ \$150,000
-  **SysAdmin**
~ \$100,000
-  **IT Security**
~ \$130,000
-  **Developers**
~ \$250,000

Time spent on operations with different deployment models

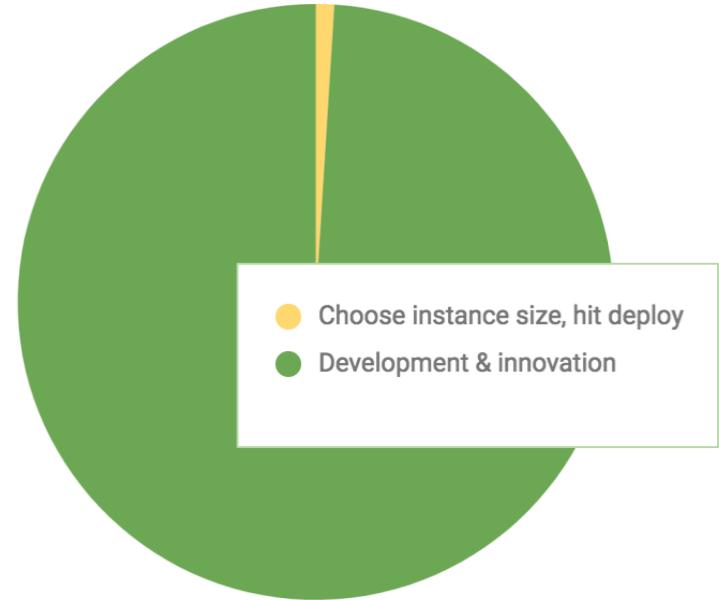
On-premises



Self-managed in the cloud

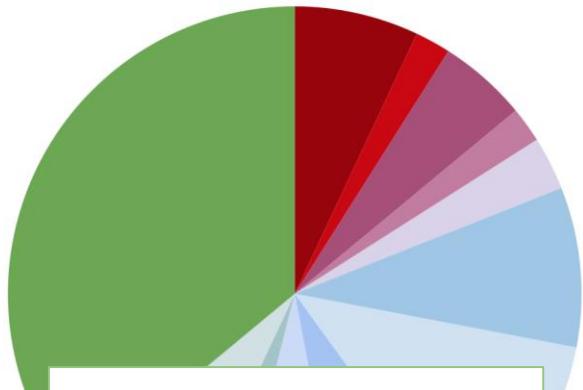


Database as a service



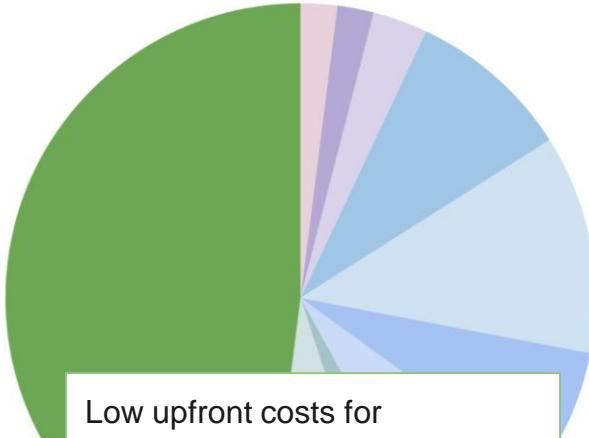
Operational & resource costs

On-premises



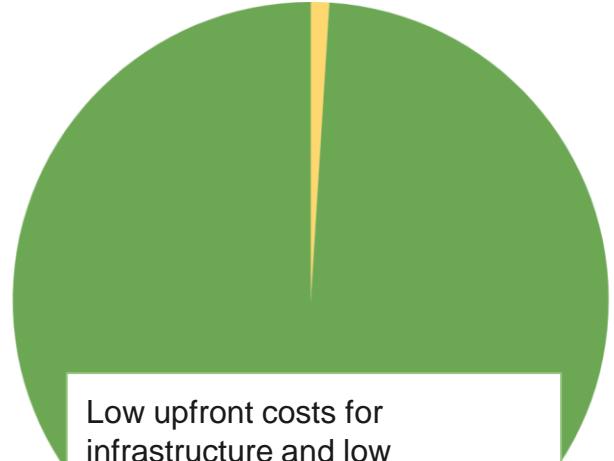
High upfront costs on infrastructure
High resource costs for installation and ongoing maintenance
Millions in operational and resource costs

Self-managed in the cloud



Low upfront costs for infrastructure
High resource costs for installation and ongoing maintenance
Significant resource costs as operational overhead is simply transferred to the cloud

Database as a service



Low upfront costs for infrastructure and low operational overhead
Self-service, elastic, and consumption-based service reduces TCO and accelerates time to value.

Aggregations

Advanced data processing pipeline for transformations and analytics using an Intelligent Operational Data Platform

```
db.orders.aggregate([
  {$match: {status: 'A'}},
  {
    $group: {
      _id: '$cust_id',
      total: {$sum: '$amount'}
    }
  },
  {
    $project: {
      customerId: '$_id',
      total: 1,
      _id: 0
    }
  }
])
```

Framework for Data Aggregation

- multiple stages

Similar to a unix pipe

- Build complex pipeline by chaining commands together

Rich Expressions

Examples using JavaScript shell

Fast at Scale

Cluster Scale		Performance Scale		Data Scale	
Entertainment Company	1,400 servers		250 Million Ticks / Sec		Petabytes
Asian Internet Company	1,000+ servers		300k Ops / Sec		10s of billions of objects
	250+ servers	Federal Agency	500k Ops / Sec		200 billion documents / 1+ PB of data

Slide Executions with Different Colored Backgrounds
Select best option that matches your projector



Flexible: Adapt to change

- **Polymorphic:** documents in a collection can contain different fields
- **Dynamic:** Modify schema without downtime: new app features or data sources
- **Govern:** JSON Schema to enforce structure

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752, -0.109081]}},  
+    {...}  
  ],  
  "phone" : [  
    { "location" : "work",  
      "number" : "+44-1234567890"},  
+    {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```



MongoDB Version 2

```
save(Map m)
{
    collection.insert(m);
}

Map fetch(String id)
{
    Map m = null;
    DBObject dbo = new BasicDBObject();
    dbo.put("id", id);
    c = collection.find(dbo);
    if(c.hasNext())
        m = (Map) c.next();
    }
    return m;
}
```



NO CHANGE

Advantages:

1. Zero time and money spent on overhead code
2. Code and database not physically linked
3. New material with more fields can be added into existing collections; backfill is optional
4. Names of fields in database precisely match key names in code layer and directly match on name, not indirectly via positional offset
1. No technical debt is created



MongoDB Version 3

```
save(Map m)
{
    collection.insert(m);
}

Map fetch(String id)
{
    Map m = null;
    DBObject dbo = new BasicDBObject();
    dbo.put("id", id);
    c = collection.find(dbo);
    if(c.hasNext())
        m = (Map) c.next();
    }
    return m;
}
```



NO CHANGE

Advantages:

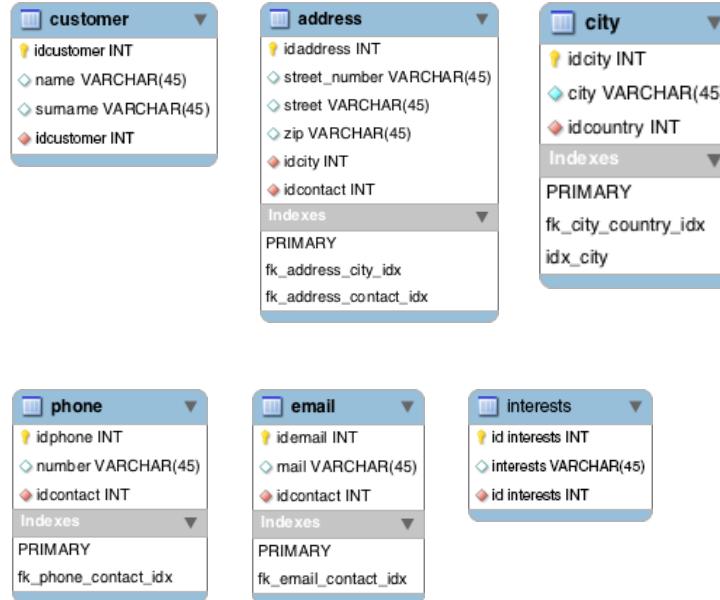
1. Zero time and money spent on overhead code
2. Code and database not physically linked
3. New material with more fields can be added into existing collections; backfill is optional
4. Names of fields in database precisely match key names in code layer and directly match on name, not indirectly via positional offset
1. No technical debt is created



Fast: To work with data

Compared to storing data across multiple tables, a single document data structure:

- Presents a single place for the database to read and write data
- Denormalized data eliminates JOINs for most operational queries
- Simplifies query development and optimization



```
_id: 12345678
> name: Object
> address: Array
> phone: Array
email: "john.doe@mongodb.com"
dob: 1966-07-30 01:00:00:000
▼ interests:Array
  0: "Cycling"
  1: "IoT"
```

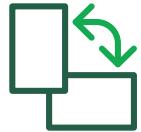


Versatile: Rich query functionality

Expressive Queries	<ul style="list-style-type: none">Find anyone with phone # “1-212...”Check if the person with number “555...” is on the “do not call” list
Geospatial	<ul style="list-style-type: none">Find the best offer for the customer at geo coordinates of 42nd St. and 6th Ave
Text Search	<ul style="list-style-type: none">Find all tweets that mention the firm within the last 2 days
Aggregation	<ul style="list-style-type: none">Count and sort number of customers by city, compute min, max, and average spend
Native Binary JSON Support	<ul style="list-style-type: none">Add an additional phone number to Mark Smith’s record without rewriting the documentUpdate just 2 phone numbers out of 10Sort on the modified date
JOIN (\$lookup)	<ul style="list-style-type: none">Query for all San Francisco residences, lookup their transactions, and sum the amount by person
Graph Queries (\$graphLookup)	<ul style="list-style-type: none">Query for all people within 3 degrees of separation from Mark

MongoDB

```
{  
  customer_id : 1,  
  first_name : "Mark",  
  last_name : "Smith",  
  city : "San Francisco",  
  phones: [      {  
    number : "1-212-777-1212",  
    type : "work"  
  },  
  {  
    number : "1-212-777-1213",  
    type : "cell"  
  }]  
..... . . . }
```



Single vs Multi-Document ACID Transactions

For many apps,
single document
transactions meet the
majority of needs

```
_id: 12345678
> name: Object
> address: Array
> phone: Array
email: "john.doe.mongodb.com"
dob: 1966-07-30 01:00:00:000
▼ interests:Array
  0: "Cycling"
  1: "IoT"
```

Related data modeled in a single, rich document against
which ACID guarantees are applied



Syntax*

```
with client.start_session() as s:  
    s.start_transaction()  
    try:  
        collection.insert_one(doc1, session=s)  
        collection.insert_one(doc2, session=s)  
        s.commit_transaction()  
    except Exception:  
        s.abort_transaction()
```

Natural for developers

- Idiomatic to the programming language
- Familiar to relational developers
- Simple

*Python. Syntax is subject to change



Comparing Syntax with Relational



```
db.start_transaction()  
cursor.execute(orderInsert, orderData)  
cursor.execute(stockUpdate, stockData)  
db.commit()
```



```
s.start_transaction()  
orders.insert_one(order, session=s)  
stock.update_one(item, stockUpdate, session=s)  
s.commit_transaction()
```

Write Everywhere:

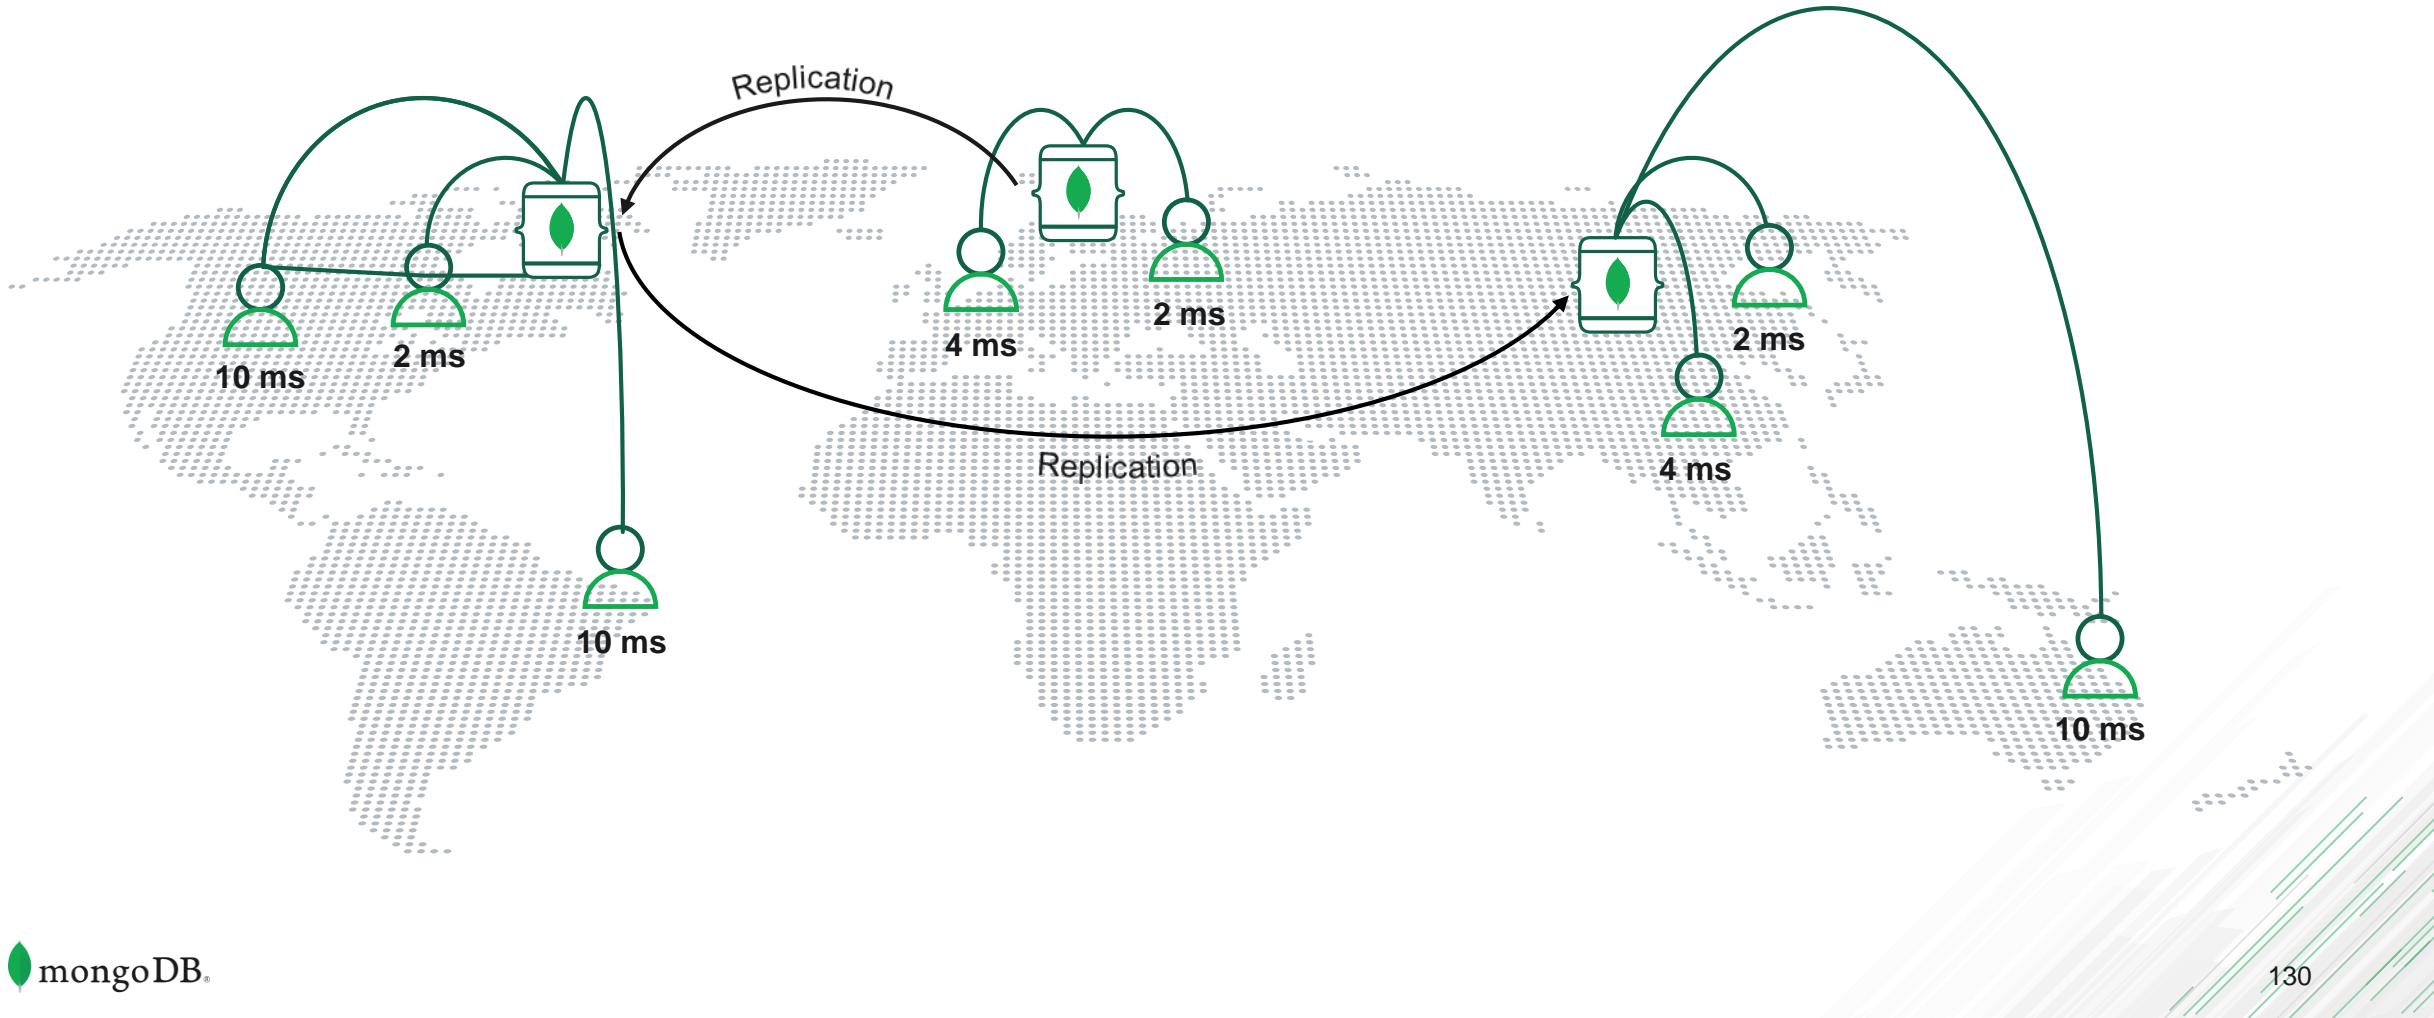
zoned sharding routes data to the regions where it's needed



Summary: MongoDB Uniquely Delivers...



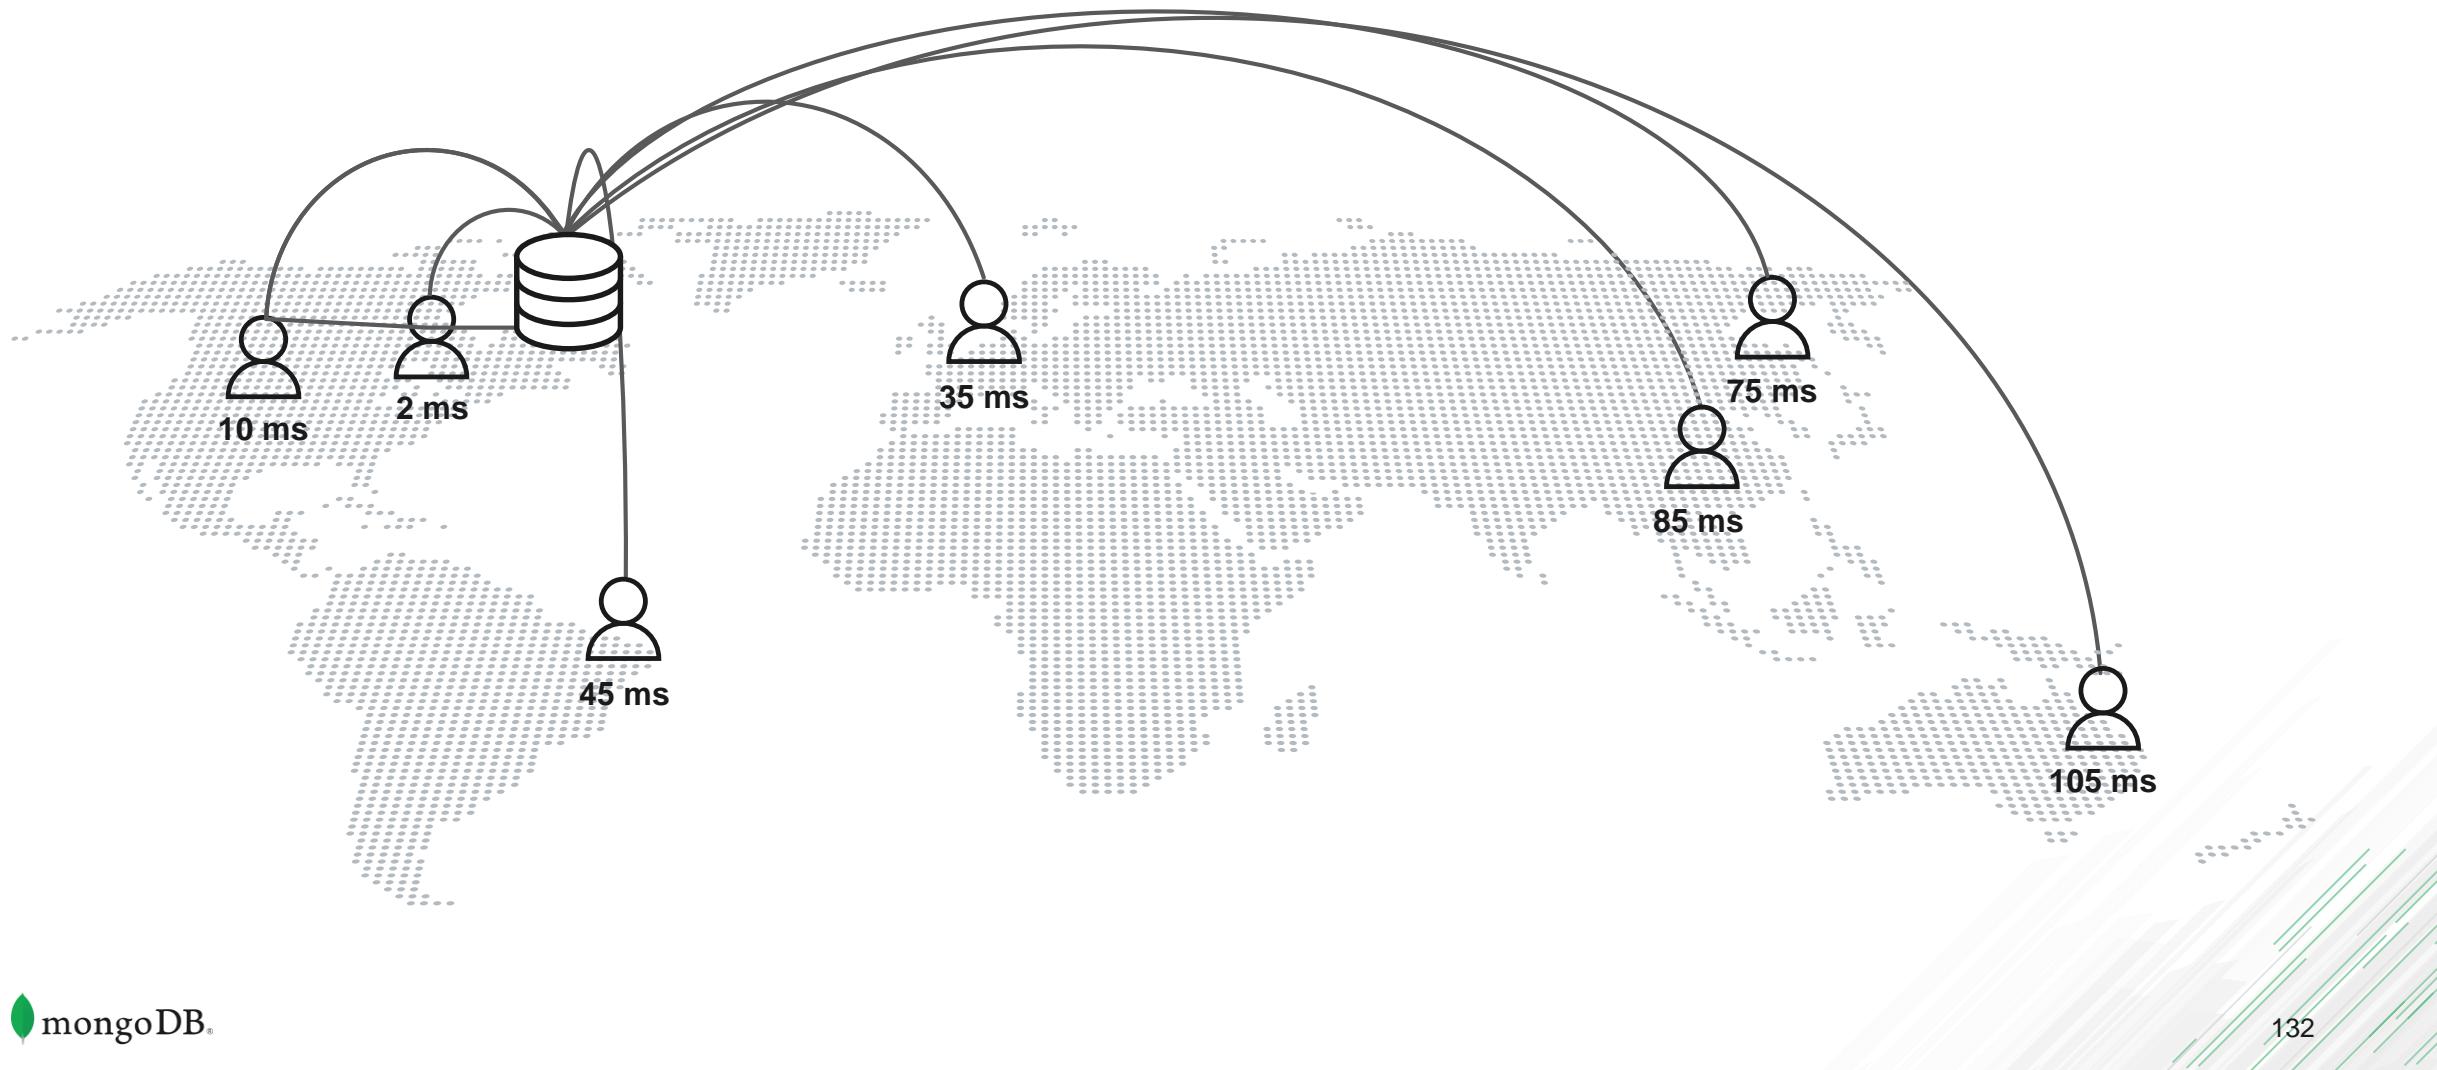
MongoDB's native replica sets puts your entire database right next to users



MongoDB's native replica sets puts your entire database right next to users



Serving global audiences with relational databases



Serving global audiences with relational databases



Write Everywhere:

zoned sharding routes data to the regions where it's needed

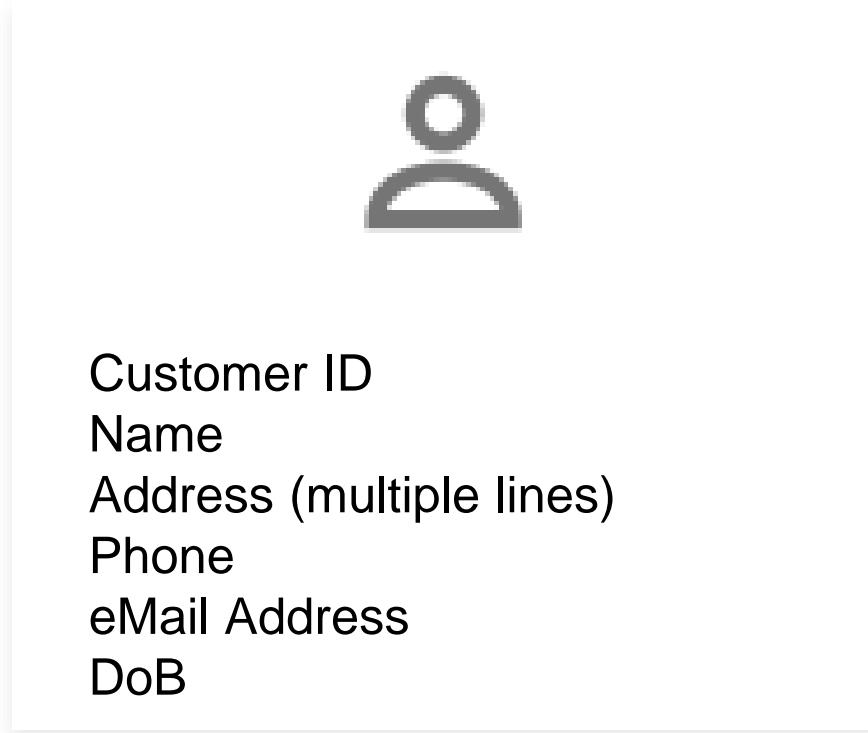


Write Everywhere:

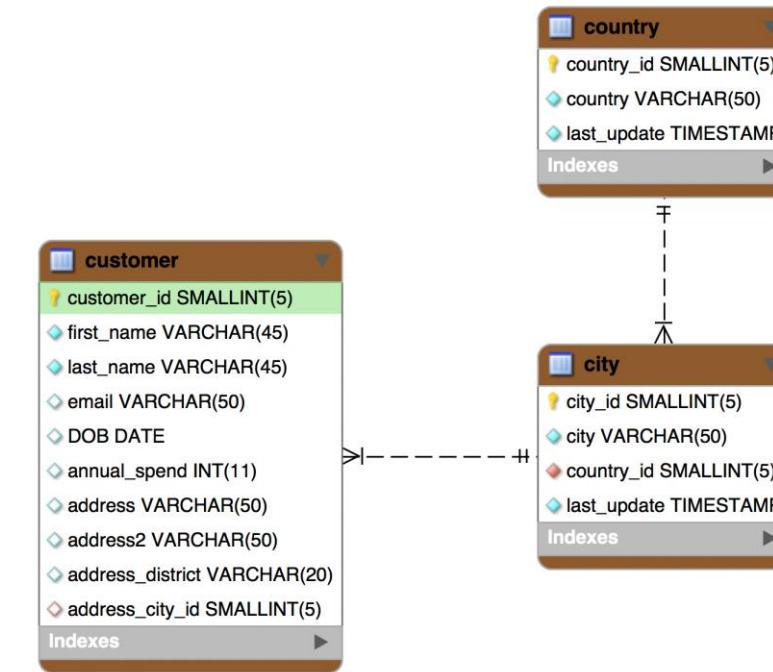
zoned sharding routes data to the regions where it's needed



Modeling a customer in a relational database

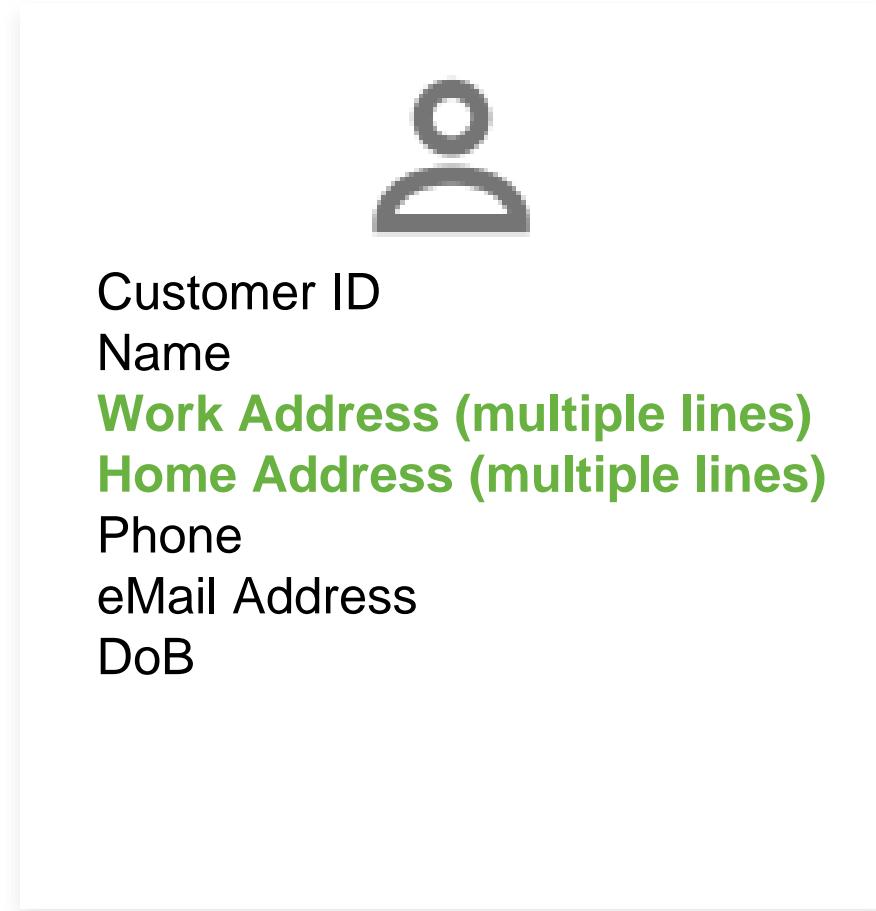


Simple Customer Data

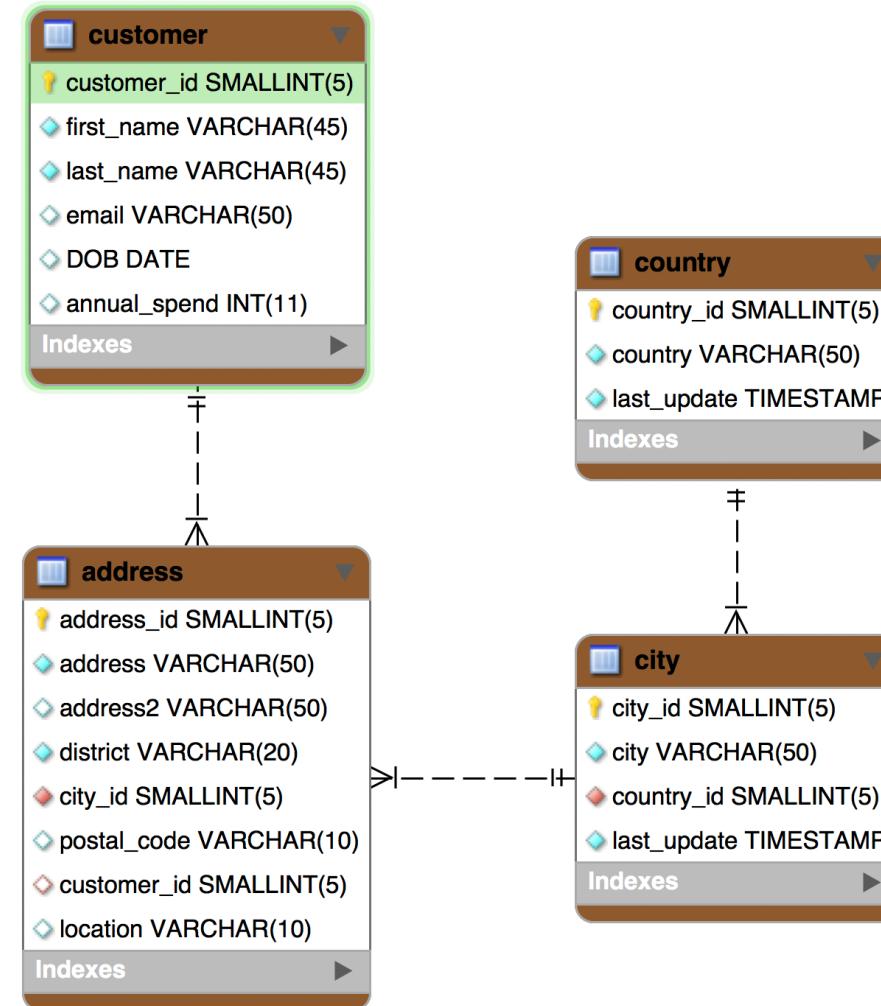


Relational Data Model

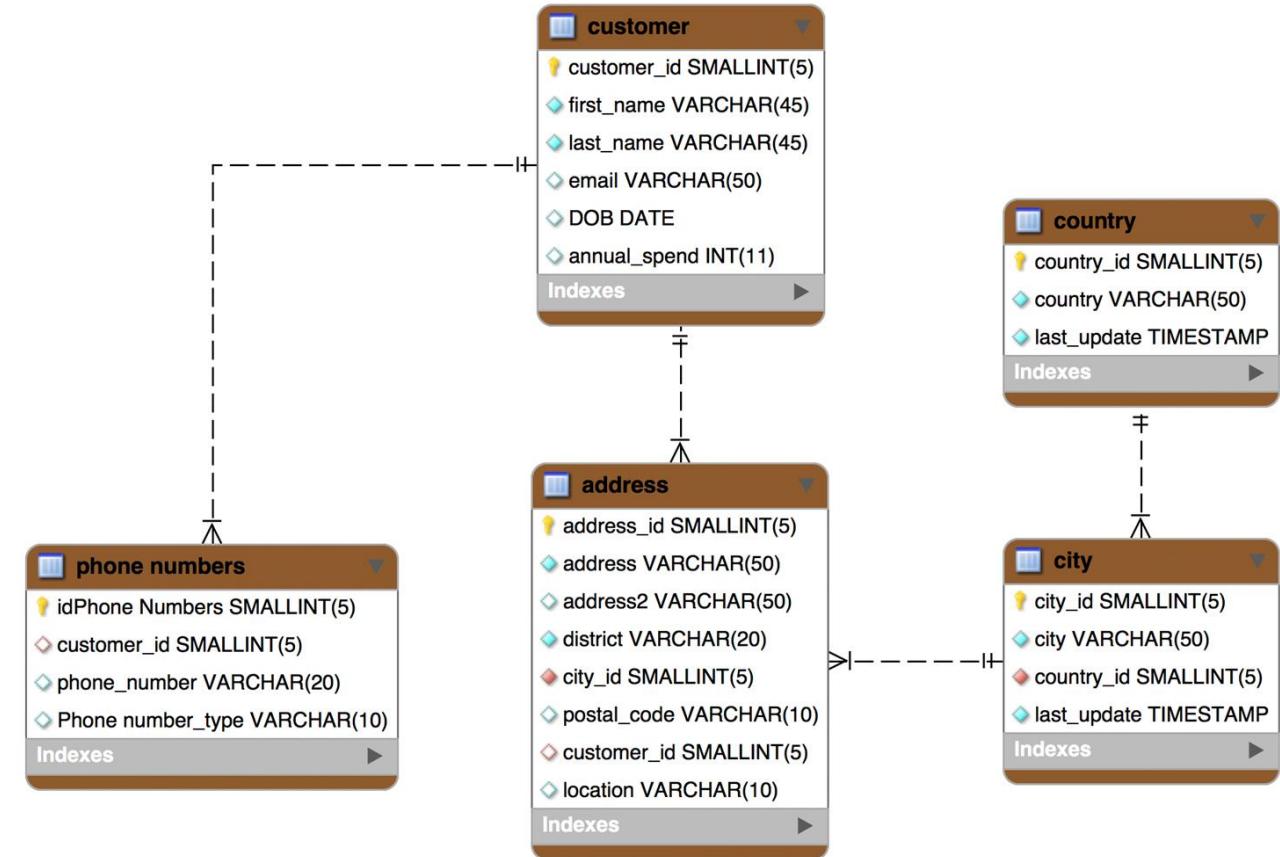
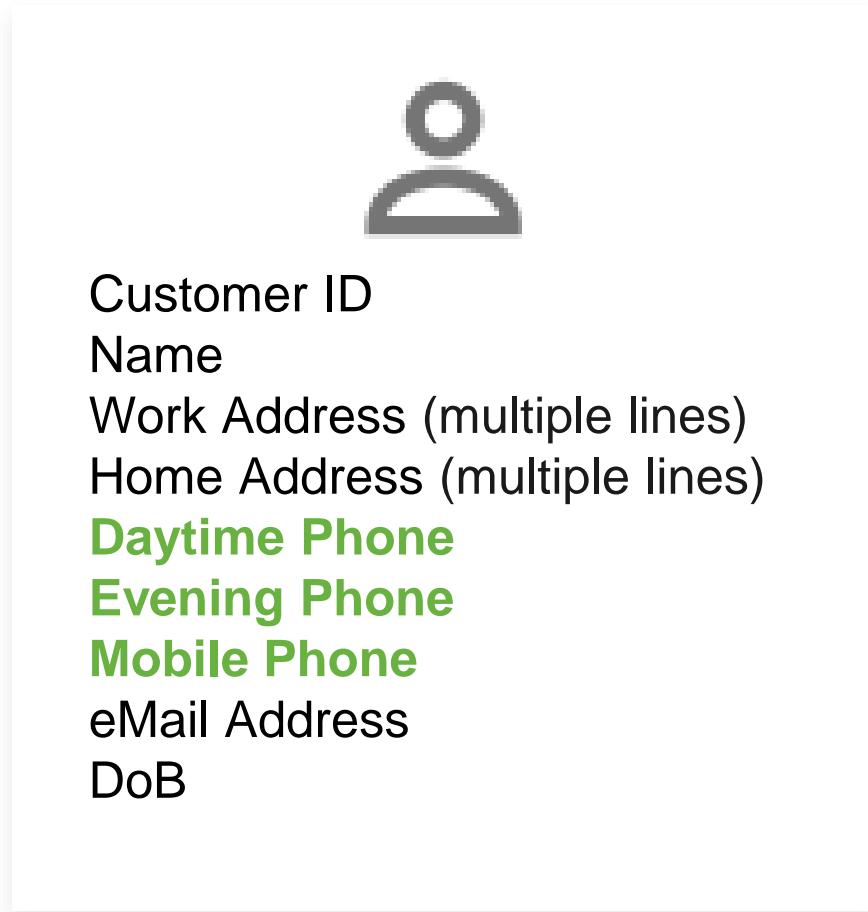
Adding basic customer data requires a new table



Enriched Customer Data



Enriching data further complicates the data model



Quickly, the relational model becomes unwieldy

